

SQL: The Query Language (Part 1)

Fall 2017

Life is just a bowl of queries.

-Anon

1

Relational Query Languages

- **A major strength of the relational model: supports simple, powerful *querying* of data.**
- **Query languages can be divided into two parts**
 - Data Manipulation Language (DML)
 - Allows for queries and updates
 - Data Definition Language (DDL)
 - Define and modify schema (at all 3 levels)
 - Permits database tables to be created or deleted. It also defines indexes (keys), specifies links between tables, and imposes constraints between tables
- **The DBMS is responsible for efficient evaluation.**
 - The key: precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.
 - Internal cost model drives use of indexes and choice of access paths and physical operators.

2

The SQL Query Language

- The most widely used relational query language.
- Originally IBM, then ANSI in 1986
- **Current standard is SQL-2008**
 - 2003 was last major update: XML, window functions, sequences, auto-generated IDs.
 - Not fully supported yet
- SQL-1999 Introduced “Object-Relational” concepts.
 - Also not fully supported yet.
- **SQL92 is a basic subset**
 - Most systems support at least this
- PostgreSQL has some “unique” aspects (as do most systems).
- SQL is not synonymous with Microsoft’s “SQL Server”

3

The SQL DML

- Single-table queries are straightforward.
- To find all 18 year old students, we can write:

```
SELECT *
  FROM Students
 WHERE age=18
```

```
SELECT *
  FROM Students
 WHERE Students.age=18
```

```
SELECT *
  FROM Students S
 WHERE S.age=18
```

4

The SQL DML

- Single-table queries are straightforward.
- To find all 18 year old students, we can write:

```
SELECT *
  FROM Students S
 WHERE S.age=18
```

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

5

Querying Multiple Relations

- Can specify a join over two tables as follows:

```
SELECT S.name, E.cid
  FROM Students S, Enrolled E
 WHERE S.sid=E.sid AND E.grade='B'
```

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

result =

S.name	E.cid
Jones	History105

Note: obviously no referential integrity constraints have been used here.

6

Basic SQL Query

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- *relation-list*: A list of relation names
 - possibly with a *range-variable* after each name
- *target-list*: A list of attributes of tables in *relation-list*
- *qualification*: Comparisons combined using AND, OR and NOT.
 - Comparisons are Attr *op* const or Attr1 *op* Attr2, where *op* is one of = ≠ < > ≤ ≥
- *DISTINCT*: optional keyword indicating that the answer should not contain duplicates.
 - In SQL SELECT, the default is that duplicates are *not* eliminated! (Result is called a “multiset”)

Query Semantics

- Semantics of an SQL query are defined in terms of the following conceptual evaluation strategy:
 1. do FROM clause: compute *cross-product* of tables (e.g., Students and Enrolled).
 2. do WHERE clause: Check conditions, discard tuples that fail. (i.e., “selection”).
 3. do SELECT clause: Delete unwanted fields. (i.e., “projection”).
 4. If DISTINCT specified, eliminate duplicate rows.

Probably the least efficient way to compute a query!

- An optimizer will find more efficient strategies to get the *same answer*.

Step 1 – Cross Product

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

9

Step 2) Discard tuples that fail predicate

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

```
SELECT S.name, E.cid
  FROM Students S, Enrolled E
 WHERE S.sid=E.sid AND E.grade='B' 10
```

Step 3) Discard Unwanted Columns

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

```
SELECT S.name, E.cid
  FROM Students S, Enrolled E
 WHERE S.sid=E.sid AND E.grade='B' 11
```

Null Values

- **Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).**
 - SQL provides a special value *null* for such situations.
- **The presence of *null* complicates many issues. E.g.:**
 - Special operators needed to check if value is/is not *null*.
 - Is *rating>8* true or false when *rating* is equal to *null*? What about AND, OR and NOT connectives?
 - We need a *3-valued logic* (true, false and *unknown*).
 - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, *outer joins*) possible/needed.

Null Values – 3 Valued Logic

(null > 0)	is null
(null + 1)	is null
(null = 0)	is null
null AND true	is null

AND	T	F	Null
T	T	F	Null
F	F	F	F
NULL	Null	F	Null

OR	T	F	Null
T	T	T	T
F	T	F	Null
NULL	T	Null	Null

13

Now the Details

We will use these instances of relations in our examples.

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

Sailors

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

14

Example Schemas (in SQL DDL)

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

CREATE TABLE Sailors (sid INTEGER, sname
CHAR(20), rating INTEGER, age REAL,
PRIMARY KEY sid)

Consider the use
Of VARCHAR instead

15

Example Schemas (in SQL DDL)

bid	bname	color	
101	Interlake	blue	
102	Interlake	red	
103	Clipper	green	
104	Marine	red	

CREATE TABLE Boats (bid INTEGER, bname CHAR (20),
color CHAR(10)
PRIMARY KEY bid)

16

Example Schemas (in SQL DDL)

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

```
CREATE TABLE Reserves (sid INTEGER, bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY sid REFERENCES Sailors,
    FOREIGN KEY bid REFERENCES Boats)
```

17

Another Join Query

```
SELECT sname
    FROM Sailors, Reserves
   WHERE Sailors.sid=Reserves.sid
         AND bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
95	Bob	3	63.5	22	101	10/10/96
95	Bob	3	63.5	95	103	11/12/96

18

Some Notes on Range Variables

- Can associate “range variables” with the tables in the FROM clause.
 - saves writing, makes queries easier to understand
- Needed when ambiguity could arise.
 - for example, if same table used multiple times in same FROM (called a “self-join”)

```
SELECT sname  
FROM Sailors,Reserves  
WHERE Sailors.sid=Reserves.sid AND bid=103
```

Can be
rewritten using
range variables as:

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid AND bid=103
```

More Notes

- Here’s an example where range variables are required (self-join example):

```
SELECT x.sname, x.age, y.sname, y.age  
FROM Sailors x, Sailors y  
WHERE x.age > y.age
```

- Note that target list can be replaced by “*” if you don’t want to do a projection:

```
SELECT *  
FROM Sailors x  
WHERE x.age > 20
```

Find sailors who've reserved at least one boat

```
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid
```

- Would adding DISTINCT to this query make a difference?
- What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause?
 - Would adding DISTINCT to this variant of the query make a difference?

21

Expressions

- Can use arithmetic expressions in SELECT clause (plus other operations we'll discuss later)
- Use AS to provide column names

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname = 'dustin'
```

- Can also have expressions in WHERE clause:

```
SELECT S1.sname AS name1, S2.sname AS name2  
FROM Sailors S1, Sailors S2  
WHERE 2*S1.rating = S2.rating - 1
```

22

String operations

- SQL also supports some string operations
- “LIKE” is used for string matching.

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

‘_’ stands for any one character and ‘%’ stands for 0 or more arbitrary characters.

23

Find sid's of sailors who've reserved a red **or** a green boat

```
SELECT DISTINCT R.sid  
FROM Boats B,Reserves R  
WHERE R.bid=B.bid AND  
(B.color='red'OR B.color='green')
```

Vs.

(note:
UNION
eliminates
duplicates
by default.
Override w/
UNION ALL)

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
UNION SELECT R.sid  
      FROM Boats B, Reserves R  
      WHERE R.bid=B.bid AND  
            B.color='green'
```

24

Find sid's of sailors who've reserved a red **and** a green boat

- If we simply replace OR by AND in the previous query, we get the wrong answer. (Why?)

```
SELECT R1.sid
FROM Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE R1.sid=R2.sid
      AND R1.bid=B1.bid
      AND R2.bid=B2.bid
      AND (B1.color='red' AND B2.color='green')
```

25

AND Continued...

- **INTERSECT:**
 - Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- **EXCEPT**
 - (sometimes called MINUS)
 - Included in the SQL/92 standard, but **many** systems (including MySQL) don't support them.

Key field!

```
SELECT S.sid
FROM Sailors S, Boats B,
      Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid
      AND B.color='red'

INTERSECT

SELECT S.sid
FROM Sailors S, Boats B,
      Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid
      AND B.color='green'
```

26

Nested Queries

- Powerful feature of SQL: WHERE clause can itself contain an SQL query!

– Actually, so can FROM and HAVING clauses.

Names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                 FROM Reserves R
                 WHERE R.bid=103)
```

- To find sailors who've *not* reserved #103, use NOT IN.
- To understand semantics of nested queries:
 - think of a *nested loops* evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
              FROM Reserves R
              WHERE R.bid=103 AND S.sid=R.sid)
```

Tests whether the set
is nonempty

- EXISTS is another set comparison operator, like IN.
- Can also specify NOT EXISTS
- If UNIQUE is used, and * is replaced by R.bid, finds sailors with at most one reservation for boat #103.
 - UNIQUE checks for duplicate tuples in a subquery;
- Subquery must be recomputed for each Sailors tuple.
 - Think of subquery as a function call that runs a query!

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op ANY*, *op ALL*
- Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY
      (SELECT S2.rating
       FROM Sailors S2
       WHERE S2.sname='Horatio')
```

29

Rewriting INTERSECT Queries Using IN

Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
      AND R.sid IN (SELECT R2.sid
                      FROM Boats B2, Reserves R2
                      WHERE R2.bid=B2.bid
                        AND B2.color='green')
```

- Similarly, EXCEPT queries re-written using NOT IN.
- How would you change this to find *names* (not *sid's*) of Sailors who've reserved both red and green boats?

30

Division in SQL

Find names of sailors who've reserved all boats.

```
SELECT S.sname  Sailors S such that ...
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
there is no boat B      FROM Boats B
                      WHERE NOT EXISTS (SELECT R.bid
that doesn't have ...          FROM Reserves R
                                         WHERE R.bid=B.bid
                                         AND R.sid=S.sid))  

a Reserves tuple showing S reserved B
```

Recall Exists Tests whether the set is nonempty

31

Division Operations in SQL (1)

Find names of sailors who've reserved all boat:

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXIST
```

((SELECT B.bid
FROM Boats B) EXCEPT

(SELECT R.bid
FROM Reserves R
WHERE R.sid = S.sid))

The sailor
reserved all boats

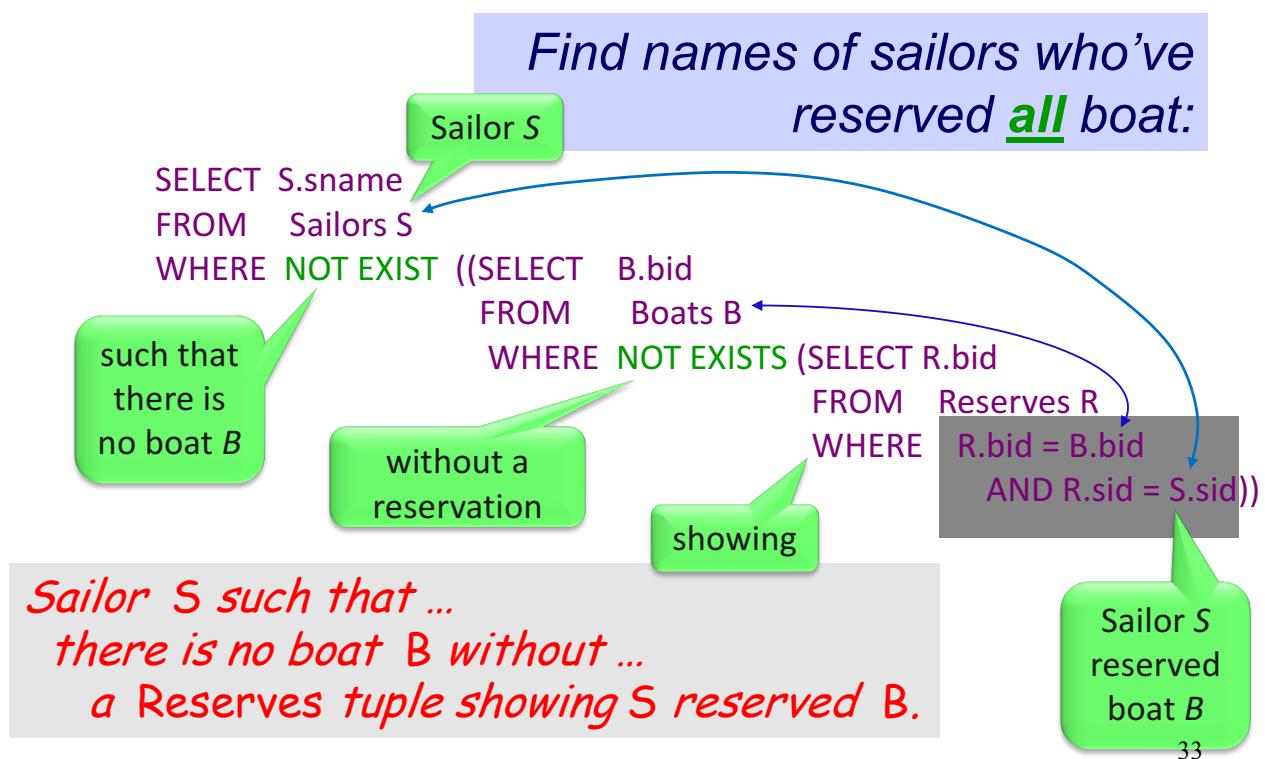
Boats not reserved
by the sailor

All boats

All boats reserved
by the sailor

32

Division Operations in SQL (2)



SQL Operators

- **BETWEEN**
- **NOT BETWEEN**
- **IN**
- **UNION [DISTINCT | ALL]**
- **EXCEPT**
 - Not Supported in MySQL

```
SELECT * FROM Products  
WHERE (Price BETWEEN 10 AND 20)  
AND NOT CategoryID IN (1,2,3);
```

More on Set-Comparison Operators

- ***op ANY, op ALL, where op: >, <, =, ≠, ≥, ≤***

Find sailors whose rating is greater than that of some sailor called Horatio

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating
                        FROM Sailors S2
                        WHERE S2.sname='Horatio')
```

>ANY means greater than at least one value
>ALL means greater than every value

35

ARGMAX?

- **The sailor with the highest rating**
 - what about ties for highest?!

```
SELECT *
FROM Sailors S
WHERE S.rating >= ALL
      (SELECT S2.rating
       FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
WHERE S.rating =
      (SELECT MAX(S2.rating)
       FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
ORDER BY rating DESC
LIMIT 1;
```

36

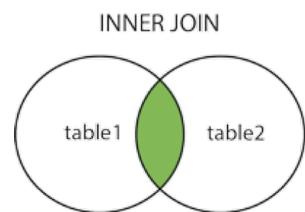
Joins

```
SELECT (column_list)
FROM table_name
[INNER | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
    ON qualification_list
WHERE ...
```

**Explicit join semantics needed unless it is an
INNER join (INNER is default)**

37

Inner Join



Selects records that have matching values in both tables.

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats

SQL-92 also allows:

```
SELECT s.sid, s.name, r.bid
FROM Sailors s NATURAL JOIN Reserves r
```

**"NATURAL" means equi-join for each pair of attributes with
the same name (may need to rename with "AS")**

38

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>s.sid</u>	<u>s.name</u>	<u>r.bid</u>
22	Dustin	101
95	Bob	103

39

Outer Joins

S1

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

No “sid = 31”



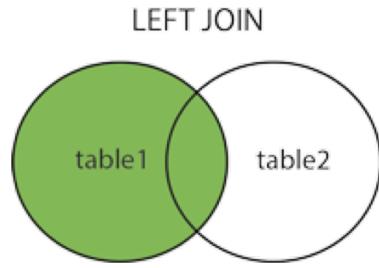
S1 $\bowtie R1$

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>	<u>bid</u>	<u>day</u>
22	dustin	7	45.0	101	10/10/96
31	lubber	8	55.55	null	null
58	rusty	10	35.0	103	11/12/96

No match
in R1

40

Left Outer Join



**Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause
(use nulls in fields of non-matching tuples)**

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

Returns all sailors & information on whether they have reserved boats

41

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

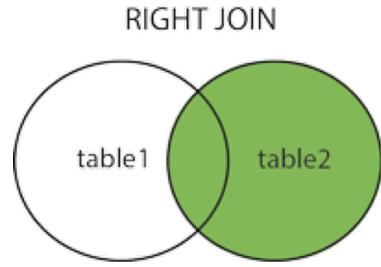
sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	null

42

Right Outer Join



Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

SELECT r.sid, b.bid, b.name

**FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid**

Returns all boats & information on which ones are reserved.

43

**SELECT r.sid, b.bid, b.name
FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid**

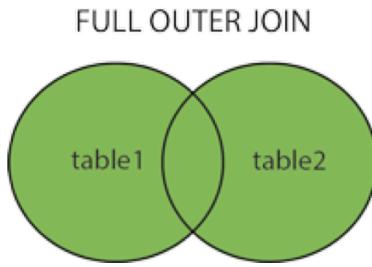
<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

bid	bname	color	
101	Interlake	blue	
102	Interlake	red	
103	Clipper	green	
104	Marine	red	

r.sid	b.bid	b.name
22	101	Interlake
null	102	Interlake
95	103	Clipper
null	104	Marine

44

Full Outer Join



Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Sailors s FULL OUTER JOIN Boats b  
ON s.sname = b.bname
```

45

```
SELECT s.sid, s.sname, b.bid, b.name  
FROM Sailors s FULL OUTER JOIN Boats b  
ON s.sname = b.bname
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

bid	bname	color
101	Interlake	blue
105	Lubber	purple

sid	sname	bid	bname
22	Dustin	<i>null</i>	<i>null</i>
31	Lubber	105	Lubber
95	Bob	<i>null</i>	<i>null</i>
<i>null</i>	<i>null</i>	101	Interlake

46

Aggregate Functions

Significant extension of relational algebra

COUNT (*)	The number of rows in the relation
COUNT ([DISTINCT] A)	The number of (unique) values in the A column
SUM ([DISTINCT] A)	The sum of all (unique) values in the A column
AVG ([DISTINCT] A)	The average of all (unique) values in the A column
MAX (A)	The maximum value in the A column
MIN (A)	The minimum value in the A column

Aggregate Operators

```
SELECT COUNT (*)
FROM Sailors S
```

Count the
number of
sailors

Find the names of
sailors with the
highest rating

```
SELECT AVG (S.age)
FROM Sailors S
WHERE S.rating=10
```

Find the average age

Count the number of
distinct ratings of
sailors called "Bob"

ing)

```
SELECT COUNT (DISTINCT S.rating)
FROM Sailors S
WHERE S.sname='Bob'
```

```
SELECT AVG (DISTINCT S.age)
FROM Sailors S
WHERE S.rating=10
```

Aggregate Operators

```
SELECT COUNT (*)
FROM Sailors S
```

```
SELECT AVG (S.age)
FROM Sailors S
WHERE S.rating=10
```

```
SELECT S.sname
FROM Sailors S
WHERE S.rating= (SELECT MAX(S2.rating)
                  FROM Sailors S2)
```

```
SELECT COUNT (DISTINCT S.rating)
FROM Sailors S
WHERE S.sname='Bob'
```

```
SELECT AVG (DISTINCT S.age)
FROM Sailors S
WHERE S.rating=10
```

Summary Aggregate Functions...

Avg ()
Count ()
First ()
Last ()
Max ()
Min ()
SQL
Sum ()
Group By
Having
Ucase ()

Lcase ()
Mid ()
Len ()
Round ()
Now ()
Format ()

```
SELECT column_name,
       aggregate_function(column_name)
  FROM table_name
 WHERE column_name operator value
 GROUP BY column_name
 HAVING
       aggregate_function(column_name)
          operator value
```

Restriction on SELECT Lists With Aggregation

- **If any aggregation is used, then each element of the SELECT list must be either:**
 1. Aggregated, or
 2. An attribute on the GROUP BY list.

51

Illegal Query Example

- **You might think you could find the bar that sells Bud the cheapest by:**

```
SELECT bar, MIN(price)
FROM Sells
WHERE beer = 'Bud' ;
```

- **But this query is illegal in SQL.**
- **Why?**

52

Find name and age of the oldest sailor(s)



Only aggregate operations allowed

```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```

Find name and age of the oldest sailor(s)

- **The first query is incorrect!**

- **Third query equivalent to second query**

– allowed in SQL/92 standard, but not supported in some systems.

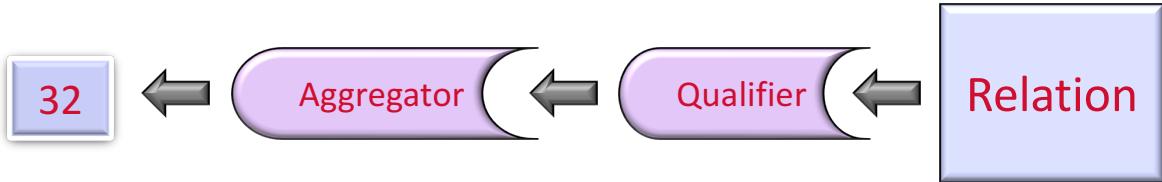
~~SELECT S.sname, MAX (S.age)
FROM Sailors S~~

SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
(SELECT MAX (S2.age)
FROM Sailors S2)

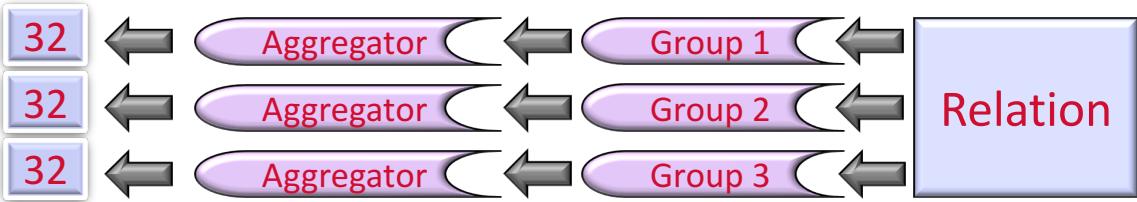
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
FROM Sailors S2)
= S.age

GROUP BY and HAVING (1)

- So far, we've applied aggregate operators to all (qualifying) tuples.



- Sometimes, we want to apply them to each of several *groups* of tuples.



55

GROUP BY and HAVING (2)

Consider: *Find the age of the youngest sailor for each rating level.* /* Min(age) for multiple groups

- If we know that rating values go from 1 to 10, we can write 10 queries that look like this:

For $i = 1, 2, \dots, 10$: $\left\{ \begin{array}{l} \text{SELECT MIN (S.age)} \\ \text{FROM Sailors S} \\ \text{WHERE S.rating} = i \end{array} \right.$

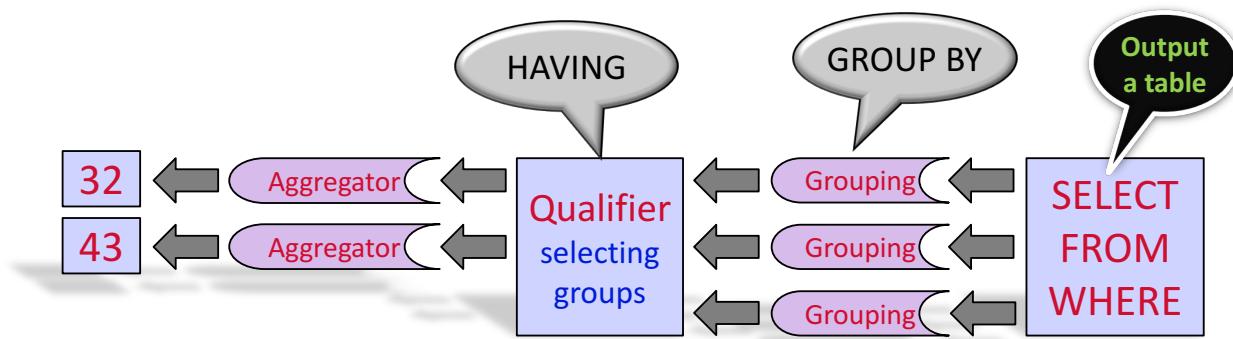
- In general, we don't know how many rating levels exist, and what the rating values for these levels are !

56

Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list  
FROM        relation-list  
WHERE       qualification  
GROUP BY   grouping-list  
HAVING     group-qualification
```

MIN(Attribute)



57

Queries With GROUP BY

- To generate values for a column based on groups of rows, use **aggregate** functions in SELECT statements with the GROUP BY clause

```
SELECT      [DISTINCT] target-list  
FROM        relation-list  
[WHERE      qualification]  
GROUP BY   grouping-list
```

The **target-list** contains

- (i) list of column names &
- (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - column name list (i) can contain only attributes from the **grouping-list**.

58

Group By Examples

For each rating, find the average age of the sailors

```
SELECT S.rating, AVG (S.age)
FROM Sailors S
GROUP BY S.rating
```

For each rating find the age of the youngest sailor with age ≥ 18

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```

59

Conceptual Evaluation

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

1. The cross-product of *relation-list* is computed
2. Tuples that fail *qualification* are discarded
3. `Unnecessary' fields are deleted
4. The remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
5. The *group-qualification* is then applied to eliminate some groups
6. One answer tuple is generated per qualifying group

60

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) > 1

```

rating	age
7	35.0
7	45.0
7	35.0
8	55.5
10	35.0

Answer

Only one group satisfies HAVING

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Input relation

Only S.rating and S.age are mentioned in SELECT

"GROUP BY and HAVING" Examples

Find the age of the youngest sailor with age ≥ 18

```

SELECT MIN(S.age)
FROM Sailors S
WHERE S.age >= 18

```

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating

```

Find the age of the youngest sailor with age ≥ 18 , for each rating

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) > 1

```

For each red boat, find the number of reservations for this boat

```
SELECT      B.bid, COUNT (*) AS scount  
FROM        Boats B, Reserves R  
WHERE       R.bid=B.bid AND B.color='red'  
GROUP BY    B.bid
```

3) Count the number of reservations for each red-boat group

1) Find all reservations for red boats

2) Group the reservations for red boats according to bid

63

For each red boat, find the number of reservations for this boat

Grouping over a join of two relations

B.color is not in the grouping-list

```
SELECT      B.bid, COUNT (*) AS scount  
FROM        Boats B, Reserves R  
WHERE       R.bid=B.bid  
GROUP BY    B.bid  
HAVING     B.color='red'
```

Illegal!

Note: HAVING clause is to select groups !

64

Find the age of the youngest sailor older than 18, for each rating level that has at least 2 sailors

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT(*)
               FROM Sailors S2
               WHERE S.rating = S2.rating)
    
```

Replacing this by
“HAVING COUNT(*) > 1”

- Shows HAVING clause can also contain a subquery.
- We can use S.rating inside the nested subquery because it has a single value for the current group of sailors.
- What if HAVING clause is replaced by “HAVING COUNT(*) > 1”
 - Find the age of the youngest sailor older than 18, for each rating level that has at least two such sailors. /* see next page

65

Find the age of the youngest sailor older than 18, for each rating level that has at least 2 sailors

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT(*)
               FROM Sailors S2
               WHERE S.rating = S2.rating)
    
```

At least 2 sailors

Counting including sailors younger than 18

“age” is not mentioned in this subquery

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING COUNT(*) > 1
    
```

Counting include only adult sailors

At least 2 such sailors, i.e., older than 18

66

Find those ratings for which the average age is the minimum over all ratings

```
SELECT S.rating  
FROM Sailors S  
WHERE S.age = (SELECT MIN (AVG (S2.age))  
                FROM Sailors S2)
```

Aggregate operations cannot be nested

67

Find those ratings for which the average age is the minimum over all ratings

Correct solution (in SQL/92):

Output this rating group and its average age

Find average age for each rating group.
This table has two columns

```
SELECT Temp.rating, Temp.avgage  
FROM (SELECT S.rating, AVG (S.age) AS avgage  
      FROM Sailors S  
      GROUP BY S.rating) AS Temp  
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)  
                      FROM Temp)
```

Average age for some rating group

Minimum over all ratings

68

Null Values

- Field values in a tuple are sometimes
 - *unknown* (e.g., a rating has not been assigned), or
 - *inapplicable* (e.g., no spouse's name).
- SQL provides a special value *null* for such situations.

69

Null Values

The presence of *null* complicates many issues:

- Special operators needed, e.g., **IS NULL** to test if a value is *null*.
- Is *rating > 8* true or false when *rating* is equal to *null*? *null*
- What about **AND**, **OR** and **NOT** ? Need a 3-valued logic (*true*, *false*, and *unknown*), e.g., *(unknown OR false) = unknown*.
- Meaning of constructs must be defined carefully, e.g., **WHERE clause eliminates rows that don't evaluate to true**.
 - Null + 5 = null; but SUM (null, 5) = 5. (nulls can cause some unexpected behavior)
- New operators (in particular, *outer joins*) possible/needed.

70

Views

```
CREATE VIEW view_name
AS select_statement
```

Makes development simpler
Often used for security
Not instantiated - makes updates tricky

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scount
      FROM Boats B, Reserves R
     WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid
```

71

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scount
      FROM Boats B, Reserves R
     WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid
```

b.bid	scount
102	1

Reds

72

Sailors who have reserved all boats

```

SELECT S.name
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.name, S.sid
HAVING COUNT(DISTINCT R.bid) =
    (Select COUNT (*) FROM Boats)
  
```

sname	sid	bid
Frodo	1	102
Bilbo	2	101
Bilbo	2	102
Frodo	1	102
Bilbo	2	103

sname	sid	bid
Frodo	1	102,102
Bilbo	2	101, 102, 103

Sailors

sid	sname	rating	age
1	Frodo	7	22
2	Bilbo	2	39
3	Sam	8	27

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

sid	bid	day
1	102	9/12
2	102	9/12
2	101	9/14
1	102	9/10
2	103	9/13

Two more important topics

- **Constraints**
- **SQL embedded in other languages**

Integrity Constraints

- **IC conditions that every legal instance of a relation must satisfy.**
 - Inserts/deletes/updates that violate ICs are disallowed.
 - Can ensure application semantics (e.g., sid is a key),
 - ...or prevent inconsistencies (e.g., sname has to be a string, age must be < 200)
- **Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.**
 - Domain constraints: Field values must be of right type.
Always enforced.
 - Primary key and foreign key constraints: coming right up.

75

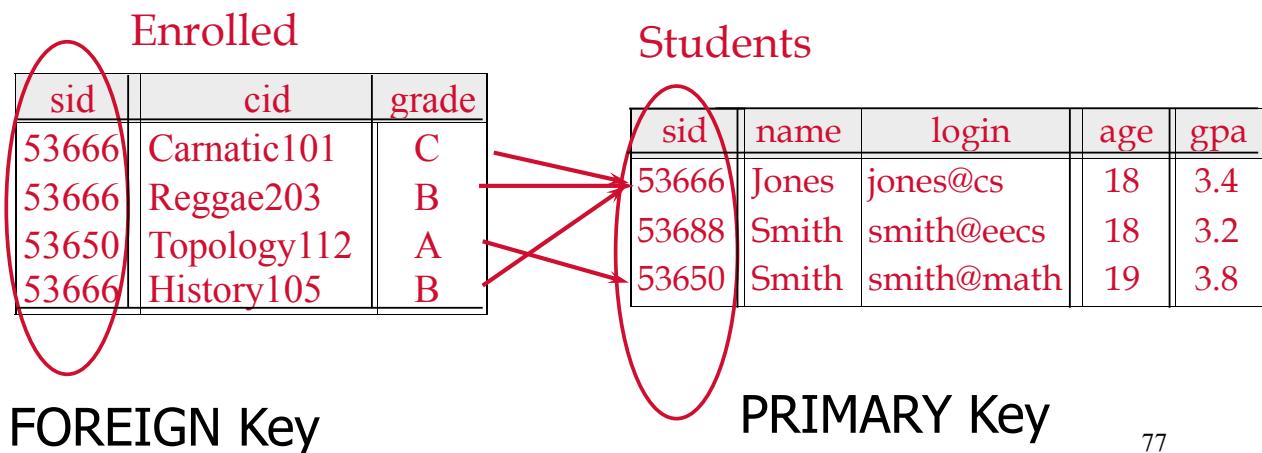
Where do ICs Come From?

- **Semantics of the real world!**
- **Note:**
 - We can check IC violation in a DB instance
 - We can NEVER infer that an IC is true by looking at an instance.
 - An IC is a statement about all possible instances!
 - From example, we know name is not a key, but the assertion that sid is a key is given to us.
- **Key and foreign key ICs are the most common**
- **More general ICs supported too.**

76

Keys

- Keys are a way to associate tuples in different relations
- Keys are one form of IC



77

Primary Keys

- A set of fields is a **superkey** if:
 - No two distinct tuples can have same values in all key fields
- A set of fields is a **key for a relation** if :
 - It is a superkey
 - No subset of the fields is a superkey
- **what if >1 key for a relation?**
 - One of the keys is chosen (by DBA) to be the **primary key**. Other keys are called **candidate keys**.
- **E.g.**
 - sid is a key for Students.
 - What about name?
 - The set {sid, gpa} is a superkey.

78

Primary and Candidate Keys

- Possibly many ***candidate keys*** (specified using **UNIQUE**), one of which is chosen as the ***primary key***.
- Keys must be used carefully!

```
CREATE TABLE Enrolled1    CREATE TABLE Enrolled2  
  (sid CHAR(20),          (sid CHAR(20),  
   cid CHAR(20),          cid CHAR(20),  
   grade CHAR(2),         grade CHAR(2),  
   PRIMARY KEY (sid,cid)) PRIMARY KEY (sid),  
                                UNIQUE (cid, grade))
```

VS.

"For a given student and course, there is a single grade."

79

Primary and Candidate Keys

```
CREATE TABLE Enrolled1  
  (sid CHAR(20),  
   cid CHAR(20),  
   grade CHAR(2),  
   PRIMARY KEY (sid,cid))
```

```
CREATE TABLE Enrolled2  
  (sid CHAR(20),  
   cid CHAR(20),  
   grade CHAR(2),  
   PRIMARY KEY (sid),  
   UNIQUE (cid, grade))
```

VS.

```
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');  
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');  
INSERT INTO enrolled1 VALUES ('1234', 'cs61c', 'A+');
```

```
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');  
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');  
INSERT INTO enrolled2 VALUES ('1234', 'cs61c', 'A+');  
INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');
```

"For a given student and course, there is a single grade."

80

Primary and Candidate Keys

```
CREATE TABLE Enrolled1  
(sid CHAR(20),  
 cid CHAR(20),  
 grade CHAR(2),  
 PRIMARY KEY (sid,cid));
```

VS.

```
CREATE TABLE Enrolled2  
(sid CHAR(20),  
 cid CHAR(20),  
 grade CHAR(2),  
 PRIMARY KEY (sid),  
 UNIQUE (cid, grade));
```

```
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');  
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');  
INSERT INTO enrolled1 VALUES ('1234', 'cs61C', 'A+');
```

```
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');  
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');  
INSERT INTO enrolled2 VALUES ('1234', 'cs61C', 'A+');  
INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');
```

"Students can take only one course, and no two students in a course receive the same grade." 81

Foreign Keys, Referential Integrity

- **Foreign key:** a “logical pointer”
 - Set of fields in a tuple in one relation that ‘refer’ to a tuple in another relation.
 - Reference to *primary key* of the other relation.
- **All foreign key constraints enforced?**
 - referential integrity!
 - i.e., no dangling references.

Foreign Keys in SQL

- **E.g. Only students listed in the Students relation should be allowed to enroll for courses.**

- *sid* is a foreign key referring to *Students*:

```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
 PRIMARY KEY (sid,cid),  
 FOREIGN KEY (sid) REFERENCES Students);
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B
11111	English102	A

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

83

Enforcing Referential Integrity

- ***sid* in Enrolled: foreign key referencing Students.**
- **Scenarios:**
 - Insert Enrolled tuple with non-existent student id?
 - Delete a Students tuple?
 - Also delete Enrolled tuples that refer to it? (CASCADE)
 - Disallow if referred to? (NO ACTION)
 - Set sid in referring Enrolled tuples to a *default* value? (SET DEFAULT)
 - Set sid in referring Enrolled tuples to *null*, denoting ‘*unknown*’ or ‘*inapplicable*’. (SET NULL)
- **Similar issues arise if primary key of Students tuple is updated.**

84

General Constraints

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
CREATE TABLE Sailors  
( sid INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY (sid),  
  CHECK ( rating >= 1  
        AND rating <= 10 ))
```

```
CREATE TABLE Reserves  
( sname CHAR(10),  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (bid,day),  
  CONSTRAINT noInterlakeRes  
  CHECK ('Interlake' <>  
         ( SELECT B.bname  
           FROM Boats B  
          WHERE B.bid=bid)))
```

Constraints Over Multiple Relations

```
CREATE TABLE Sailors  
( sid INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY (sid),  
  CHECK  
  ( (SELECT COUNT (S.sid) FROM Sailors S)  
    + (SELECT COUNT (B.bid) FROM  
      Boats B) < 100 )
```

Number of boats plus number of sailors is < 100

Constraints Over Multiple Relations

```
CREATE TABLE Sailors
```

- Awkward and wrong!
 - Only checks sailors!
- ASSERTION is the right solution; not associated with either table.
 - Unfortunately, not supported in many DBMS.
 - *Triggers* are another solution.

```
( sid INTEGER,  
sname CHAR(10),  
rating INTEGER,  
age REAL,  
PRIMARY KEY (sid),  
CHECK  
( (SELECT COUNT (S.sid) FROM Sailors S)  
+ (SELECT COUNT (B.bid) FROM  
Boats B) < 100 )
```

Number of boats plus number of sailors is < 100

```
CREATE ASSERTION smallClub  
CHECK  
( (SELECT COUNT (S.sid) FROM Sailors S)  
+ (SELECT COUNT (B.bid)  
FROM Boats B) < 100 )
```

Two more important topics

- **Constraints**
- **SQL embedded in other languages**

Writing Applications with SQL

- **SQL is not a general purpose programming language.**
 - + Tailored for data retrieval and manipulation
 - + Relatively easy to optimize and parallelize
 - Can't write entire apps in SQL alone
- **Options:**
 - Make the query language “Turing complete”
 - Avoids the “impedance mismatch”
 - makes “simple” relational language complex
 - Allow SQL to be embedded in regular programming languages.
 - Q: What needs to be solved to make the latter approach work?

89

Cursors

- **Can declare a cursor on a relation or query**
- **Can *open* a cursor**
- **Can repeatedly *fetch* a tuple (moving the cursor)**
- **Special return value when all tuples have been retrieved.**
- ***ORDER BY* allows control over the order tuples are returned.**
 - Fields in ORDER BY clause must also appear in SELECT clause.
- ***LIMIT* controls the number of rows returned (good fit w/*ORDER BY*)**
- **Can also modify/delete tuple pointed to by a cursor**
 - A “non-relational” way to get a handle to a particular tuple

90

Database APIs

- **A library with database calls (API)**
 - special objects/methods
 - passes SQL strings from language, presents **result sets** in a language-friendly way
 - *ODBC* a C/C++ standard started on Windows
 - *JDBC* a Java equivalent
 - Most scripting languages have similar things
 - E.g. in Ruby there's the "pg" gem for Postgres
- **ODBC/JCDB try to be DBMS-neutral**
 - at least try to hide distinctions across different DBMSs

91

Summary

- Relational model has **well-defined query semantics**
- SQL provides functionality close to basic relational model

(some differences in duplicate handling, null values, set operators, ...)
- Typically, many ways to write a query
 - DBMS figures out a fast way to execute a query, regardless of how it is written.

92

ADVANCED EXAMPLES

93

Getting Serious

- **Two “fancy” queries for different applications**
 - Clustering Coefficient for Social Network graphs
 - Medians for “robust” estimates of the central value

94

Serious SQL: Social Nets Example

```
-- An undirected friend graph. Store each link once
CREATE TABLE Friends(
    fromID integer,
    toID integer,
    since date,
    PRIMARY KEY (fromID, toID),
    FOREIGN KEY (fromID) REFERENCES Users,
    FOREIGN KEY (toID) REFERENCES Users,
    CHECK (fromID < toID));

-- Return both directions
CREATE VIEW BothFriends AS
    SELECT * FROM Friends
    UNION ALL
    SELECT F.toID AS fromID, F.fromID AS toID, F.since
    FROM Friends F;
```

95

6 degrees of friends

```
SELECT F1.fromID, F5.toID
    FROM BothFriends F1, BothFriends F2, BothFriends F3,
        BothFriends F4, BothFriends F5
   WHERE F1.toID = F2.fromID
     AND F2.toID = F3.fromID
     AND F3.toID = F4.fromID
     AND F4.toID = F5.fromID;
```

Clustering Coefficient of a Node

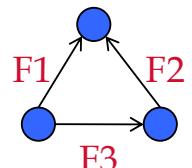
$$C_i = \frac{2|\{e_{jk}\}|}{k_i(k_i-1)}$$

- where:
 - k_i is the number of neighbors of node I
 - e_{jk} is an edge between nodes j and k neighbors of i , ($j < k$). (A triangle!)
- I.e. Cliquishness: the fraction of your friends that are friends with each other!
- Clustering Coefficient of a graph is the average CC of all nodes.

97

In SQL

$$C_i = \frac{2|\{e_{jk}\}|}{k_i(k_i-1)}$$

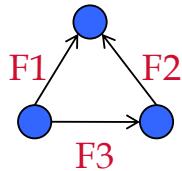


```
CREATE VIEW NEIGHBOR_CNT AS
SELECT fromID AS nodeID, count(*) AS friend_cnt
FROM BothFriends
GROUP BY nodeID;

CREATE VIEW TRIANGLES AS
SELECT F1.toID as root, F1.fromID AS friend1,
       F2.fromID AS friend2
  FROM BothFriends F1, BothFriends F2, Friends F3
 WHERE F1.toID = F2.toID      /* Both point to root */
   AND F1.fromID = F3.fromID /* Same origin as F1 */
   AND F3.toID = F2.fromID  /* points to origin of F2 */
;
```

98

In SQL



$$C_i = 2|\{e_{jk}\}| / k_i(k_i-1)$$

```
CREATE VIEW NEIGHBOR_EDGE_CNT AS
SELECT root, COUNT(*) as cnt FROM TRIANGLES
GROUP BY root;

CREATE VIEW CC_PER_NODE AS
SELECT NE.root, 2.0*NE.cnt /
       (N.friend_cnt*(N.friend_cnt-1)) AS CC
  FROM NEIGHBOR_EDGE_CNT NE, NEIGHBOR_CNT N
 WHERE NE.root = N.nodeID;

SELECT AVG(cc) FROM CC_PER_NODE;
```

99

Median

- **Given n values in sorted order, the one at position n/2**
 - Assumes an odd # of items
 - For an even #, can take the lower of the middle 2
- **A much more “robust” statistic than average**
 - Q: Suppose you want the mean to be 1,000,000. What fraction of values do you have to corrupt?
 - Q2: Suppose you want the median to be 1,000,000. Same question.
 - This is called the *breakdown point* of a statistic.
 - Important for dealing with data *outliers*
 - E.g. dirty data
 - Even with real data: “overfitting”

100

Median in SQL

```
SELECT c AS median FROM T
WHERE
(SELECT COUNT(*) from T AS T1
 WHERE T1.c < T.c)
=
(SELECT COUNT(*) from T AS T2
 WHERE T2.c > T.c);
```

101

Median in SQL

```
SELECT c AS median FROM T
WHERE
(SELECT COUNT(*) from T AS T1
 WHERE T1.c < T.c)
=
(SELECT COUNT(*) from T AS T2
 WHERE T2.c > T.c);
```

102

Faster Median in SQL

```
SELECT x.c as median
  FROM T x, T y
 GROUP BY x.c
HAVING
  SUM(CASE WHEN y.c <= x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)+1)/2
AND
  SUM(CASE WHEN y.c >= x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)/2)+1
```

Why faster?

Note: handles even # of items!

103

EXAMPLES

Essential SQL Statements

```
CREATE DATABASE database_name
```

```
DROP DATABASE database_name
```

```
CREATE TABLE table_name  
(  
column_name1 data_type,  
column_name2 data_type,  
column_name3 data_type,  
...  
)
```

```
DROP TABLE table_name  
DELETE FROM table_name  
DELETE * FROM table_name
```

```
CREATE INDEX index_name  
ON table_name (column_name)
```

```
DROP INDEX index_name
```

```
CREATE UNIQUE INDEX index_name  
ON table_name (column_name)
```

```
INSERT INTO table_name  
VALUES (value1, value2, ....)  
INSERT INTO table_name  
(column1, column2,...)  
VALUES (value1, value2, ....)
```

```
DELETE FROM table_name  
WHERE some_column=some_value
```

More SQL Statements

```
ALTER TABLE table_name  
ADD column_name datatype
```

```
ALTER TABLE table_name  
DROP COLUMN column_name
```

```
SELECT column_name AS column_alias  
FROM table_name
```

```
SELECT column_name  
FROM table_name AS table_alias
```

```
SELECT column_name(s)  
FROM table_name  
WHERE column_name  
BETWEEN value1 AND value2  
SELECT column_name(s)  
FROM table_name  
WHERE condition  
AND|OR condition
```

Other SQL Statements

- **AUTO INCREMENT Field**
- **SELECT INTO**
 - Selects data from one table and inserts it into a new table
- **LIMIT**
 - Specify the number of records to return
- **CREATE VIEW**
 - Create a virtual table based on the result-set of an SQL statement
- **TRUNCATE TABLE**
 - Delete all table contents

```
CREATE VIEW view_name AS      SELECT column_name(s)      SELECT column_name(s)
SELECT column_name(s)        INTO newtable [IN          FROM table_name
FROM table_name           externaldb]          LIMIT number;
WHERE condition            FROM table1;
```

107

Example: Relation (Table)

Row/Tuple/Record

Column/Attribute/Field

name	birth	gpa	grad
Anderson	1987-10-22	3.9	2009
Jones	1990-4-16	2.4	2012
Hernandez	1989-8-12	3.1	2011
Chen	1990-2-4	3.2	2011

Column Types → VARCHAR(30) DATE FLOAT INT

108

Example: Primary Key

Unique For Each Row

The diagram illustrates a primary key constraint. A vertical arrow points from the text "Unique For Each Row" down to the "id" column header of the table below. The table has five columns: id, name, birth, gpa, and grad. The "id" column is highlighted with a light blue background, indicating it is the primary key. Below the table, the data types are listed: INT for id, VARCHAR(30) for name, DATE for birth, FLOAT for gpa, and INT for grad.

id	name	birth	gpa	grad
14	Anderson	1987-10-22	3.9	2009
38	Jones	1990-4-16	2.4	2012
77	Hernandez	1989-8-12	3.1	2011
104	Chen	1990-2-4	3.2	2011

INT VARCHAR(30) DATE FLOAT INT

109

Basic Table Operations

```
CREATE TABLE students (
    id INT AUTO_INCREMENT,
    name VARCHAR(30),
    birth DATE,
    gpa FLOAT,
    grad INT,
    PRIMARY KEY(id));

INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Anderson', '1987-10-22', 3.9, 2009);
INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Jones', '1990-4-16', 2.4, 2012);

DELETE FROM students WHERE name='Anderson';

DROP TABLE students;
```

110

Query: Display Entire Table

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.4	2012
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

```
SELECT * FROM students;
```

```
+-----+-----+-----+-----+
| id | name      | birth       | gpa     | grad   |
+-----+-----+-----+-----+
| 1  | Anderson  | 1987-10-22 | 3.9    | 2009  |
| 2  | Jones     | 1990-04-16 | 2.4    | 2012  |
| 3  | Hernandez | 1989-08-12 | 3.1    | 2011  |
| 4  | Chen      | 1990-02-04 | 3.2    | 2011  |
+-----+-----+-----+-----+
```

111

Query: Select Columns name and gpa

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.4	2012
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

```
SELECT name, gpa  
FROM students;
```

```
+-----+-----+
| name      | gpa   |
+-----+-----+
| Anderson  | 3.9  |
| Jones     | 2.4  |
| Hernandez | 3.1  |
| Chen      | 3.2  |
+-----+-----+
```

112

Query: Filter Rows

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.4	2012
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

```
SELECT name, gpa FROM students  
WHERE gpa > 3.0;
```

```
+-----+-----+  
| name | gpa |  
+-----+-----+  
| Anderson | 3.9 |  
| Hernandez | 3.1 |  
| Chen | 3.2 |  
+-----+-----+
```

113

Query: Sort Output

- **The ORDER BY keyword is used to sort the result-set by one or more columns**

```
SELECT column_name, column_name  
FROM table_name  
ORDER BY column_name ASC | DESC, column_name ASC | DESC;
```

114

Query: Sort Output by gpa

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.4	2012
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

```
SELECT gpa, name, grad FROM students
WHERE gpa > 3.0
ORDER BY gpa DESC;
```

```
+-----+-----+-----+
| gpa | name      | grad |
+-----+-----+-----+
| 3.9 | Anderson  | 2009 |
| 3.2 | Chen      | 2011 |
| 3.1 | Hernandez | 2011 |
+-----+-----+-----+
```

115

Update Value(s)

- **The UPDATE statement is used to update existing records in a table.**

```
UPDATE table_name
SET column1=value1, column2=value2, ...
WHERE some_column=some_value;
```

116

Update Value(s)

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.4	2012
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

```
UPDATE students
SET gpa = 2.6, grad = 2013
WHERE id = 2
```

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-4-16	2.6	2013
3	Hernandez	1989-8-12	3.1	2011
4	Chen	1990-2-4	3.2	2011

117

Foreign Key

students	id	name	birth	gpa	grad	advisor_id
1	Anderson	1987-10-22	3.9	2009	2	
2	Jones	1990-4-16	2.4	2012	1	
3	Hernandez	1989-8-12	3.1	2011	1	
4	Chen	1990-2-4	3.2	2011	1	

advisors	id	name	title
1	Fujimura	assocprof	
2	Bolosky	prof	

```
SELECT s.name, s.gpa
FROM students s, advisors p
WHERE s.advisor_id = p.id AND p.name = 'Fujimura';
```

s.id	s.name	s.birth	s.gpa	s.grad	s.advisor_id	p.id	p.name	p.title
1	Anderson	1987-10-22	3.9	2009	2	1	Fujimura	assocprof
1	Anderson	1987-10-22	3.9	2009	2	2	Bolosky	prof
2	Jones	1990-4-16	2.4	2012	1	1	Fujimura	assocprof
2	Jones	1990-4-16	2.4	2012	1	2	Bolosky	prof
3	Hernandez	1989-8-12	3.1	2011	1	1	Fujimura	assocprof
3	Hernandez	1989-8-12	3.1	2011	1	2	Bolosky	prof
4	Chen	1990-2-4	3.2	2011	1	1	Fujimura	assocprof
4	Chen	1990-2-4	3.2	2011	1	2	Bolosky	prof

Slide 118

students	id	name	birth	gpa	grad	advisor_id	advisors	id	name	title
	1	Anderson	1987-10-22	3.9	2009	2		1	Fujimura	assocprof
	2	Jones	1990-4-16	2.4	2012	1		2	Bolosky	prof
	3	Hernandez	1989-8-12	3.1	2011	1				
	4	Chen	1990-2-4	3.2	2011	1				

```
SELECT s.name, s.gpa
FROM students s, advisors p
WHERE s.advisor_id = p.id AND p.name = 'Fujimura';
```

	s.id	s.name	s.birth	s.gpa	s.grad	s.advisor_id		p.id	p.name	p.title
	1	Anderson	1987-10-22	3.9	2009	2		1	Fujimura	assocprof
	1	Anderson	1987-10-22	3.9	2009	2		2	Bolosky	prof
	2	Jones	1990-4-16	2.4	2012	1		1	Fujimura	assocprof
	2	Jones	1990-4-16	2.4	2012	1		2	Bolosky	prof
	3	Hernandez	1989-8-12	3.1	2011	1		1	Fujimura	assocprof
	3	Hernandez	1989-8-12	3.1	2011	1		2	Bolosky	prof
	4	Chen	1990-2-4	3.2	2011	1		1	Fujimura	assocprof
	4	Chen	1990-2-4	3.2	2011	1		2	Bolosky	prof

students	id	name	birth	gpa	grad	advisor_id	advisors	id	name	title
	1	Anderson	1987-10-22	3.9	2009	2		1	Fujimura	assocprof
	2	Jones	1990-4-16	2.4	2012	1		2	Bolosky	prof
	3	Hernandez	1989-8-12	3.1	2011	1				
	4	Chen	1990-2-4	3.2	2011	1				

```
SELECT s.name, s.gpa
FROM students s, advisors p
WHERE s.advisor_id = p.id AND p.name = 'Fujimura';
```

```
+-----+-----+
| name      | gpa   |
+-----+-----+
| Jones     | 2.4   |
| Hernandez | 3.1   |
| Chen      | 3.2   |
+-----+-----+
```

	id	name	birth	gpa	grad
students	1	Anderson	1987-10-22	3.9	2009
Courses	1	CS142	Web stuff	Winter 2009	
2	ART101	Finger painting	Fall 2008		
3	ART101	Finger painting	Winter 2009		
4	PE204	Mud wrestling	Winter 2009		

	course_id	student_id
	1	1
	3	1
	4	1
	1	2
	2	2
	1	3
	2	4
	4	4

```
SELECT s.name, c.quarter
FROM students s, courses c, courses_students cs
WHERE c.id = cs.course_id AND s.id = cs.student_id
AND c.number = 'ART101';
```

```
+-----+-----+
| name      | quarter    |
+-----+-----+
| Jones     | Fall 2008  |
| Chen      | Fall 2008  |
| Anderson   | Winter 2009 |
+-----+-----+
```

Slide 121

Back to Our Running Example ...

```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age REAL,
    PRIMARY KEY sid);
```

```
CREATE TABLE Boats (
    bid INTEGER,
    bname CHAR (20),
    color CHAR(10)
    PRIMARY KEY bid);
```

```
CREATE TABLE Reserves (
    sid INTEGER,
    bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY sid REFERENCES Sailors,
    FOREIGN KEY bid REFERENCES Boats);
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

sid	bid	day
1	102	9/12
2	102	9/13

Back to Our Running Example ...

<i>Reserves</i>			<i>Sailors</i>
<u>sid</u>	<u>bid</u>	<u>day</u>	<u>sid</u>
22	101	10/10/96	22 Dustin 7 45.0
95	103	11/12/96	31 Lubber 8 55.5
			95 Bob 3 63.5

<i>Boats</i>		
<u>bid</u>	<u>bname</u>	<u>color</u>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

APPENDIX

DDL – Create Table

- **CREATE TABLE *table_name***
({ *column_name data_type*
[DEFAULT *default_expr*] [
***column_constraint* [, ...]] |**
***table_constraint* } [, ...])**
- **Data Types (PostgreSQL) include:**
character(n) – fixed-length character string
character varying(n) – variable-length character string
smallint, integer, bigint, numeric, real, double precision
date, time, timestamp, ...
serial - unique ID for indexing and cross reference
...
- **PostgreSQL also allows OIDs and other “system types”, arrays, inheritance, rules...**
conformance to the SQL-1999 standard is variable.

125

Constraints

- **Recall that the schema defines the legal instances of the relations.**
- **Data types are a way to limit the kind of data that can be stored in a table, but they are often insufficient.**
 - e.g., prices must be positive values
 - uniqueness, referential integrity, etc.
- **Can specify constraints on individual columns or on tables.**

126

Column constraints

```
[ CONSTRAINT constraint_name ]
  { NOT NULL | NULL | UNIQUE | PRIMARY KEY |
    CHECK (expression) |
    REFERENCES reftable [ ( refcolumn ) ] [ ON
      DELETE action ] [ ON UPDATE action ] }
```

primary key = unique + not null; also used as default target for references. (can have at most 1)
expression must produce a boolean result and reference that column's value only.
references is for foreign keys; action is one of:
NO ACTION, CASCADE, SET NULL, SET DEFAULT

127

Table constraints

- **CREATE TABLE *table_name* ({ *column_name* *data_type* [DEFAULT *default_expr*] [*column_constraint* [, ...]] | *table_constraint* } [, ...])**

Table Constraints:

- [CONSTRAINT *constraint_name*]
 { UNIQUE (*column_name* [, ...]) |
 PRIMARY KEY (*column_name* [, ...]) |
 CHECK (*expression*) |
 FOREIGN KEY (*column_name* [, ...]) |
 REFERENCES *reftable* [(*refcolumn* [, ...])] [ON
 DELETE *action*] [ON UPDATE *action*] }

128

Create Table (Examples)

```
CREATE TABLE films (
    code      CHAR(5) PRIMARY KEY,
    title     VARCHAR(40),
    did       DECIMAL(3),
    date_prod DATE,
    kind      VARCHAR(10),
    CONSTRAINT production UNIQUE(date_prod)
    FOREIGN KEY did REFERENCES distributors
        ON DELETE NO ACTION
);
CREATE TABLE distributors (
    did      DECIMAL(3) PRIMARY KEY,
    name    VARCHAR(40)
    CONSTRAINT con1 CHECK (did > 100 AND name <> '129')
);
```

Other DDL Statements

- **Alter Table**
 - use to add/remove columns, constraints, rename things ...
- **Drop Table**
 - Compare to “Delete * From Table”
- **Create/Drop View**
- **Create/Drop Index**
- **Grant/Revoke privileges**
 - SQL has an authorization model for saying who can read/modify/delete etc. data and who can grant and revoke privileges!