

Database Management Systems


Buffer and File Management

Fall 2015

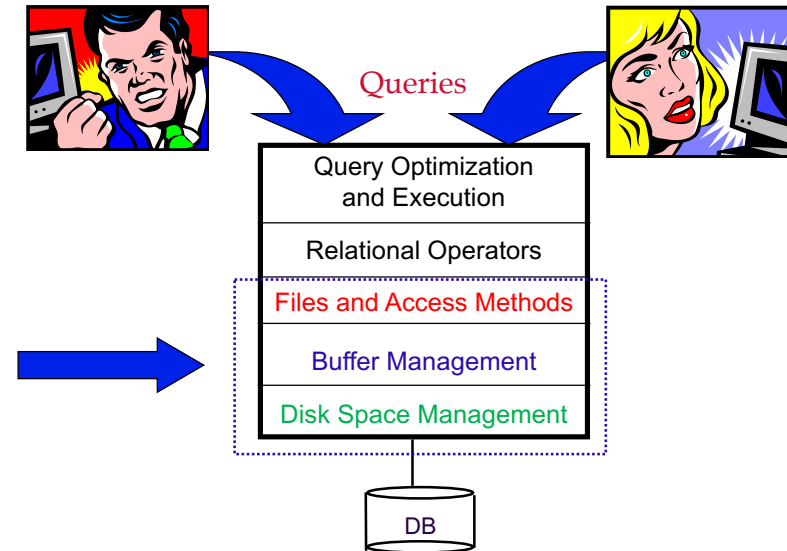
"Yea, from the table of my memory
I'll wipe away all trivial fond records."
-- Shakespeare, *Hamlet*

Disks and Files

- DBMS stores information on disks.
 - In an electronic world, disks are a mechanical anachronism!
- This has major implications for DBMS design!
 - **READ**: transfer data from disk to main memory (RAM).
 - **WRITE**: transfer data from RAM to disk.
 - Both are high-cost operations, relative to in-memory operations, so must be planned carefully!



The BIG Picture



Why Not Store It All in Main Memory?

- *Costs too much.* \$60 will buy you either around 8 GB of RAM or around 1000 GB (1 TB) of disk today.
 - High-end Databases today can be in the Petabyte (1000TB) range.
 - Approx 60% of the cost of a production system is in the disks.
- *Main memory is volatile.* We want data to be saved between runs. (Obviously!)
- Note, some specialized systems do store entire database in main memory.
 - Vendors claim 10x speed up vs. traditional DBMS running in main memory.

Storage Access

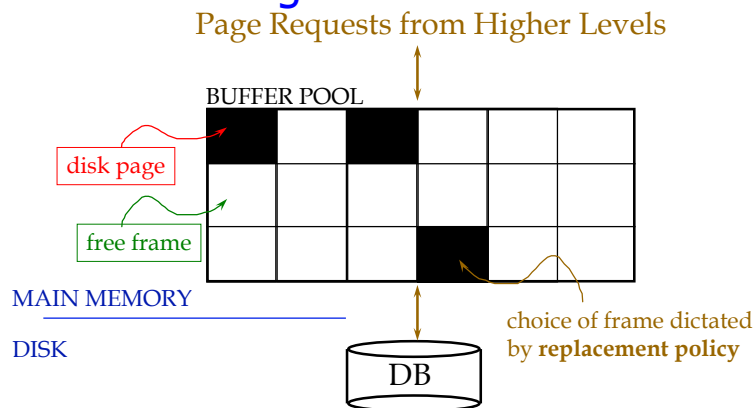
- A database file is partitioned into fixed-length storage units called **blocks**
 - Units of both storage allocation and data transfer.
 - Database system seeks to minimize the number of block transfers between the disk and memory.
 - Can reduce the number of disk accesses by keeping as many blocks as possible in main memory.
- **Buffer** – portion of main memory available to store copies of disk blocks.
- **Buffer manager** – subsystem responsible for allocating buffer space in main memory.

More Terminology...

- **Disk Page** – the unit of transfer between the disk and memory
 - Typically set as a config parameter for the DBMS.
 - Typical value between 4 KBytes to 32 KBytes.
- **Frame** – a unit of memory
 - Typically the same size as the Disk Page Size
- **Buffer Pool** –
 - An area of memory into which database pages are read, modified, and held during processing
 - A collection of frames used by the DBMS to temporarily keep data for use by the query processor.
 - note: We will sometime use the term “buffer” and “frame” synonymously.
- **Pinned block** – memory block that is not allowed to be written back to disk.

Question: When would you use a larger page size rather than a smaller one?

Buffer Management in a DBMS



- *Data must be in RAM for DBMS to operate on it!*
 - The query processor refers to data using virtual memory addresses.
- *Buffer Mgr hides the fact that not all data is in RAM*

When a Page is Requested ...

- If requested page IS in the pool:
 - *Pin* the page and return its address.
- Else, if requested page IS NOT in the pool:
 - If a free frame exists, choose it, Else:
 - Choose a frame for *replacement (only un-pinned pages are candidates)*
 - If chosen frame is “dirty”, write it to disk
 - Read requested page into chosen frame
 - *Pin* the page and return its address.

Buffer Control Blocks (BCBs):

<frame#, pageid, pin_count, dirty>

- A page may be requested many times, so
 - a *pin count* is used.
 - To pin a page, `pin_count++`
 - A page is a candidate for replacement iff `pin_count == 0` ("*unpinned*")
 - Requestor of page must eventually unpin it.
 - `pin_count--`
 - Must also indicate if page has been modified:
 - *dirty* bit is used for this.
- Q: Why is this important?

Buffer Replacement Policy

- Frame is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU), MRU, Clock, etc.
- This policy can have big impact on the number of disk reads and writes.
 - Remember, these are sloooooooooooooow.
- **BIG IDEA** – throw out the page that you are least likely to need in the future.
 - Q: How do you predict the future?
- Efficacy depends on the *access pattern*.

Additional Buffer Manager Notes

- BCB's are hash indexed by pageID
- Concurrency Control & Recovery may entail additional I/O when a frame is chosen for replacement.
(*Write-Ahead Log* protocol; more later.)
- If requests can be predicted (e.g., sequential scans) pages can be pre-fetched several pages at a time.

LRU Replacement Policy

- Least Recently Used (LRU)
 - 1) for each page in buffer pool, keep track of time last *unpinned*
 - *What else might you keep track off?*
 - *How would that impact performance?*
 - 2) Replace the frame that has the oldest (earliest) time
 - Most common policy: intuitive and simple
 - Based on notion of "Temporal Locality"
 - Works well to keep "working set" in buffers.
 - Implemented through doubly linked list of BCBs
 - Requires list manipulation on unpin


Some issues with LRU

- Problem: Sequential flooding
 - LRU + repeated sequential scans.
 - # buffer frames < # pages in file means each page request causes an I/O. MRU much better in this situation (but not in all situations, of course).
- Problem: "cold" pages can hang around a long time before they are replaced
 - Cold pages are pages that have been touched only once recently

"2Q" Replacement Policy

- One Queue (A1) has pages that have been referenced only once.
 - new pages enter here
- A second, LRU Queue (A_m) has pages that have been referenced (pinned) multiple times.
 - pages get promoted from A1 to here
- Replacement victims are usually taken from A1
 - Q: Why????

"Clock" Replacement Policy



- An approximation of LRU
- Arrange frames into a cycle, store one *reference bit* per frame
 - Can think of this as the *2nd chance* bit
- When pin count reduces to 0, turn on ref. bit
- When replacement necessary
 - do for each page in cycle {
 - if (pincount == 0 && ref bit is on)
 - turn off ref bit;
 - else if (pincount == 0 && ref bit is off)
 - choose this page for replacement;
 - } until a page is chosen;

Questions:
How like LRU?
Problems?

DBMS vs. OS File System

OS does disk space & buffer mgmt:
why not let OS manage these tasks?

- Some limitations, e.g., files can't span disks.
 - Note, this is changing --- OS File systems are getting smarter (i.e., more like databases!)
- Buffer management in DBMS requires ability to:
 - *pin a page* in buffer pool, *force a page* to disk & *order writes* (important for implementing CC & recovery)
 - adjust *replacement policy*, and *pre-fetch pages* based on access patterns in typical DB operations.
- Q: Compare DBMS Buffer Mgmt to OS Virtual Memory? to Processor Cache?

Files of Records

- Blocks interface for I/O, but...
- Higher levels of DBMS operate on *records*, and *files of records*.
- **FILE**: A collection of **pages**, each containing a collection of records. Must support:
 - insert/delete/modify** record
 - fetch** a particular record (specified using *record id*)
 - scan** all records (possibly with some conditions on the records to be retrieved)
- Note: typically
 $\text{page size} = \text{block size} = \text{frame size}.$

Catalogs are Stored as Relations!

attr_name	rel_name	type	position
attr_name	Attribute_Cat	string	1
rel_name	Attribute_Cat	string	2
type	Attribute_Cat	string	3
position	Attribute_Cat	integer	4
sid	Students	string	1
name	Students	string	2
login	Students	string	3
age	Students	integer	4
gpa	Students	real	5
fid	Faculty	string	1
fname	Faculty	string	2
sal	Faculty	real	3

Attr_Cat(attr_name, rel_name, type, position)

"MetaData" - System Catalogs


- How to impose structure on all those bytes??
- MetaData: "Data about Data"
- For each relation:
 - name, file location, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
- For each index:
 - structure (e.g., B+ tree) and search key fields
- For each view:
 - view name and definition
- Plus statistics, authorization, buffer pool size, etc.
 - ➡ *Q: But how to store the catalogs????*

It's a bit more complicated...

```

Terminal — psql — 99x28
joeh=# \dt pg_attribute
No matching relations found.
joeh=# \d pg_attribute
Table "pg_catalog.pg_attribute"
  Column      | Type      | Modifiers
-----+-----+-----
attrelid      | oid       | not null
attname       | name      | not null
atttypid      | oid       | not null
attstattarget | integer   | not null
attlen        | smallint  | not null
attnum        | smallint  | not null
attndims      | integer   | not null
attcacheoff   | integer   | not null
atttypmod     | integer   | not null
attbyval      | boolean   | not null
attstorage    | "char"    | not null
attalign      | "char"    | not null
attnotnull    | boolean   | not null
atthasdef     | boolean   | not null
attisdropped  | boolean   | not null
attislocal    | boolean   | not null
attinhcount   | integer   | not null
Indexes:
    "pg_attribute_relid_attname_index" UNIQUE, btree (attrelid, attname)
    "pg_attribute_relid_attnum_index"  UNIQUE, btree (attrelid, attnum)
joeh=#
  
```

Record Formats: Fixed Length





- Information about field types same for all records in a file; stored in *system catalogs*.
- Finding *i*'th field done via arithmetic.

How to Identify a Record?

- The Relational Model doesn't expose "pointers", but that doesn't mean that the DBMS doesn't use them internally.
- Q: Can we use memory addresses to "point" to records?
- Systems use a "Record ID" or "RecID"


Record Formats: Variable Length

- Two alternative formats (# fields is fixed):




- Second offers direct access to *i*'th field, efficient storage of *nulls* (special *don't know* value); small directory overhead.

Page Formats: Fixed Length Records




- *Record id* = *<page id, slot #>*. In first alternative, moving records for free space management changes *rid*; may not be acceptable.

"Slotted Page" for Variable Length Records



- Slot contains: [offset (from start of page), length]
 - in bytes
- Record id = <page id, slot #>
- Page is full when data space and slot array meet.

Slotted page (continued)



- What's the biggest record you can add to the above page?
 - Need 2 bytes for slot: [offset, length] plus record.
- What happens when a record needs to move to a different page?
 - Leave a "tombstone" behind, pointing to new page and slot.
 - Record id remains unchanged – no more than one hop needed.

Slotted Page (continued)

- When need to allocate:
 - If enough room in free space, use it and update free space pointer.
 - Else, try to compact, if successful, use the freed space.
 - Else, tell caller that page is full.
- Advantages:
 - Can move records around in page without changing their record ID
 - Allows lazy space management within the page, with opportunity for clean up later

So far we've organized:

- Fields into Records (fixed and variable length)
- Records into Pages (fixed and variable length)

Now we need to organize Pages into Files

Alternative File Organizations


Many alternatives exist, *each good for some situations, and not so good in others:*

Heap files: Unordered. Suitable when typical access is a file scan retrieving all records. Easy to maintain.

Sorted Files: Best for retrieval in *search key* order, or if only a 'range' of records is needed. Expensive to maintain.

Clustered Files (with Indexes): A compromise between the above two extremes.

Heap File Implemented as a List



- The Heap file name and header page id must be stored persistently.
The catalog is a good place for this.
- Each page contains 2 'pointers' plus data.

Unordered (Heap) Files

- Simplest file structure contains records in *no particular order*.
- As file grows and shrinks, pages are allocated and de-allocated.
- To support record level operations, we must:
 - keep track of the *pages* in a file
 - keep track of *free space* on *pages*
 - keep track of the *records* on a page
- Can organize as a list, as a directory, a tree, ...

Cost Model for Analysis

We ignore CPU costs, for simplicity:

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** (Average) time to read or write disk block



- Measuring number of block I/O's ignores gains of pre-fetching and *sequential access*; thus, even I/O cost is only loosely approximated.
- Average-case analysis; based on several simplistic assumptions.
 - Often called a "*back of the envelope*" calculation.
☛ *Good enough to show the overall trends!*

Some Assumptions in the Analysis

- Single record insert and delete.
- Equality selection - exactly one match (what if more or less???)
- For Heap Files we'll assume:
 - Insert always appends to end of file.
 - Delete just leaves free space in the page.
 - Empty pages are not deallocated.

Sorted Files

- Heap files are **lazy** on **update** - you end up paying on searches.
- Sorted files **eagerly** maintain the file on **update**.
 - The opposite choice in the trade-off
- Let's consider an extreme version
 - No gaps allowed, pages fully packed always
 - Q: How might you relax these assumptions?
- Assumptions for our BotE Analysis:
 - Files compacted after deletions.
 - Searches are on sort key field(s).

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered File
Scan all records	BD		
Equality Search (unique key)	0.5 BD		
Range Search	BD		
Insert	2D		
Delete	$(0.5B+1)D$		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered File
Scan all records	BD	BD	
Equality Search (unique key)	0.5 BD	$(\log_2 B) * D$	
Range Search	BD	$[(\log_2 B) + \text{\#match pg}] * D$	
Insert	2D	$((\log_2 B) + B)D$ <small>(because rd, w0.5 File)</small>	
Delete	$(0.5B+1)D$	Same cost as Insert	