

MongoDB

Modified from slides provided by S. Parikh, A. Im, G. Cai, H. Tunc, J. Stevens, Y. Barve, S. Hei



Motivations

- Problems with SQL
 - Rigid schema
 - Not easily scalable (designed for 90's technology or worse)
 - Requires unintuitive joins
- Perks of mongoDB
 - Easy interface with common languages (Java, Javascript, PHP, etc.)
 - DB tech should run anywhere (VM's, cloud, etc.)
 - Keeps essential features of RDBMS's while learning from key-value noSQL systems



History

- mongoDB = “**Humongous** DB”
 - Open-source
 - Document-based
 - “High performance, high availability”
 - Automatic scaling
 - C-P on CAP



Design Goals

- Scale horizontally over commodity systems
- Incorporate what works for RDBMSs
 - Rich data models, ad-hoc queries, full indexes
- Move away from what doesn't scale easily
 - Multi-row transactions, complex joins
- Use idiomatic development APIs
- Match agile development and deployment workflows



Key Features

- Data stored as documents (JSON)
 - Dynamic-schema
- Full CRUD support (Create, Read, Update, Delete)
 - Ad-hoc queries: Equality, RegEx, Ranges, Geospatial
 - Atomic in-place updates
- Full secondary indexes
 - Unique, sparse, TTL
- Replication – redundancy, failover
- Sharding – partitioning for read/write scalability



Getting Started with Mongo



MongoDB Drivers and Shell

Drivers

Drivers for most popular programming languages and frameworks



Shell

Command-line shell for interacting directly with database

```
> db.collection.insert({product:"MongoDB",
> type:"Document Database"})
> db.collection.findOne()
{
  "_id" : ObjectId("5106c1c2fc629bfe52792e86"),
  "product" : "MongoDB",
  "type" : "Document Database"
}
```



Installation

- Install Mongo from: <http://www.mongodb.org/downloads>
 - Extract the files
 - Create a data directory for Mongo to use
- Open your mongodb/bin directory and run the binary file (name depends on the architecture) to start the database server.
- To establish a connection to the server, open another command prompt window and go to the same directory, entering in mongo.exe or mongo for macs and Linuxes.
- This engages the mongodb shell—it's that easy!



MongoDB Design Model



Mongo Data Model

- Document-Based (max 16 MB)
- Documents are in BSON format, consisting of field-value pairs
- Each document stored in a collection
- Collections
 - Have index set in common
 - Like tables of relational db's.
 - Documents do not have to have uniform structure



JSON

- “JavaScript Object Notation”
- Easy for humans to write/read, easy for computers to parse/generate
- Objects can be nested
- Built on
 - name/value pairs
 - Ordered list of values



BSON

- “Binary JSON”
- Binary-encoded serialization of JSON-like docs
- Also allows “referencing”
- Embedded structure reduces need for joins
- Goals
 - Lightweight
 - Traversable
 - Efficient (decoding and encoding)



BSON Example

```
{  
  "_id": "37010"  
  "city": "ADAMS",  
  "pop": 2660,  
  "state": "TN",  
  "councilman": {  
    "name": "John Smith"  
    "address": "13 Scenic Way"  
  }  
}
```



The _id Field

- By default, each document contains an `_id` field. This field has a number of special characteristics:
 - Value serves as primary key for collection.
 - Value is unique, immutable, and may be any non-array type.
 - Default data type is `ObjectId`, which is “small, likely unique, fast to generate, and ordered.” Sorting on an `ObjectId` value is roughly equivalent to sorting on creation time.



BSON Types

Type	Number
Double	1
String	2
Object	3
Array	4
Binary data	5
Object id	7
Boolean	8
Date	9
Null	10
Regular Expression	11
JavaScript	13
Symbol	14
JavaScript (with scope)	15
32-bit integer	16
Timestamp	17
64-bit integer	18
Min key	255
Max key	127

The number can
be used with the
`$type` operator to
query by type!



MongoDB vs. Relational Databases



Why Databases Exist in the First Place?

- Why can't we just write programs that operate on objects?
 - Memory limit
 - We cannot swap back from disk merely by OS for the page based memory management mechanism
- Why can't we have the database operating on the same data structure as in program?
 - That is where Mongo comes in



RDBMS	→	MongoDB
Database	→	Database
Table	→	Collection
Row	→	Document
Index	→	Index
Join	→	Embedded Document
Foreign Key	→	Reference



Mongo is basically schema-free

- The purpose of schema in SQL is for meeting the requirements of tables and quirky SQL implementation
- Every “row” in a database “table” is a data structure, much like a “struct” in C, or a “class” in Java.
 - A table is then an array (or list) of such data structures
- So what we design in Mongo is basically similar to how we design a compound data type binding in JSON

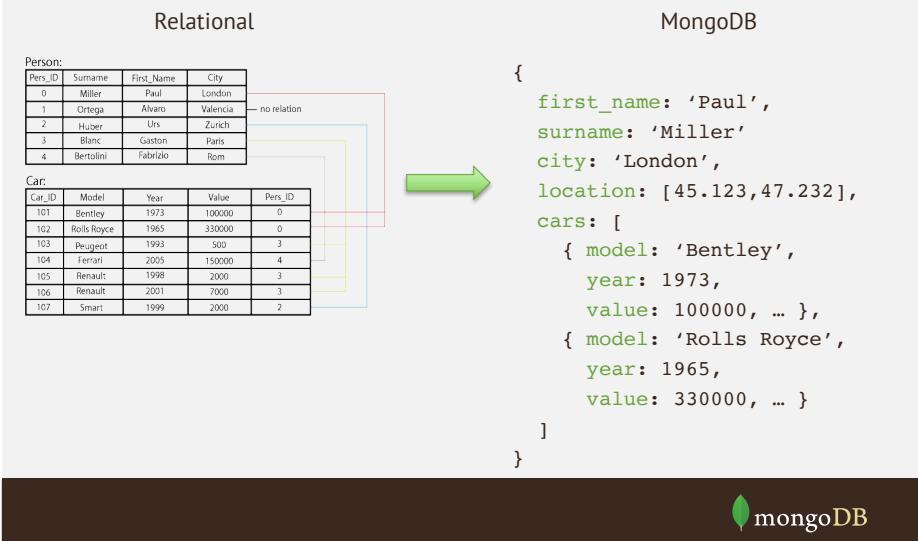


mongoDB vs. SQL

MongoDB	SQL
Document	Tuple
Collection	Table/View
PK: _id Field	PK: Any Attribute(s)
Uniformity not Required	Uniform Relation Schema
Index	Index
Embedded Structure	Joins
Shard	Partition



Document Oriented, Dynamic Schema



CRUD:

Create, Read, Update, Delete

CRUD: Using the Shell

- To check which db you're using → db
- Show all databases → show dbs
- Switch db's/make a new one → use <name>
- See what collections exist → show collections
- Note: db's are not actually created until you insert data!

CRUD: Using the Shell (cont.)

- To insert documents into a collection/make a new collection:

• db.<collection>.insert(<document>)

• <=>

• INSERT INTO <table>

• VALUES(<attributevalues>);

CRUD: Inserting Data

- Insert one document
 - db.<collection>.insert({<field>:<value>})
- Inserting a document with a field name new to the collection is inherently supported by the BSON model.
- To insert multiple documents, use an array.



CRUD: Querying

To match a specific value:

```
db.<collection>.find({<field>:<value>})
"AND"
db.<collection>.find({<field1>:<value1>, <field2>:<value2>
})
```

```
SELECT *
FROM <table>
WHERE <field1> = <value1> AND <field2> = <value2>;
```



CRUD: Querying

- Done on collections.
- Get all docs: db.<collection>.find()
 - Returns a cursor, which is iterated over shell to display first 20 results.
 - Add .limit(<number>) to limit results
 - SELECT * FROM <table>;
- Get one doc: db.<collection>.findOne()



CRUD: Querying

OR
db.<collection>.find({ \$or: [
<field>:<value1>
<field>:<value2>
]}

```
SELECT *
FROM <table>
WHERE <field> = <value1> OR <field> = <value2>;
```

Checking for multiple values of same field

```
db.<collection>.find({<field>: {$in [<value>, <value>]}})
```



CRUD: Querying

Including/excluding document fields

```
db.<collection>.find({<field1>:<value>}, {<field2>: 0})
```

```
SELECT field1  
FROM <table>;
```

```
db.<collection>.find({<field>:<value>}, {<field2>: 1})
```

Find documents with or w/o field

```
db.<collection>.find({<field>: { $exists: true}})
```



CRUD: Updating

To remove a field

```
db.<collection>.update({<field>:<value>},  
{ $unset: { <field>: 1}})
```

Replace all field-value pairs

```
db.<collection>.update({<field>:<value>},  
{ <field>:<value>, <field>:<value>})
```

*NOTE: This overwrites ALL the contents of a document, even removing fields.



CRUD: Updating

```
db.<collection>.update(  
{<field1>:<value1>}, //all docs in which field = value  
{$set: {<field2>:<value2>}}, //set field to value  
{multi:true}) //update multiple docs
```

Bulk.find.upsert(): if true, creates a new doc when none matches search criteria.

```
UPDATE <table>  
SET <field2> = <value2>  
WHERE <field1> = <value1>;
```



CRUD: Removal

Remove all records where field = value

```
db.<collection>.remove({<field>:<value>})
```

```
DELETE FROM <table>  
WHERE <field> = <value>;
```

As above, but only remove first document

```
db.<collection>.remove({<field>:<value>}, true)
```



CRUD: Isolation

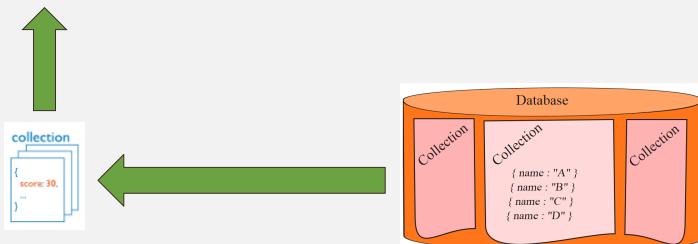
- By default, all writes are atomic only on the level of a single document.
- This means that, by default, all writes can be interleaved with other operations.
- You can isolate writes on an unsharded collection by adding `$isolated:1` in the query area:
 - `db.<collection>.remove({<field>:<value>, $isolated: 1})`



Before Index

- What does database normally do when we query?
 - MongoDB must scan every document.
 - Inefficient because process large volume of data

```
db.users.find( { score: { "$lt": 30} } )
```



Index in MongoDB



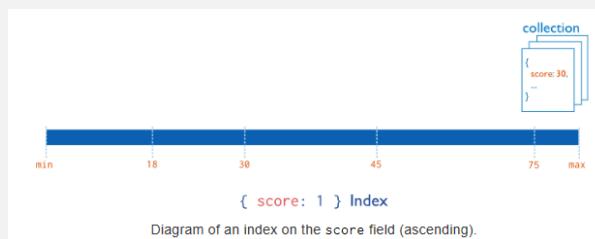
Index in MongoDB: Operations

- Creation index
 - `db.users.ensureIndex({ score: 1 })`
- Show existing indexes
 - `db.users.getIndexes()`
- Drop index
 - `db.users.dropIndex({score:1})`
- Explain—Explain
 - `db.users.find().explain()`
 - Returns a document that describes the process and indexes
- Hint
 - `db.users.find().hint({score: 1})`
 - Override MongoDB's default index selection



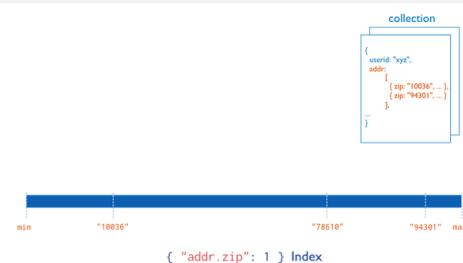
Index in MongoDB

- Types
 - Single Field Indexes
 - Compound Field Indexes
 - Multikey Indexes
- Single Field Indexes
 - db.users.ensureIndex({ score: 1 })



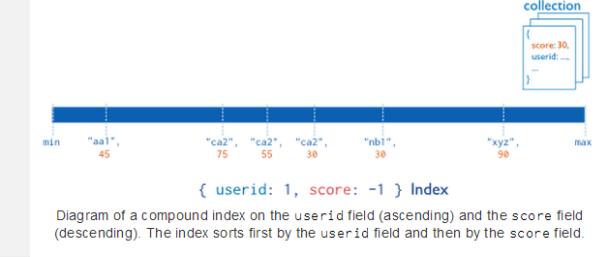
Index in MongoDB

- Types
 - Single Field Indexes
 - Compound Field Indexes
 - Multikey Indexes
- Compound Field Indexes
 - db.users.ensureIndex({ userid:1,score:-1 })



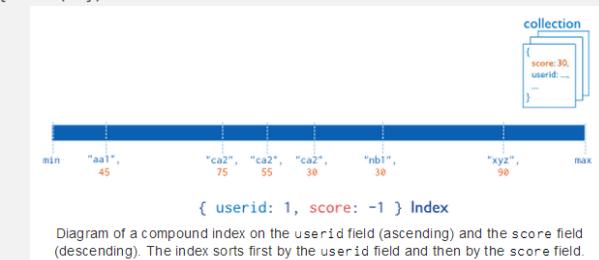
Index in MongoDB

- Types
 - Single Field Indexes
 - Compound Field Indexes
 - Multikey Indexes
- Compound Field Indexes
 - db.users.ensureIndex({ userid:1,score:-1 })



Index in MongoDB

- Types
 - Single Field Indexes
 - Compound Field Indexes
 - Multikey Indexes
- Multikey Indexes
 - db.users.ensureIndex({ addr.zip:1})



Mongo Examples



Querying

```
> db.addressBook.find()
{
  _id: ObjectId("4c4ba5c0672c685e5e8aabf3"),
  firstname: "John",
  lastname: "Smith",
  age: 25,
  address: {
    street: "21 2nd Street", city: "New York",
    state: "NY", zipcode: 10021
  }
}
// _id is unique but can be anything you like
```



Documents

```
> var new_entry = {
  firstname: "John",
  lastname: "Smith",
  age: 25,
  address: {
    street: "21 2nd Street",
    city: "New York",
    state: "NY",
    zipcode: 10021
  }
}
> db.addressBook.save(new_entry)
```



Indexes

```
// create an ascending index on "state"
> db.addressBook.ensureIndex({state:1})

> db.addressBook.find({state:"NY"})
{
  _id: ObjectId("4c4ba5c0672c685e5e8aabf3"),
  firstname: "John",
  ...
}

> db.addressBook.find({state:"NY", zip: 10021})
```



Queries

```
// Query Operators:  
// $all, $exists, $mod, $ne, $in, $nin, $nor, $or,  
// $size, $type, $lt, $lte, $gt, $gte  
  
// find contacts with any age  
> db.addressBook.find({age: {$exists: true}})  
  
// find entries matching a regular expression  
> db.addressBook.find( {lastname: /^smi*/i } )  
  
// count entries with "John"  
> db.addressBook.find( {firstname: 'John'} ).count()
```



Nested Documents

```
{  
  _id: ObjectId("4c4ba5c0672c685e5e8aabf3"),  
  firstname: "John", lastname: "Smith",  
  age: 25,  
  address: {  
    street: "21 2nd Street", city: "New York",  
    state: "NY", zipcode: 10021  
  }  
  phonenumbers : [ {  
    type: "mobile", number: "646-555-4567"  
  } ]  
}
```



Updates

```
// Update operators  
// $set, $unset, $inc, $push, $pushAll, $pull,  
// $pullAll, $bit  
  
> var new_phonenumber = {  
  type: "mobile",  
  number: "646-555-4567"  
}  
  
> db.addressBook.update({ _id: "..." }, {  
  $push: {phonenumbers: new_phonenumber}  
});
```



Secondary Indexes

```
// Index nested documents  
> db.addressBook.ensureIndex({"phonenumbers.type":1})  
  
// Geospatial indexes, 2d or 2dsphere  
> db.addressBook.ensureIndex({location: "2d"})  
> db.addressBook.find({location: {$near: [22,42]}})  
  
// Unique and Sparse indexes  
> db.addressBook.ensureIndex({field:1}, {unique:true})  
> db.addressBook.ensureIndex({field:1}, {sparse:true})
```



Additional Features

- Geospatial queries
 - Simple 2D plane
 - Or accounting for the surface of the earth (ellipsoid)
- Full Text Search
- Aggregation Framework
 - Similar to SQL GROUP BY operator
- Javascript MapReduce
 - Complex aggregation tasks



Open Source

- MongoDB source code is on Github
 - <https://github.com/mongodb/mongo>
- Issue tracking for MongoDB and drivers
 - <http://jira.mongodb.org>



MongoDB Development



Support

- Tickets are created by
 - Customer support
 - Community support (Google Groups, StackOverflow)
 - Community members
 - MongoDB employees
- Tickets can be voted on and watched to track progress
- Follow-the-Sun support
- All technical folks spend time doing community and customer support



Development

- Issues are triaged by CTO and engineering managers
- Then assigned into buckets, like
 - Specific version (ex. 2.7.1)
 - Desired version (ex. 2.7 desired)
 - Planning buckets
 - Unscheduled
- Engineers assign themselves tickets
- Once code is committed, a code review is needed



QA and Testing

- Code reviewer nominates for QA
- Unit tests are done by engineer
- Integration tests are done by QA team
- Support/Consulting/Architect teams do
 - Internal feature reviews/presentations
 - Beta testing with community and customers
- Documentation updates are linked to QA tickets



Questions?

- Sandeep Parikh
 - sap@mongodb.com
 - @crcsmnky
- MongoDB
 - MongoDB, drivers, documentation
 - <http://www.mongodb.org>
 - <http://docs.mongodb.org>
 - Free online training, presentations, whitepapers
 - <http://www.mongodb.com>

