

CSC 447: Parallel Programming for Multi-Core and Cluster Systems

Shared Parallel Programming Using OpenMP

Instructor: Haidar M. Harmanani

Spring 2021



OpenMP Recap

- Approaches to Parallelism Using OpenMP:
 - Fine-grained parallelism (loop-level)
 - Coarse-grained parallelism (task and section level)
- Fine-grained parallelism
 - Individual loops parallelized where each thread is assigned a unique range of the loop index
 - Multiple threads are spawned inside a parallel loop after parallel loop execution is serial
 - Easy to implement
- Coarse-grained parallelism
 - Based on parallel regions parallelism where multiple threads are started for parallel regions



Examples

Spring 2021

Parallel Programming for Multicore and Cluster Systems

3 | LAU
للمسيحية الأمريكية الجامعة Lebanese American University

Simple Example 1

Write a program that prints either “A race car” or “A car race” and maximize the parallelism

Spring 2021

Parallel Programming for Multicore and Cluster Systems

4 | LAU
للمسيحية الأمريكية الجامعة Lebanese American University

Simple Example 1

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("A ");
    printf("race ");
    printf("car ");

    printf("\n");
    return(0);
}
```

Output?

Parallel Simple Example 1

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main(int argc, char *argv[]) {
#pragma omp parallel
{
    printf("A ");
    printf("race ");
    printf("car ");
}
printf("\n");  return(0);
}
```

8 threads

A race car A race car A race car A
race car A race car A race car A
race car A race car

Parallel Simple Example 1 Using Single

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main(int argc, char *argv[]) {
#pragma omp parallel
{
    #pragma omp single
    {
        printf("A ");
        printf("race ");
        printf("car ");
    }
    printf("\n");  return(0);
}
```

Output?

Parallel Simple Example 1 Using Tasks

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main(int argc, char *argv[]) {
#pragma omp parallel
{
    #pragma omp single
    {
        printf("A ");
        #pragma omp task
        printf("race ");
        #pragma omp task
        printf("car ");
    }
}
printf("\n");  return(0);
}
```

Output?

A race car

Parallel Simple Example 1 Using Tasks

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main(int argc, char *argv[]) {
#pragma omp parallel
{
    #pragma omp single
    printf("A ");

    #pragma omp task
    printf("race ");

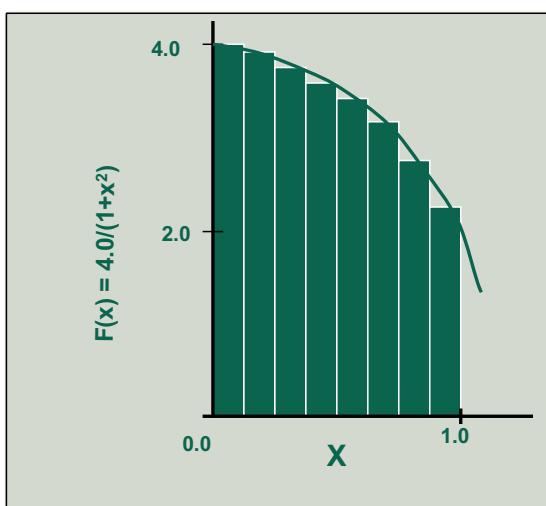
    #pragma omp task
    printf("car ");

    printf("\n");
return(0);
}
```

And Now?

A race car race car race race
race race race race car car car
car car car

Example 2: Compute PI



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i.

$$\Pi = \int_0^1 \frac{4}{(1+x^2)} dx$$

Compute PI : Serial Code

```

static long num_steps = 100000;
double step;
void main ()
{
    int i;
    double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;

    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}

```

$$\Pi = \int_0^1 \frac{4}{(1+x^2)} dx$$

Compute PI : Serial Code

```

double f(double x) {
    return (double)4.0 / ((double)1.0 + (x*x));
}
void computePi() {
    double h = (double) 1.0 / (double) iNumIntervals;
    double sum = 0, x;

    ...

    for (int i = 1; i <= iNumIntervals; i++) {
        x = h * ((double)i - (double)0.5);
        sum += f(x);
    }

    myPi = h * sum;
}

```

$$\Pi = \int_0^1 \frac{4}{(1+x^2)} dx$$

Compute PI : || Code

```

double f(double x) {
    return (double)4.0 / ((double)1.0 + (x*x));
}
void computePi() {
    double h = (double)1.0 / (double)iNumIntervals;
    double sum = 0, x;

#pragma omp parallel for private(x) reduction(+:sum)

    for (int i = 1; i <= iNumIntervals; i++) {
        x = h * ((double)i - (double)0.5);
        sum += f(x);
    }

    myPi = h * sum;
}

```

Compute PI : || Code

```

long num_steps=100000; double step;

void main()
{ int i;
  double x, sum = 0.0, pi;

  step = 1./(double)num_steps;
  start = clock();
#pragma omp parallel for private(x) reduction (+:sum)
  for (i=0; i<num_steps; i++)
  {
    x = (i + .5)*step;
    sum = sum + 4.0/(1.+ x*x);
  }

  pi = sum*step;
  stop = clock();

  printf("The value of PI is %15.12f\n",pi);
  printf("Time to calculate PI was %f seconds\n",((double)(stop - start)/1000.0));
  return 0;
}

```

$$\Pi = \int_0^1 \frac{4}{(1+x^2)} dx$$

Calculate Pi by integration

# Threads	Runtime [sec.]	Speedup
1	0.002877	1.00
2	0.001777	1.62
4	0.002988	0.96
8	0.002050	1.40
16	0.105787	1.26

Number of Iterations: 1000,000

PI value: 3.141593

Architecture: Intel i7, 3 GHz, running Mac OS 10.11.3 with 16GB RAM

Useful MacOS/Linux Commands

To compile

```
gcc -Wall -fopenmp -o pi pi.c
```

To set the number of threads to 4 using OMP_NUM_THREADS:

In the bash shell, type: `export OMP_NUM_THREADS=4`
in the c shell, type: `setenv OMP_NUM_THREADS 4`

You can set the number of threads to different values (2, 8, etc) using the same command
To run the OpenMP example code, simply type `./pi`

You can use the time command to evaluate the program's runtime:

Not very accurate but will do!

```
voyager-2:~ haidar$ export OMP_NUM_THREADS=4
voyager-2:~ haidar$ /usr/bin/time -p ./pi
The value of PI is 3.141592653590
The time to calculate PI was 18037.175000 seconds
real 6.52
user 17.97
sys 0.06
```

You can compare the running time of the OpenMP version with the serial version by compiling the serial version and repeating the above analysis

Useful MacOS/Linux Commands

- From within a shell, global adjustment of the number of threads:
 - `export OMP_NUM_THREADS=4`
 - `./pi`
- From within a shell, one-time adjustment of the number of threads:
 - `OMP_NUM_THREADS=4 ./pi`
- Intel Compiler on Linux: asking for more information:
 - `export KMP_AFFINITY=verbose`
 - `export OMP_NUM_THREADS=4`
 - `./pi`

Example 2: Recursive Fibonacci

```
int main(int argc, char* argv[])
{
    [...]
    fib(input);
    [...]
```

```
int fib(int n)  {
    if (n < 2) return n;
    int x = fib(n - 1);
    int y = fib(n - 2);
    return
        }  x+y;
```

Recursive Fibonacci: Discussion

```
int main(int argc, char* argv[])
{
    [...]
    fib(input);
    [...]
```

```
int fib(int n)  {
    if (n < 2) return n;
    int x = fib(n - 1);
    int y = fib(n - 2);
    return
} x+y;
```

- Only one Task / Thread should enter `fib()` from `main()`, it is responsible for creating the two initial work tasks
- `taskwait` is required, as otherwise `x` and `y` would be lost

Recursive Fibonacci: Attempt 0

```
int main(int argc, char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)  {
    if (n < 2) return n;
    int x, y;
    #pragma omp task
    {
        x = fib(n - 1);
    }
    #pragma omp task
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

What's wrong here?

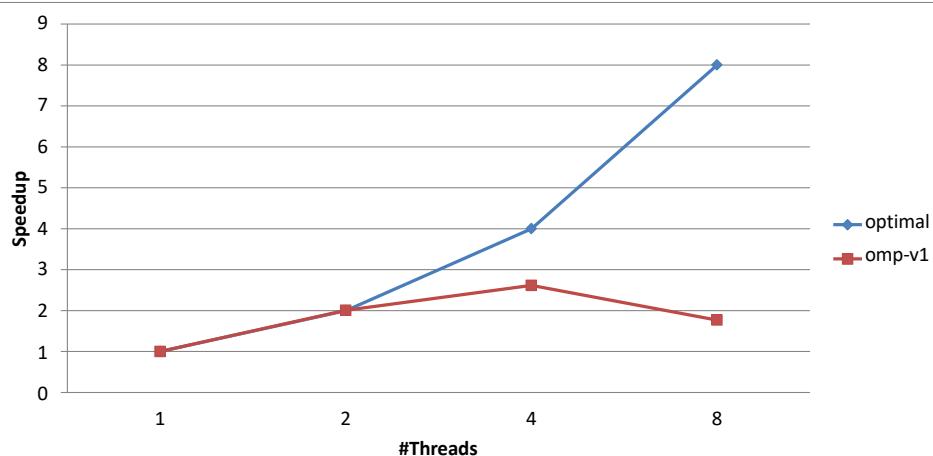
**x and y are private.
Can't use values of
private variables
outside of tasks**

Recursive Fibonacci: Attempt 1

```
int main(int argc, char* argv[])
{
    [...]
#pragma omp parallel
{
#pragma omp single
{
    fib(input);
}
}
[...]
}

int fib(int n) {
    if (n < 2) return n;
    int x, y;
#pragma omp task shared(x)
{
    x = fib(n - 1);
}
#pragma omp task shared(y)
{
    y = fib(n - 2);
}
#pragma omp taskwait
    return x+y;
}
```

Recursive Fibonacci: Attempt 1



Task creation overhead prevents better scalability

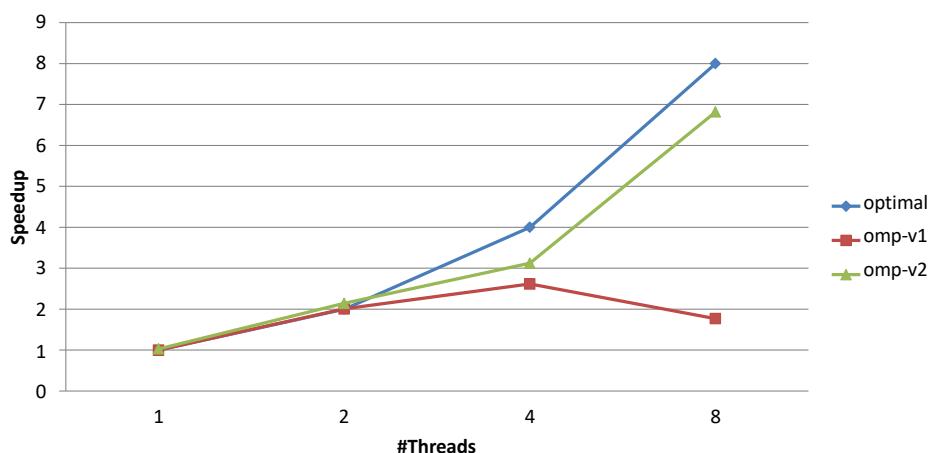
Recursive Fibonacci: Attempt 2

```
int main(int argc, char* argv[])
{
    [...]
#pragma omp parallel
{
#pragma omp single
{
    fib(input);
}
}
[...]
}

int fib(int n)  {
    if (n < 2) return n;
    int x, y;
#pragma omp task shared(x) if(n > 30)
{
    x = fib(n - 1);
}
#pragma omp task shared(y) if(n > 30)
{
    y = fib(n - 2);
}
#pragma omp taskwait
return x+y;
}
```

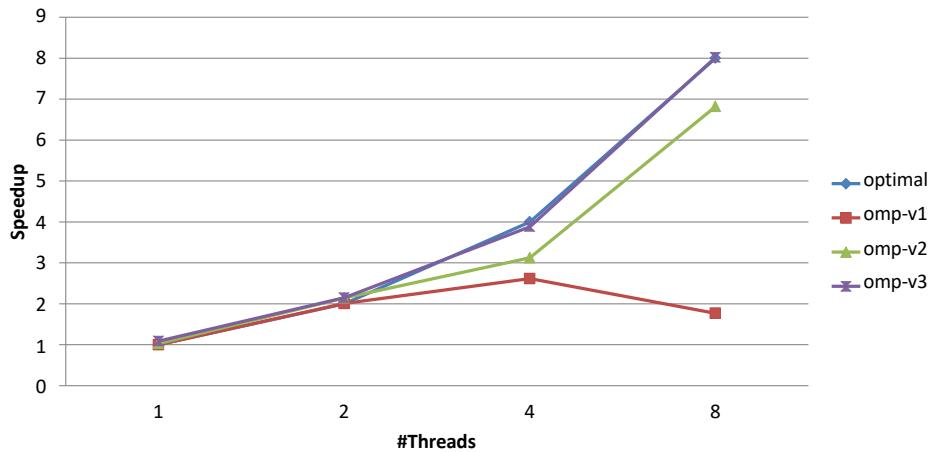
Don't create yet another task once a certain (small enough) n is reached

Recursive Fibonacci: Attempt 2



Overhead persists when running with 4 or 8 threads

Recursive Fibonacci: Attempt 3



Recursive Fibonacci: Attempt 3

```
int main(int argc,  char* argv[])
{
    [...]
#pragma omp parallel
{
#pragma omp single
{
    fib(input);
}
}
[...]
}

int fib(int n)  {
    if (n < 2)  return n;
if (n <= 30)
    return serfib(n);
int x, y;
#pragma omp task shared(x)
{
    x = fib(n - 1);
}
#pragma omp task shared(y)
{
    y = fib(n - 2);
}
#pragma omp taskwait
return x+y;
```

Skip the OpenMP overhead once a certain n is reached (no issue w/ production compilers)

Example: Linked List using Tasks

- A typical C code that implements a linked list pointer chasing code is:

```
while(p != NULL) {
    do_work(p->data);
    p = p->next;
}
```

- Can we implement the above code using tasks in order to parallelize the application?

Parallel Tree Traversal: What is Wrong?

```
void traverse (Tree *tree)
{
    #pragma omp task
    if(tree->left)
        traverse(tree->left);

    #pragma omp task
    if(tree->right)
        traverse(tree->right);
    process(tree);
}
```

Simple Example: || Linked List

```
List ml; //my_list
Element *e;
#pragma omp parallel
#pragma omp single
{
    for(e=ml->first;e;e=e->next)
        #pragma omp task
            process(e);
}
```

What's wrong here?

Possible data race !
Shared variable e
updated by multiple tasks

List Traversal

```
List ml; //my_list
Element *e;
#pragma omp parallel
#pragma omp single
{
    for(e=ml->first;e;e=e->next)
        #pragma omp task firstprivate(e)
            process(e);
}
```

Good solution – e is
firstprivate

List Traversal

```
List ml; //my_list
Element *e;
#pragma omp parallel
#pragma omp single private(e)
{
    for(e=ml->first;e;e=e->next)
        #pragma omp task
            process(e);
}
```

Good solution – e is private

List Traversal

```
List ml; //my_list
#pragma omp parallel
{
    Element *e;
    for(e=ml->first;e;e=e->next)
#pragma omp task
        process(e);
}
```

Good solution – e is private

Simple Example: || Linked List

```
#pragma omp parallel
// assume 8 threads
{
    #pragma omp single private(p)
    {
        ...
        while (p) {
            #pragma omp task
            {
                processwork(p);
            }
            p = p->next;
        }
    }
}
```

A pool of 8 threads is created here

One thread gets to execute the while loop

The single “while loop” thread creates a task for each instance of processwork()

Task Construct: Linked List

- A team of threads is created at the omp parallel construct
- A single thread is chosen to execute the while loop
 - o Lets call this thread “L”
- Thread L operates the while loop, creates tasks, and fetches next pointers
- Each time L crosses the omp task, construct it generates a new task and has a thread assigned to it
- Each task runs in its own thread
- All tasks complete at the barrier at the end of the parallel region’s single construct

```
#pragma omp parallel
{
    #pragma omp single
    { // block 1
        node * p = head;
        while (p) {
            //block 2
            #pragma omp task private(p)
            process(p);
            p = p->next; //block 3
        }
    }
}
```

Example: Linked List using Tasks

```
#pragma omp parallel
{
    #pragma omp single nowait
    {
        while (p != NULL) {
            #pragma omp task firstprivate(p)
            {
                processwork(p);
            }
            p = p->next;
        }
    }
}
```

Using OpenMP

PORTABLE SHARED MEMORY PARALLEL PROGRAMMING

Environment Variables

Environment Variables in OpenMP

- **OMP_NUM_THREADS**: This environment variable specifies the default number of threads created upon entering a parallel region.
- **OMP_SET_DYNAMIC**: Determines if the number of threads can be dynamically changed.
- **OMP_NESTED**: Turns on nested parallelism.
- **OMP_SCHEDULE**: Scheduling of for-loops if the clause specifies runtime

Environment Variables

Name	Possible Values	Most Common Default
OMP_NUM_THREADS	Non-negative Integer	1 or #cores
OMP_SCHEDULE	“schedule [, chunk]”	„static, (N/P)“
OMP_DYNAMIC	{TRUE FALSE}	TRUE
OMP_NESTED	{TRUE FALSE}	FALSE
OMP_STACKSIZE	“size [B K M G]”	-
OMP_WAIT_POLICY	{ACTIVE PASSIVE}	PASSIVE
OMP_MAX_ACTIVE_LEVELS	Non-negative Integer	-
OMP_THREAD_LIMIT	Non-negative Integer	1024
OMP_PROC_BIND	{TRUE FALSE}	FALSE
OMP_PLACES	Place List	-
OMP_CANCELLATION	{TRUE FALSE}	FALSE
OMP_DISPLAY_ENV	{TRUE FALSE}	FALSE
OMP_DEFAULT_DEVICE	Non-negative Integer	-

Using OpenMP

PORTABLE SHARED MEMORY PARALLEL PROGRAMMING

OMP False Sharing

Spring 2021

Parallel Programming for Multicore and Cluster Systems

39



False Sharing: OMP Example

```
double sum=0.0, sum_local[NUM_THREADS];
#pragma omp parallel num_threads(NUM_THREADS)
{
    int me = omp_get_thread_num();
    sum_local[me] = 0.0;
    #pragma omp for
    for (i = 0; i < N; i++)
        sum_local[me] += x[i] * y[i];
    #pragma omp atomic
    sum += sum_local[me];
}
```

Potential for false sharing on array sum_local

Spring 2021

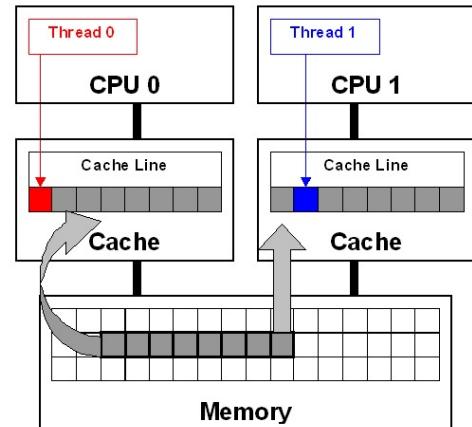
Parallel Programming for Multicore and Cluster Systems

40



False Sharing: OMP Example

- The `sum_local` array is dimensioned according to the number of threads and is small enough to fit in a single cache line
- When executed in parallel, the threads modify different, but adjacent, elements of `sum_local`
- The cache line is invalidated for all processors



False Sharing: Solution

- Solution: ensure that variables causing false sharing are spaced far enough apart in memory that they cannot reside on the same cache line
- See: [https://software.intel.com/sites/default/files/m/d/4/1/d/8/3-4-MemMgt_-
Avoiding_and_Identifying_False_Sharing_Among_Threads.pdf](https://software.intel.com/sites/default/files/m/d/4/1/d/8/3-4-MemMgt_-Avoiding_and_Identifying_False_Sharing_Among_Threads.pdf)

Nesting parallel Directives

- Nested parallelism can be enabled using the `OMP_NESTED` environment variable.
- If the `OMP_NESTED` environment variable is set to `TRUE`, nested parallelism is enabled.
- In this case, each parallel directive creates a new team of threads.

Using OpenMP

PORTABLE SHARED MEMORY PARALLEL PROGRAMMING

OMP Library Functions

API?

- You can use OpenMP set of pragmas
 - Safe as your code can be analyzed as regular serial code by your compiler if you do not use the -openmp flag.
 - It's flexible as you can define a lot of parameters in the code or from command line.
- But in some cases, you may want to write code aware of the OpenMP execution : use the OpenMP API on top of pragmas.

```
#include <omp.h>
int omp_get_thread_num(void);
int omp_get_num_threads(void);
```

OpenMP Library Functions

- Simple Lock routines:
 - A simple lock is available if it is unset.
 - `omp_init_lock()`, `omp_set_lock()`, `omp_unset_lock()`, `omp_test_lock()`, `omp_destroy_lock()`
- Nested Locks
 - A nested lock is available if it is unset or if it is set but owned by the thread executing the nested lock function
 - `omp_init_nest_lock()`, `omp_set_nest_lock()`, `omp_unset_nest_lock()`, `omp_test_nest_lock()`, `omp_destroy_nest_lock()`

A lock implies a memory fence of all thread visible variables

Synchronization: Lock Functions

- Protect resources with locks.

```
omp_lock_t lck;
omp_init_lock(&lck);
#pragma omp parallel private (tmp, id)
{
    id = omp_get_thread_num();
    tmp = do_lots_of_work(id);
    omp_set_lock(&lck);
    printf("%d %d", id, tmp);
    omp_unset_lock(&lck);
}
omp_destroy_lock(&lck);
```

Wait here for your turn.

Release the lock so the next thread gets a turn.

Free-up storage when done.

OpenMP Library Functions

- In addition to directives, OpenMP also supports a number of functions that allow a programmer to control the execution of threaded programs.

```
/* thread and processor count */
void omp_set_num_threads (int num_threads);
int omp_get_num_threads ();
int omp_get_max_threads ();
int omp_get_thread_num ();
int omp_get_num_procs ();
int omp_in_parallel();
```

OpenMP Library Functions

```
/* controlling and monitoring thread creation */
void omp_set_dynamic (int dynamic_threads);
int omp_get_dynamic ();
void omp_set_nested (int nested);
int omp_get_nested ();
```

- In addition, all lock routines also have a nested lock counterpart for recursive mutexes.