

CSC 611: Analysis of Algorithms

Lecture 8

Heaps, Priority Queues, and HeapSort

A Job Scheduling Application

- Job scheduling
 - The key is the priority of the jobs in the queue
 - The job with the highest priority needs to be executed next
- Operations
 - Insert, remove maximum
- Data structures
 - **Priority queues**
 - Ordered array/list, unordered array/list

PQ Implementations & Cost

Worst-case asymptotic costs for a PQ with N items


	Insert	Remove max
ordered array	N	1
ordered list	N	1
unordered array	1	N
unordered list	1	N

Can we implement both operations efficiently?

CSC611/ Lecture08

Background on Trees


- *Def:* Binary tree = structure composed of a finite set of nodes that either:
 - Contains no nodes, or
 - Is composed of three disjoint sets of nodes: a **root** node, a **left subtree** and a **right subtree**




CSC611/ Lecture08

Special Types of Trees

- **Def: Full binary tree** = a binary tree in which each node is either a leaf or has degree (number of children) exactly 2.




- **Def: Complete binary tree** = a binary tree in which all leaves have the same depth and all internal nodes have degree 2.



CSC611/ Lecture08

The Heap Data Structure

- **Def:** A **heap** is a nearly complete binary tree with the following two properties:
 - **Structural property:** all levels are full, except possibly the last one, which is filled from left to right
 - **Order (heap) property:** for any node x
 $\text{Parent}(x) \geq x$




It doesn't matter that 4 in level 1 is smaller than 5 in level 2

CSC611/ Lecture08

Definitions



- **Height** of a node = the number of edges on a longest simple path from the node down to a leaf
- **Depth** of a node = the length of a path from the root to the node
- **Height** of tree = height of root node
 $= \lfloor \lg n \rfloor$, for a heap of n elements



CSC611/ Lecture08

Array Representation of Heaps

- A heap can be stored as an array A .
 - Root of tree is $A[1]$
 - Left child of $A[i] = A[2i]$
 - Right child of $A[i] = A[2i + 1]$
 - Parent of $A[i] = A[\lfloor i/2 \rfloor]$
 - $\text{Heapsize}[A] \leq \text{length}[A]$
- The elements in the subarray $A[(\lfloor n/2 \rfloor + 1) .. n]$ are leaves
- The root is the maximum element of the heap



CSC611/ Lecture08

Heap Types

- **Max-heaps** (largest element at root), have the *max-heap property*:

- for all nodes i , excluding the root:

$$A[\text{PARENT}(i)] \geq A[i]$$

- **Min-heaps** (smallest element at root), have the *min-heap property*:

- for all nodes i , excluding the root:

$$A[\text{PARENT}(i)] \leq A[i]$$

CSC611/ Lecture08

Operations on Heaps

- Maintain the max-heap property
 - MAX-HEAPIFY
- Create a max-heap from an unordered array
 - BUILD-MAX-HEAP
- Sort an array in place
 - HEAPSORT
- Priority queue operations

CSC611/ Lecture08


Operations on Priority Queues

- Max-priority queues support the following operations:
 - **INSERT(S, x)**: inserts element x into set S
 - **EXTRACT-MAX(S)**: removes and returns element of S with largest key
 - **MAXIMUM(S)**: returns element of S with largest key
 - **INCREASE-KEY(S, x, k)**: increases value of element x 's key to k (assume $k \geq$ current key value at x)

CSC611/ Lecture08

Maintaining the Heap Property

- Suppose a node is smaller than a child
 - Left and Right subtrees of i are max-heaps
- Invariant:
 - the heap condition is violated only at that node
- To eliminate the violation:
 - Exchange with larger child
 - Move down the tree
 - Continue until node is not smaller than children




CSC611/ Lecture08

Maintaining the Heap Property

• Assumptions: \mathcal{Alg} : MAX-HEAPIFY(A, i, n)

- Left and Right subtrees of i are max-heaps
- $A[i]$ may be smaller than its children




1. $l \leftarrow \text{LEFT}(i)$
2. $r \leftarrow \text{RIGHT}(i)$
3. **if** $l \leq n$ and $A[l] > A[i]$
4. **then** $\text{largest} \leftarrow l$
5. **else** $\text{largest} \leftarrow i$
6. **if** $r \leq n$ and $A[r] > A[\text{largest}]$
7. **then** $\text{largest} \leftarrow r$
8. **if** $\text{largest} \neq i$
9. **then** exchange $A[i] \leftrightarrow A[\text{largest}]$
10. MAX-HEAPIFY($A, \text{largest}, n$)

CSC611/ Lecture08


Example

MAX-HEAPIFY($A, 2, 10$)




$A[2]$ violates the heap property

$A[2] \leftrightarrow A[4]$



$A[4]$ violates the heap property

$A[4] \leftrightarrow A[9]$



Heap property restored

CSC611/ Lecture08

MAX-HEAPIFY Running Time

- Intuitively:
 - A heap is an almost complete binary tree \Rightarrow must process $O(\lg n)$ levels, with constant work at each level
- Running time of MAX-HEAPIFY is $O(\lg n)$
- Can be written in terms of the height of the heap, as being $O(h)$
 - Since the height of the heap is $\lfloor \lg n \rfloor$


CSC611/ Lecture08

Building a Heap

- Convert an array $A[1 \dots n]$ into a max-heap ($n = \text{length}[A]$)
- The elements in the subarray $A[\lfloor n/2 \rfloor + 1 \dots n]$ are leaves
- Apply MAX-HEAPIFY on elements between 1 and $\lfloor n/2 \rfloor$

Alg: BUILD-MAX-HEAP(A)

1. $n = \text{length}[A]$
2. **for** $i \leftarrow \lfloor n/2 \rfloor$ **downto** 1
3. **do** MAX-HEAPIFY(A, i, n)




A:

4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---

Example:

A

4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---



CSC611/ Lecture08


Correctness of BUILD-MAX-HEAP

- **Loop invariant:**

- At the start of each iteration of the **for** loop, each node $i + 1, i + 2, \dots, n$ is the root of a max-heap

- **Initialization:**

- $i = \lfloor n/2 \rfloor$: Nodes $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$ are leaves
 \Rightarrow they are the root of trivial max-heaps




CSC611/ Lecture08

Correctness of BUILD-MAX-HEAP

- **Maintenance:**

- MAX-HEAPIFY makes node i a max-heap root and preserves the property that nodes $i + 1, i + 2, \dots, n$ are roots of max-heaps
- Decrementing i in the for loop reestablishes the loop invariant



- **Termination:**

- $i = 0 \Rightarrow$ each node $1, 2, \dots, n$ is the root of a max-heap (by the loop invariant)

CSC611/ Lecture08

Running Time of BUILD MAX HEAP

Alg: BUILD-MAX-HEAP(A)

1. $n = \text{length}[A]$
 2. **for** $i \leftarrow \lfloor n/2 \rfloor$ **downto** 1
 3. **do** MAX-HEAPIFY(A, i , n)
- $O(\lg n) \} O(n)$

\Rightarrow It would seem that running time is $O(n \lg n)$


- This is not an asymptotically tight upper bound

CSC611/ Lecture08

Running Time of BUILD MAX HEAP

- HEAPIFY takes $O(h) \Rightarrow$ the cost of HEAPIFY on a node i is proportional to the height of the node i in the tree

$$\Rightarrow T(n) = \sum_{i=0}^h n_i h_i = \sum_{i=0}^h 2^i (h-i) = O(n)$$



$h_i = h - i$ height of the heap rooted at level i

$n_i = 2^i$ number of nodes at level i

CSC611/ Lecture08

Running Time of BUILD MAX HEAP

$$\begin{aligned}
 T(n) &= \sum_{i=0}^h n_i h_i && \text{Cost of HEAPIFY at level } i \times \text{number of nodes at that level} \\
 &= \sum_{i=0}^h 2^i (h-i) && \text{Replace the values of } n_i \text{ and } h_i \text{ computed before} \\
 &= \sum_{i=0}^h \frac{h-i}{2^{h-i}} 2^h && \text{Multiply by } 2^h \text{ both at the numerator and denominator} \\
 &&& \text{and write } 2^i \text{ as } \frac{1}{2^{-i}} \\
 &= 2^h \sum_{k=0}^h \frac{k}{2^k} && \text{Change variables: } k = h - i \\
 &\leq n \sum_{k=0}^{\infty} \frac{k}{2^k} && \text{The sum above is smaller than the sum of all elements to } \infty \\
 &&& \text{and } h = \lg n \\
 &= O(n) && \text{The sum above is smaller than 2}
 \end{aligned}$$

$$\sum_{k=0}^{\infty} kx^k = \frac{x}{(1-x)^2}$$

for $|x| < 1$.

Running time of BUILD-MAX-HEAP: $T(n) = O(n)$

CSC611/ Lecture08

Operations on Priority Queues

- Max-priority queues support the following operations:
 - **INSERT(S, x)**: inserts element x into set S
 - **EXTRACT-MAX(S)**: removes and returns element of S with largest key
 - **MAXIMUM(S)**: returns element of S with largest key
 - **INCREASE-KEY(S, x, k)**: increases value of element x 's key to k (assume $k \geq$ current key value at x)

CSC611/ Lecture08

HEAP-MAXIMUM

Goal:


- Return the largest element of the heap

Alg: **HEAP-MAXIMUM(A)**

Running time: $O(1)$

1. **return $A[1]$**

Heap A :



Heap-Maximum(A) returns 7

CSC611/ Lecture08

HEAP-EXTRACT-MAX


Goal:

- Extract the largest element of the heap (i.e., return the max value and also remove that element from the heap)

Idea:

- Exchange the root element with the last
- Decrease the size of the heap by 1 element
- Call MAX-HEAPIFY on the new root, on a heap of size $n-1$


Heap A:



HEAP-EXTRACT-MAX


Alg: HEAP-EXTRACT-MAX(A, n)

1. **if** $n < 1$
2. **then error** "heap underflow"
3. $\text{max} \leftarrow A[1]$
4. $A[1] \leftarrow A[n]$
5. MAX-HEAPIFY($A, 1, n-1$)
6. **return** max




Running time: $O(\lg n)$

Example: HEAP-EXTRACT-MAX




max = 16



Heap size decreased with 1


Call MAX-HEAPIFY(A, 1, n-1)



CSC611/ Lecture08

HEAP-INCREASE-KEY

- Goal:
 - Increases the key of an element i in the heap
- Idea:
 - Increment the key of $A[i]$ to its new value
 - If the max-heap property does not hold anymore: traverse a path toward the root to find the proper place for the newly increased key




CSC611/ Lecture08

HEAP-INCREASE-KEY

Alg: HEAP-INCREASE-KEY(A, i, key)





1. **if** $\text{key} < A[i]$
2. **then error** "new key is smaller than current key"
3. $A[i] \leftarrow \text{key}$
4. **while** $i > 1$ and $A[\text{PARENT}(i)] < A[i]$
5. **do** exchange $A[i] \Leftrightarrow A[\text{PARENT}(i)]$
6. $i \leftarrow \text{PARENT}(i)$

- Running time: $O(\lg n)$



CSC611/ Lecture08

Example: HEAP-INCREASE-KEY



CSC611/ Lecture08



MAX-HEAP-INSERT

- Goal:

- Inserts a new element into a max-heap

- Idea:

- Expand the max-heap with a new element whose key is $-\infty$
- Calls HEAP-INCREASE-KEY to set the key of the new node to its correct value and maintain the max-heap property




CSC611/ Lecture08

MAX-HEAP-INSERT

Alg: MAX-HEAP-INSERT(A , key , n)

1. $heap-size[A] \leftarrow n + 1$
2. $A[n + 1] \leftarrow -\infty$
3. HEAP-INCREASE-KEY(A , $n + 1$, key)




Running time: $O(\lg n)$


CSC611/ Lecture08

Example: MAX-HEAP-INSERT


Insert value 15:
- Start by inserting $-\infty$




Increase the key to 15
Call HEAP-INCREASE-KEY on $A[11] = 15$



The restored heap containing
the newly added element




CSC611/ Lecture08



Summary

- We can perform the following operations on heaps:
 - MAX-HEAPIFY $O(\lg n)$
 - BUILD-MAX-HEAP $O(n)$
 - HEAP-SORT $O(n \lg n)$
 - MAX-HEAP-INSERT $O(\lg n)$
 - HEAP-EXTRACT-MAX $O(\lg n)$
 - HEAP-INCREASE-KEY $O(\lg n)$
 - HEAP-MAXIMUM $O(1)$

Readings



- Chapters 14