


CSC 611: Analysis of Algorithms

Lecture 9

Red-Black Trees, OS Trees, and Interval Trees

Binary Search Trees

- Tree representation:
 - A linked data structure in which each node is an object
- Node representation:
 - **Key** field
 - Satellite data
 - **Left**: pointer to left child
 - **Right**: pointer to right child
 - **p**: pointer to parent ($p[\text{root}[T]] = \text{NIL}$)
- Satisfies the binary search tree property




Binary Search Tree Example

- Binary search tree property:

- If y is in left subtree of x ,
then $\text{key}[y] \leq \text{key}[x]$

- If y is in right subtree of x ,
then $\text{key}[y] \geq \text{key}[x]$



CSC611/ Lecture09

Binary Search Trees

- Support many dynamic set operations
 - SEARCH, MINIMUM, MAXIMUM, PREDECESSOR, SUCCESSOR, INSERT, DELETE
- Running time of basic operations on binary search trees
 - On average: $\Theta(\lg n)$
 - The expected height of the tree is $\lg n$
 - In the worst case: $\Theta(n)$
 - The tree is a linear chain of n nodes

CSC611/ Lecture09

Red-Black Trees

- “Balanced” binary trees guarantee an $O(\lg n)$ running time on the basic dynamic-set operations
- Red-black tree
 - Binary tree with an additional attribute for its nodes: **color** which can be **red** or **black**
 - Constrains the way nodes can be colored on any path from the root to a leaf
 - Ensures that no path is more than twice as long as another \Rightarrow the tree is balanced
 - The nodes inherit all the other attributes from the binary-search trees: **key**, **left**, **right**, **p**


CSC611/ Lecture09

Red-Black Trees Properties

1. Every **node** is either **red** or **black**
2. The **root** is **black**
3. Every **leaf** (**NIL**) is **black**
4. If a node is red, then both its children are black
 - No two red nodes in a row on a simple path from the root to a leaf
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes

CSC611/ Lecture09


Example: RED-BLACK TREE



- For convenience we use a sentinel **NIL[T]** to represent all the **NIL** nodes at the leafs
 - **NIL[T]** has the same fields as an ordinary node
 - **Color[NIL[T]] = BLACK**
 - The other fields may be set to arbitrary values

CSC611/ Lecture09

Black-Height of a Node



- **Height of a node:** the number of edges in a longest path to a leaf
- **Black-height** of a node **x**: **bh(x)** is the number of black nodes (including NIL) on a path from **x** to leaf, not counting **x**

CSC611/ Lecture09


Properties of Red-Black Trees

- **Claim**

- Any node with height h has black-height $\geq h/2$

- **Proof**

- By property 4, there are at most $h/2$ **red** nodes on the path from the node to a leaf
- Hence at least $h/2$ are **black**



Property 4: if a node is **red** then both its children are black

CSC611/ Lecture09

Properties of Red-Black Trees

Claim

The subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes

Proof: By induction on height of x

Basis: $height[x] = 0 \Rightarrow$

x is a leaf ($NIL[T]$) \Rightarrow

$bh(x) = 0 \Rightarrow$

of internal nodes: $2^0 - 1 = 0$




CSC611/ Lecture09

Properties of Red-Black Trees

Inductive step:


- Let $\text{height}(x) = h$ and $\text{bh}(x) = b$
- Any child y of x has:
 - $\text{bh}(y) = b$ (if the child is **red**), or
 - $\text{bh}(y) = b - 1$ (if the child is **black**)



CSC611/ Lecture09

Properties of Red-Black Trees

- Want to prove:
 - The subtree rooted at any node x contains at least $2^{\text{bh}(x)} - 1$ internal nodes
- Assume true for children of x :
 - Their subtrees contain at least $2^{\text{bh}(x) - 1} - 1$ internal nodes
- The subtree rooted at x contains at least:
$$(2^{\text{bh}(x) - 1} - 1) + (2^{\text{bh}(x) - 1} - 1) + 1 =$$
$$2 \cdot (2^{\text{bh}(x) - 1} - 1) + 1 =$$
$$2^{\text{bh}(x)} - 1 \text{ internal nodes}$$




CSC611/ Lecture09

Properties of Red-Black Trees

Lemma: A red-black tree with n internal nodes has height at most $2\lg(n + 1)$.

Proof:

$$\begin{array}{lcl} n & \geq 2^b - 1 & \geq 2^{h/2} - 1 \\ \text{number } n & & \text{since } b \geq h/2 \\ \text{of internal} & & \\ \text{nodes} & & \end{array}$$



- Add 1 to all sides and then take logs:

$$\begin{aligned} n + 1 &\geq 2^b \geq 2^{h/2} \\ \lg(n + 1) &\geq h/2 \Rightarrow \\ h &\leq 2 \lg(n + 1) \end{aligned}$$

CSC611/ Lecture09

Operations on Red-Black Trees

- The non-modifying binary-search tree operations **MINIMUM**, **MAXIMUM**, **SUCCESSOR**, **PREDECESSOR**, and **SEARCH** run in $O(h)$ time
 - They take $O(\lg n)$ time on red-black trees
- What about TREE-INSERT and TREE-DELETE?
 - They will still run in $O(\lg n)$
 - We have to guarantee that the modified tree will still be a red-black tree

CSC611/ Lecture09

Insertion


- Goal:
 - Insert a new node z into a red-black tree
- Idea:
 - Insert node z into the tree as for an ordinary binary search tree
 - Color the node **red**
 - Restore the red-black tree properties
 - Use an auxiliary procedure RB-INSERT-FIXUP

CSC611/ Lecture09

INSERT


INSERT: what color to make the new node?

- Red? Let's insert 35!
 - Property 4: if a node is red, then both its children are black
- Black? Let's insert 14!
 - Property 5: all paths from a node to its leaves contain the same number of black nodes



CSC611/ Lecture09

DELETE



DELETE: what color was the node that was removed? **Red?**

1. Every **node** is either **red** or **black** OK!
2. The **root** is **black** OK!
3. Every **leaf** (NIL) is **black** OK!
4. If a node is red, then both its children are black


OK! Does not change any black heights

OK! Does not create two red nodes in a row

5. For each node, all paths from the node to descendant leaves contain the same number of black nodes

CSC611/ Lecture09

DELETE



DELETE: what color was the node that was removed? **Black?**

1. Every **node** is either **red** or **black** OK!
2. The **root** is **black** **Not OK!** If removing the root and the child that replaces it is **red**
3. Every **leaf** (NIL) is **black** OK!
4. If a node is red, then both its children are black

Not OK! Could change the black heights of some nodes

Not OK! Could create two red nodes in a row

5. For each node, all paths from the node to descendant leaves contain the same number of black nodes

CSC611/ Lecture09

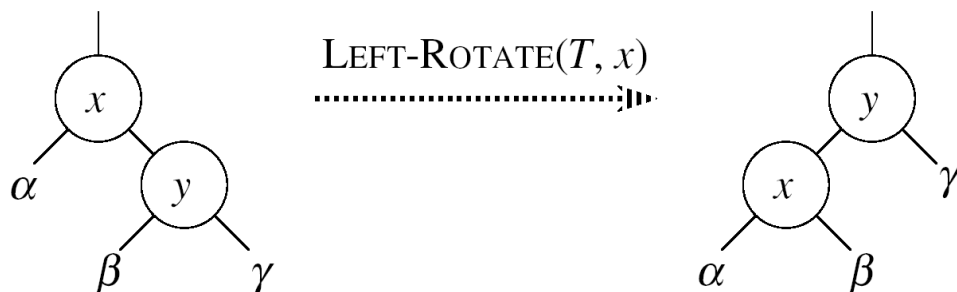
Rotations

- Operations for restructuring the tree after insert and delete operations on red-black trees
- Rotations take a red-black tree and a node within the tree and:
 - Together with some node re-coloring they help restore the red-black tree property
 - Change some of the pointer structure
 - Do not change the binary-search tree property
- Two types of rotations:
 - Left & right rotations

CSC611/ Lecture09

Left Rotations


- Assumption for a left rotation on a node x :
 - The right child of x (y) is not NIL



- Idea:
 - Pivots around the link from x to y
 - Makes y the new root of the subtree
 - x becomes y 's left child
 - y 's left child becomes x 's right child

CSC611/ Lecture09

Example: LEFT-ROTATE




LEFT-ROTATE(T, x)

1. $y \leftarrow \text{right}[x]$ ► Set y
2. $\text{right}[x] \leftarrow \text{left}[y]$ ► y 's left subtree becomes x 's right subtree
3. **if** $\text{left}[y] \neq \text{NIL}$
4. **then** $p[\text{left}[y]] \leftarrow x$ ► Set the parent relation from $\text{left}[y]$ to x
5. $p[y] \leftarrow p[x]$ ► The parent of x becomes the parent of y
6. **if** $p[x] = \text{NIL}$
7. **then** $\text{root}[T] \leftarrow y$
8. **else if** $x = \text{left}[p[x]]$
9. **then** $\text{left}[p[x]] \leftarrow y$
10. **else** $\text{right}[p[x]] \leftarrow y$
11. $\text{left}[y] \leftarrow x$ ► Put x on y 's left
12. $p[x] \leftarrow y$ ► y becomes x 's parent

Right Rotations

- Assumption for a right rotation on a node x :
 - The left child of y (x) is not NIL



- Idea:
 - Pivots around the link from y to x
 - Makes x the new root of the subtree
 - y becomes x 's right child
 - x 's right child becomes y 's left child


CSC611/ Lecture09

Insertion

- Goal:
 - Insert a new node z into a red-black tree
- Idea:
 - Insert node z into the tree as for an ordinary binary search tree
 - Color the node **red**
 - Restore the red-black tree properties
 - Use an auxiliary procedure RB-INSERT-FIXUP

CSC611/ Lecture09


RB-INSERT(T, z)



1. $y \leftarrow \text{NIL}$
 2. $x \leftarrow \text{root}[T]$
 3. **while** $x \neq \text{NIL}$
 4. **do** $y \leftarrow x$
 5. **if** $\text{key}[z] < \text{key}[x]$
 6. **then** $x \leftarrow \text{left}[x]$
 7. **else** $x \leftarrow \text{right}[x]$
 8. $p[z] \leftarrow y$
- Initialize nodes x and y
 - Throughout the algorithm y points to the parent of x
 - Go down the tree until reaching a leaf
 - At that point y is the parent of the node to be inserted
 - Sets the parent of z to be y

CSC611/ Lecture09

RB-INSERT(T, z)



9. **if** $y = \text{NIL}$
 10. **then** $\text{root}[T] \leftarrow z$
 11. **else if** $\text{key}[z] < \text{key}[y]$
 12. **then** $\text{left}[y] \leftarrow z$
 13. **else** $\text{right}[y] \leftarrow z$
 14. $\text{left}[z] \leftarrow \text{NIL}$
 15. $\text{right}[z] \leftarrow \text{NIL}$
 16. $\text{color}[z] \leftarrow \text{RED}$
 17. $\text{RB-INSERT-FIXUP}(T, z)$
- The tree was empty: set the new node to be the root
 - Otherwise, set z to be the left or right child of y , depending on whether the inserted node is smaller or larger than y 's key
 - Set the fields of the newly added node
 - Fix any inconsistencies that could have been introduced by adding this new red node


CSC611/ Lecture09

RB Properties Affected by Insert

1. Every **node** is either **red** or **black** OK!
2. The **root** is **black** If z is the root \Rightarrow **not OK**
3. Every **leaf** (NIL) is **black** OK!
4. If a node is red, then both its children are black

If $p(z)$ is red \Rightarrow **not OK**
 z and $p(z)$ are both red

- OK!
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes




CSC611/ Lecture09

RB-INSERT-FIXUP – Case 1

z 's "uncle" (y) is **red**

Idea: (z is a right child)

- $p[p[z]]$ (z 's grandparent) must be black: z and $p[z]$ are both red
- Color $p[z]$ **black**
- Color y **black**
- Color $p[p[z]]$ **red**
 - Push the **red** node up the tree
- Make $z = p[p[z]]$



CSC611/ Lecture09

RB-INSERT-FIXUP – Case 1




z's "uncle" (y) is **red**

Idea: (z is a left child)

- $p[p[z]]$ (z's grandparent) must be black: z and $p[z]$ are both red

- $color[p[z]] \leftarrow \text{black}$
- $color[y] \leftarrow \text{black}$
- $color[p[p[z]]] \leftarrow \text{red}$
- $z = p[p[z]]$ Case1

– Push the **red** node up the tree



CSC611/ Lecture09




RB-INSERT-FIXUP – Case 3

Case 3:

- z's "uncle" (y) is **black**
- z is a left child

Idea:

- $color[p[z]] \leftarrow \text{black}$
- $color[p[p[z]]] \leftarrow \text{red}$
- $\text{RIGHT-ROTATE}(T, p[p[z]])$ Case3
- No longer have 2 reds in a row
- $p[z]$ is now black



CSC611/ Lecture09

RB-INSERT-FIXUP – Case 2


Case 2:

- z 's "uncle" (y) is **black**
- z is a right child

Idea:

- $z \leftarrow p[z]$
- $\text{LEFT-ROTATE}(T, z)$ Case2

\Rightarrow now z is a left child, and both z and $p[z]$ are red \Rightarrow case 3




CSC611/ Lecture09

RB-INSERT-FIXUP(T, z)

1. **while** $\text{color}[p[z]] = \text{RED}$ \leftarrow The while loop repeats only when case1 is executed: $O(\lg n)$ times
2. **do if** $p[z] = \text{left}[p[p[z]]]$
3. **then** $y \leftarrow \text{right}[p[p[z]]]$ } Set the value of x 's "uncle"
4. **if** $\text{color}[y] = \text{RED}$
5. **then Case1**
6. **else if** $z = \text{right}[p[z]]$
7. **then Case2**
8. **Case3**
9. **else** (same as **then** clause with "right" and "left" exchanged)
10. $\text{color}[\text{root}[T]] \leftarrow \text{BLACK}$ \leftarrow We just inserted the root, or the red node reached the root

CSC611/ Lecture09

Example



Analysis of RB-INSERT

- Inserting the new element into the tree
 $O(\lg n)$
- RB-INSERT-FIXUP
 - The while loop repeats only if CASE 1 is executed
 - The number of times the while loop can be executed is $O(\lg n)$
- Total running time of RB-INSERT: $O(\lg n)$

Red-Black Trees - Summary

- Operations on red-black trees:
 - SEARCH $O(h)$
 - PREDECESSOR $O(h)$
 - SUCCESSION $O(h)$
 - MINIMUM $O(h)$
 - MAXIMUM $O(h)$
 - INSERT $O(h)$
 - DELETE $O(h)$
- Red-black trees guarantee that the height of the tree will be $O(\lg n)$

CSC611/ Lecture09

Augmenting Data Structures

- Let's look at two new problems:
 - Dynamic order statistic
 - Interval search
- It is unusual to have to design all-new data structures from scratch
 - Typically: store additional information in an already known data structure
 - The augmented data structure can support new operations
- We need to correctly maintain the new information without loss of efficiency

CSC611/ Lecture09

Dynamic Order Statistics

- *Def.:* the i -th order statistic of a set of n elements, where $i \in \{1, 2, \dots, n\}$ is the element with the i -th smallest key.
- We can retrieve an order statistic from an unordered set:
 - Using: RANDOMIZED-SELECT
 - In: $O(n)$ time
- We will show that:
 - With red-black trees we can achieve this in $O(\lg n)$
 - Finding the **rank** of an element takes also $O(\lg n)$

CSC611/ Lecture09


Order-Statistic Tree

- *Def.:* **Order-statistic tree:** a red-black tree with additional information stored in each node
- Node representation:
 - Usual fields: `key[x]`, `color[x]`, `p[x]`, `left[x]`, `right[x]`
 - Additional field: **size[x]** that contains the number of (internal) nodes in the subtree rooted at x (including x itself)
- For any internal node of the tree:

$$\text{size}[x] = \text{size}[\text{left}[x]] + \text{size}[\text{right}[x]] + 1$$

CSC611/ Lecture09

Example: Order-Statistic Tree



CSC611/ Lecture09


OS-SELECT

Goal:

- Given an order-statistic tree, return a pointer to the node containing the i -th smallest key in the subtree rooted at x

Idea:

- $\text{size}[\text{left}[x]]$ = the number of nodes that are smaller than x
- $\text{rank}'[x] = \text{size}[\text{left}[x]] + 1$ in the subtree rooted at x
- If $i = \text{rank}'[x]$ Done!
- If $i < \text{rank}'[x]$: look left
- If $i > \text{rank}'[x]$: look right



CSC611/ Lecture09

OS-SELECT(x, i)

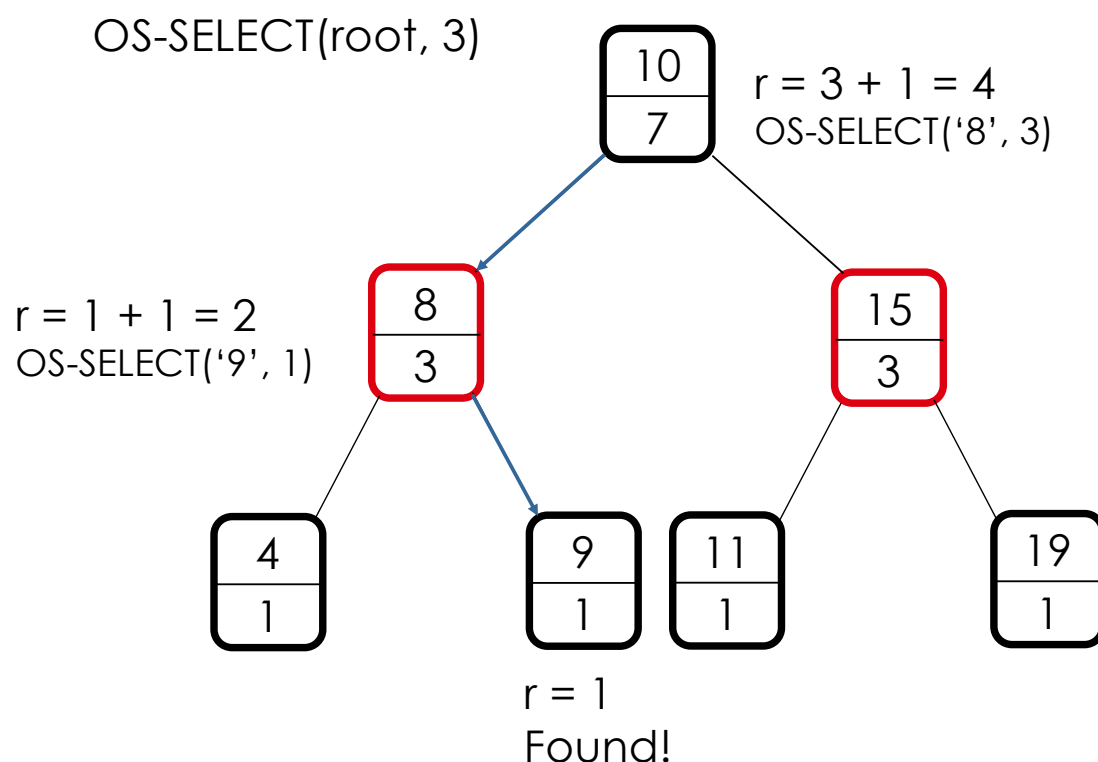
1. $r \leftarrow \text{size}[\text{left}[x]] + 1$ ► compute the rank of x within the subtree rooted at x
2. **if** $i = r$
3. **then return** x
4. **elseif** $i < r$
5. **then return** OS-SELECT($\text{left}[x], i$)
6. **else return** OS-SELECT($\text{right}[x], i - r$)

Initial call: OS-SELECT($\text{root}[T], i$)

Running time: $O(\lg n)$

CSC611/ Lecture09

Example: OS-SELECT



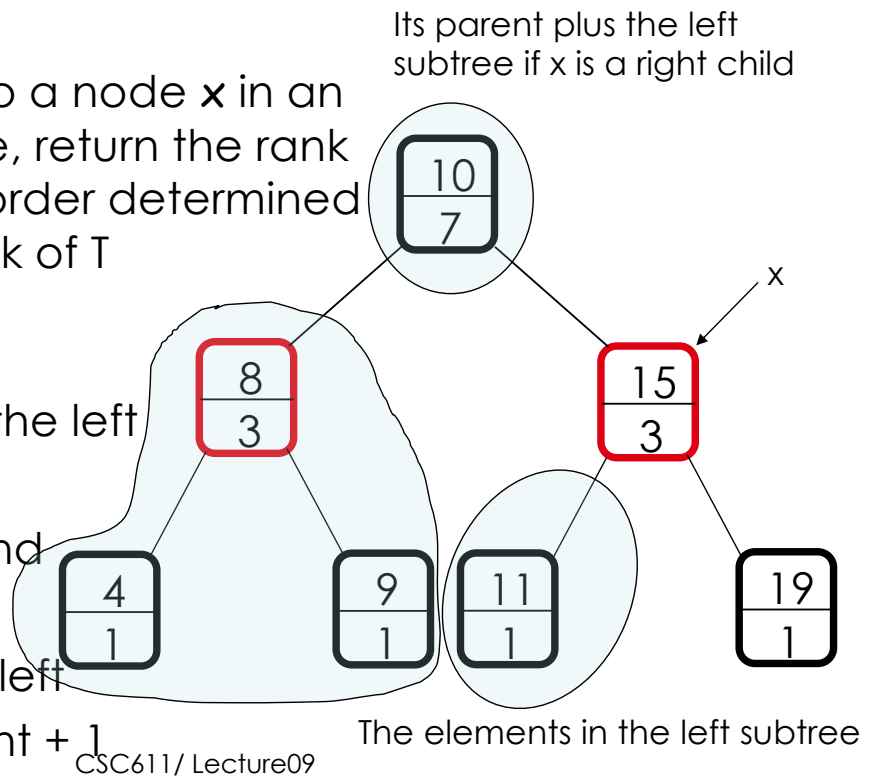
OS-RANK

Goal:

- Given a pointer to a node x in an order-statistic tree, return the rank of x in the linear order determined by an inorder walk of T

Idea:

- Add elements in the left subtree
- Go up the tree and if a right child: add the elements in the left subtree of the parent + 1



OS-RANK(T, x)

1. $r \leftarrow \text{size}[\text{left}[x]] + 1$

Add to the rank the elements in its left subtree + 1 for itself

2. $y \leftarrow x$

Set y as a pointer that will traverse the tree

3. **while** $y \neq \text{root}[T]$

4. **do if** $y = \text{right}[p[y]]$

5. **then** $r \leftarrow r + \text{size}[\text{left}[p[y]]] + 1$

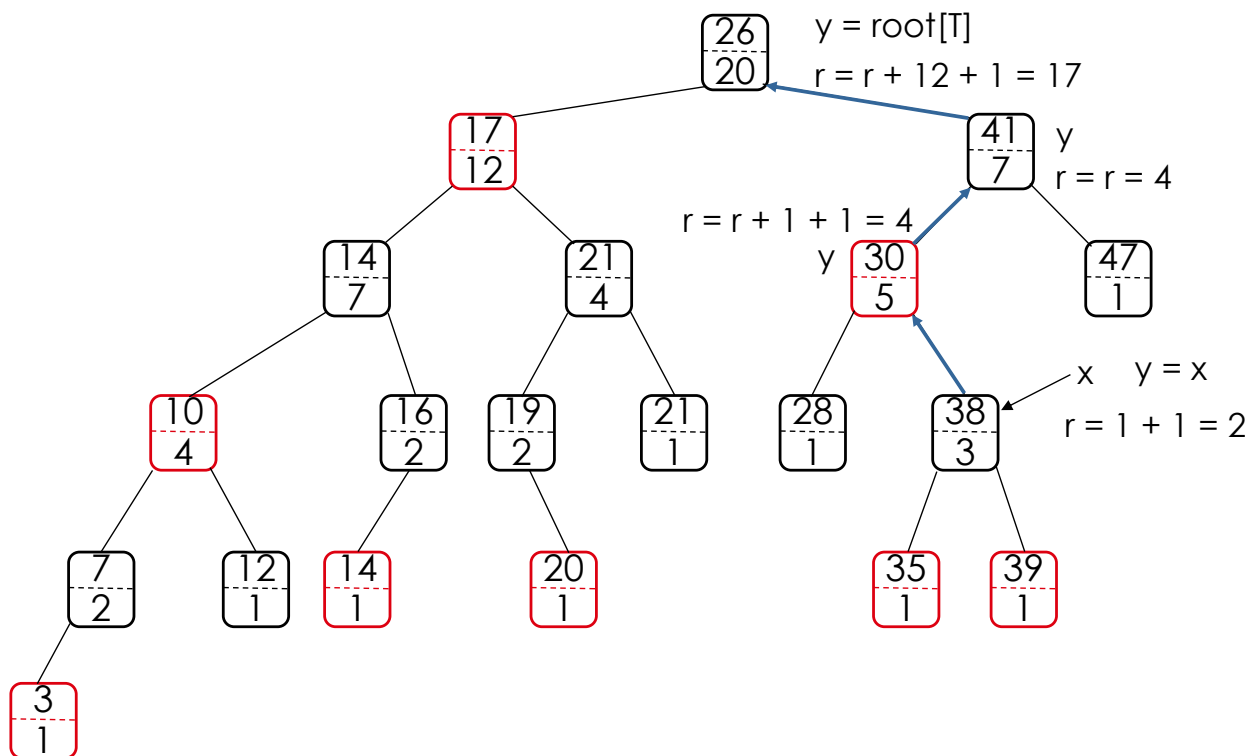
6. $y \leftarrow p[y]$

7. **return** r

If a right child add the size of the parent's left subtree + 1 for the parent

Running time: $O(\lg n)$

Example: OS-RANK



CSC611/ Lecture09

Maintaining Subtree Sizes

- We need to maintain the **size** field during INSERT and DELETE operations
- Need to maintain them efficiently
- Otherwise, might have to recompute all **size** fields, at a cost of $\Omega(n)$

CSC611/ Lecture09

Maintaining *Size* for OS-INSERT

- Insert in a red-black tree has two stages
 1. Perform a binary-search tree insert
 2. Perform rotations and change node colors to restore red-black tree properties


CSC611/ Lecture09

OS-INSERT

Idea for maintaining the **size** field during insert

Phase 1 (going down):

- Increment $\text{size}[x]$ for each node x on the traversed path from the root to the leaves
- The new node gets a size of 1
- Constant work at each node, so still $O(\lg n)$

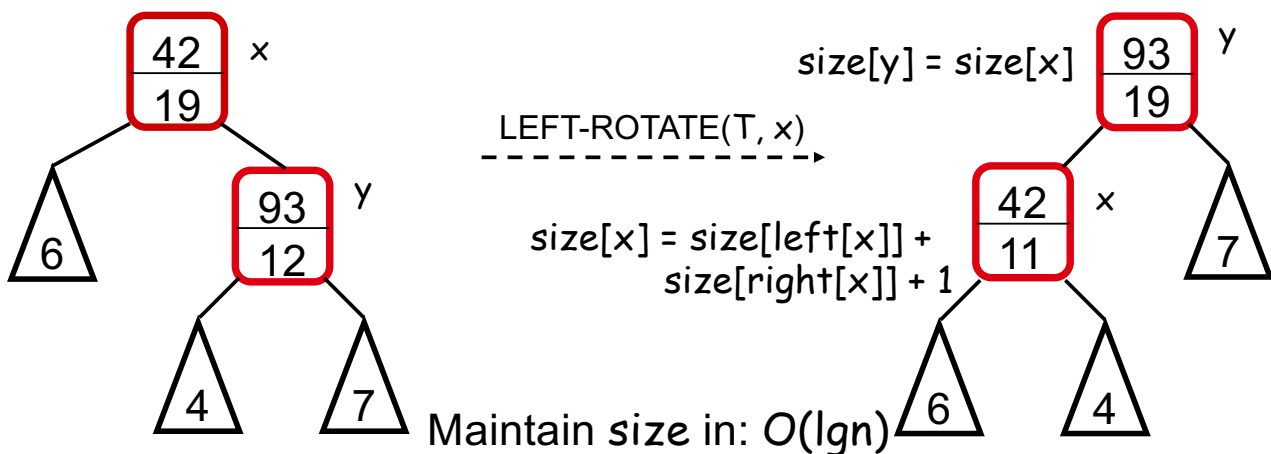


CSC611/ Lecture09

OS-INSERT

Idea for maintaining the **size** field during insert
Phase 2 (going up):

- During RB-INSERT-FIXUP there are:
 - $O(\lg n)$ changes in node colors
 - At most two rotations **Rotations affect the subtree sizes !!**



Augmenting a Data Structure

1. Choose an underlying data structure
⇒ Red-black trees
2. Determine additional information to maintain
⇒ $\text{size}[x]$
3. Verify that we can maintain additional information for existing data structure operations
⇒ Shown how to maintain size during modifying operations
4. Develop new operations
⇒ Developed OS-RANK and OS-SELECT

Augmenting Red-Black Trees

Theorem: Let f be a field that augments a red-black tree. If the contents of f for a node can be computed using only the information in x , $\text{left}[x]$, $\text{right}[x]$ \Rightarrow we can maintain the values of f in all nodes during insertion and deletion, without affecting their $O(\lg n)$ running time.

CSC611/ Lecture09

Examples

1. Can we augment a RBT with $\text{size}[x]$?

Yes: $\text{size}[x] = \text{size}[\text{left}[x]] + \text{size}[\text{right}[x]] + 1$

2. Can we augment a RBT with $\text{height}[x]$?

Yes: $\text{height}[x] = 1 + \max(\text{height}[\text{left}[x]], \text{height}[\text{right}[x]])$

3. Can we augment a RBT with $\text{rank}[x]$?

No, inserting a new minimum will cause all n rank values to change

CSC611/ Lecture09

Interval Trees

Def.: **Interval tree** = a red-black tree that maintains a dynamic set of elements, each element x having associated an interval $\text{int}[x]$.

- Operations on interval trees:
 - $\text{INTERVAL-INSERT}(T, x)$
 - $\text{INTERVAL-DELETE}(T, x)$
 - $\text{INTERVAL-SEARCH}(T, i)$

CSC611/ Lecture09

Interval Properties

- Intervals i and j overlap iff:
 $\text{low}[i] \leq \text{high}[j]$ and $\text{low}[j] \leq \text{high}[i]$



- Intervals i and j do not overlap iff:
 $\text{high}[i] < \text{low}[j]$ or $\text{high}[j] < \text{low}[i]$



CSC611/ Lecture09

Interval Trichotomy

- Any two intervals i and j satisfy the **interval trichotomy**: exactly one of the following three properties holds:
 - i and j overlap,
 - i is to the left of j ($\text{high}[i] < \text{low}[j]$)
 - i is to the right of j ($\text{high}[j] < \text{low}[i]$)

CSC611/ Lecture09

Designing Interval Trees

- Underlying data structure
 - Red-black trees
 - Each node x contains: an interval $\text{int}[x]$, and the key: $\text{low}[\text{int}[x]]$
 - An inorder tree walk will list intervals sorted by their low endpoint


- Additional information

- $\text{max}[x]$ = maximum endpoint value in subtree rooted at x

- Maintaining the information

$$\text{max}[x] = \max \begin{cases} \text{high}[\text{int}[x]] \\ \text{max}[\text{left}[x]] \\ \text{max}[\text{right}[x]] \end{cases}$$

Constant work at each node, so still $O(\lg n)$ time



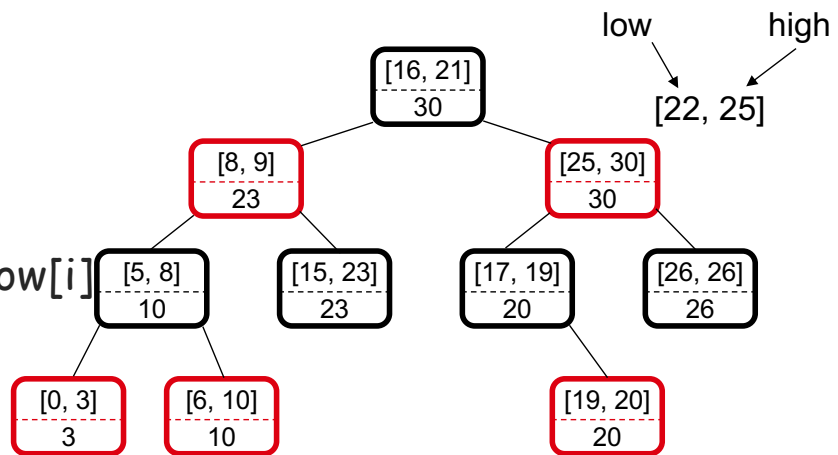
CSC611/ Lecture09

Designing Interval Trees

4. Develop new operations

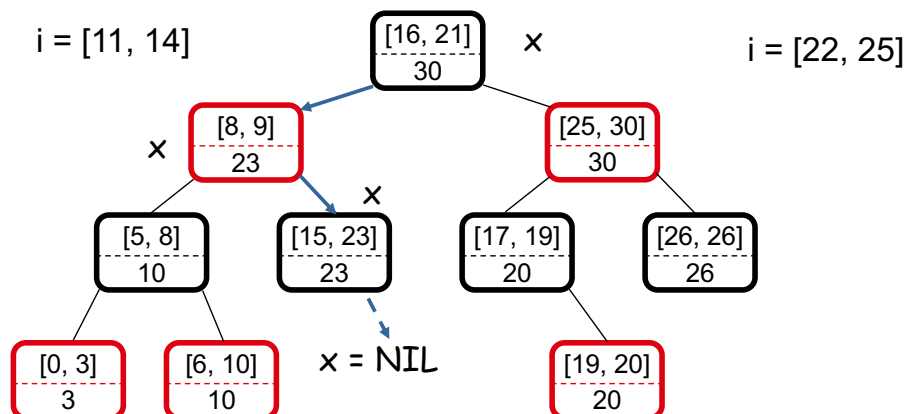
- **INTERVAL-SEARCH(T, i):**
 - Returns a pointer to an element x in the interval tree T , such that $\text{int}[x]$ overlaps with i , or **NIL** otherwise

- Idea:
- Check if $\text{int}[x]$ overlaps with i
- **$\text{Max}[\text{left}[x]] \geq \text{low}[i]$**
 - Go left
- Otherwise, go right



CSC611/ Lecture09

Example



CSC611/ Lecture09

INTERVAL-SEARCH(T, i)

1. $x \leftarrow \text{root}[T]$
2. **while** $x \neq \text{nil}[T]$ and i does not overlap $\text{int}[x]$
3. **do if** $\text{left}[x] \neq \text{nil}[T]$ and
 $\text{max}[\text{left}[x]] \geq \text{low}[i]$
4. **then** $x \leftarrow \text{left}[x]$
5. **else** $x \leftarrow \text{right}[x]$
6. **return** x

CSC611/ Lecture09

Theorem

At the execution of interval search: if the search goes right, then either:

- There is an overlap in right subtree, or
- There is no overlap in either subtree
- Similar when the search goes left
- **It is safe to always proceed in only one direction**


CSC611/ Lecture09

Theorem

- **Proof:** If search goes right:
 - If there is an overlap in right subtree, done
 - If there is no overlap in right \Rightarrow show there is no overlap in left
 - Went right because:

$\text{left}[x] = \text{nil}[T] \Rightarrow$ no overlap in left, or

$\text{max}[\text{left}[x]] < \text{low}[i] \Rightarrow$ no overlap in left



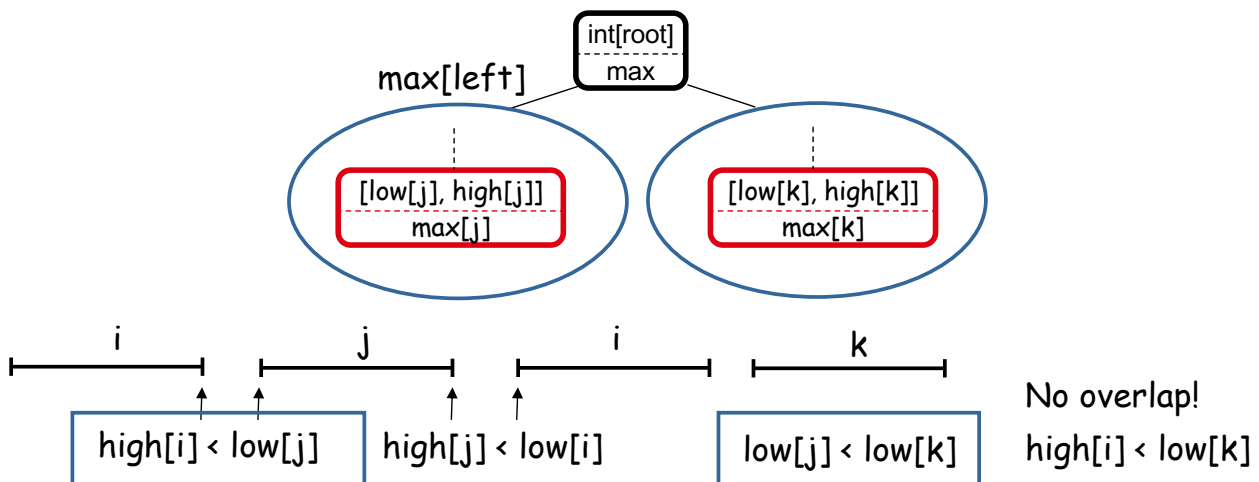
CSC611/ Lecture09

Theorem - Proof

If search goes left:

- If there is an overlap in left subtree, done
- If there is no overlap in left, show there is no overlap in right
- Went left because:

$\text{low}[i] \leq \text{max}[\text{left}[x]] = \text{high}[j]$ for some j in left subtree



CSC611/ Lecture09