

CSC 631: High-Performance Computer Architecture

Spring 2017

Lecture 2: Instruction Set Architectures


Analog Computers

- Analog computer represents problem variables as some physical quantity (e.g., mechanical displacement, voltage on a capacitor) and uses scaled physical behavior to calculate results



[Marsyas, Creative Commons BY-SA 3.0]

Antikythera mechanism c.100BC




[BenFrantzDale, Creative Commons BY-SA 3.0]

Wingtip vortices off Cessna tail in wind tunnel

Digital Computers

- Represent problem variables as numbers encoded using discrete steps
 - Discrete steps provide noise immunity
- Enables accurate and deterministic calculations
 - Same inputs give same outputs exactly
- Not constrained by physically realizable functions
- Programmable digital computers are the focus of computer architectures

Charles Babbage (1791-1871)



*[Copyright expired and in public domain.
Image obtained from Wikimedia Commons.]*

- Lucasian Professor of Mathematics, Cambridge University, 1828-1839
- A true “polymath” with interests in many areas
- Frustrated by errors in printed tables, wanted to build machines to evaluate and print accurate tables
- Inspired by earlier work organizing human “computers” to methodically calculate tables by hand

Difference Engine 1822

- Continuous functions can be approximated by polynomials, which can be computed from difference tables:


$$\begin{aligned}f(n) &= n^2 + n + 41 \\d1(n) &= f(n) - f(n-1) = 2n \\d2(n) &= d1(n) - d1(n-1) = 2\end{aligned}$$

- Can calculate using only a single adder:

n	0	1	2	3	4
d2(n)			2	2	2
d1(n)		2	4	6	8
f(n)	41	43	47	53	61

Realizing the Difference Engine

- Mechanical calculator, hand-cranked, using decimal digits
- Babbage did not complete the DE, moving on to the Analytical Engine (but used ideas from AE in improved DE 2 plan)
- Scheutz completed working version in 1855, sold copy to British Government



[Geni, Creative Commons BY-SA 3.0]

- Modern day recreation of DE2, including printer, showed entire design possible using original technology
 - first at British Science Museum
 - copy at Computer History Museum in San Jose

Analytical Engine 1837

- Recognized as first general-purpose digital computer
 - Many iterations of the design (multiple Analytical Engines)
- Contains the major components of modern computers:
 - “Store”: Main memory where numbers and intermediate results were held (1,000 decimal words, 40-digits each)
 - “Mill”: Arithmetic unit where processing was performed including addition, multiplication, and division
 - Also supported conditional branching and looping, and exceptions on overflow (machine jams and bell rings)
 - Had a form of microcode (the “Barrel”)
- Program, input and output data on punched cards
- Instruction cards hold opcode and address of operands in store
 - 3-address format with two sources and one destination, all in store
- Branches implemented by mechanically changing order cards were inserted into machine
- Only small pieces were ever built

Analytical Engine Design Choices

- Decimal, because storage on mechanical gears
 - Babbage considered binary and other bases, but no clear advantage over human-friendly decimal
- 40-digit precision (equivalent to >133 bits)
 - To reduce impact of scaling given lack of floating-point hardware
- Used “locking” or mechanical amplification to overcome noise in transferring mechanical motion around machine
 - Similar to non-linear gain in digital electronic circuits
- Had a fast “anticipating” carry
 - Mechanical version of pass-transistor carry propagate used in CMOS adders (and earlier in relay adders)

Ada Lovelace (1815-1852)



*[By Margaret Sarah Carpenter,
Copyright expired and in public domain]*

- Translated lectures of Luigi Menabrea who published notes of Babbage's lectures in Italy
- Lovelace considerably embellished notes and described Analytical Engine program to calculate Bernoulli numbers that would have worked if AE was built
 - The first program!
- Imagined many uses of computers beyond calculations of tables
- Was interested in modeling the brain

Early Programmable Calculators

- Analog computing was popular in first half of 20th century as digital computing was too expensive
- But during late 30s and 40s, several programmable digital calculators were built (date when operational)
 - Atanasoff Linear Equation Solver (1939)
 - Zuse Z3 (1941)
 - Harvard Mark I (1944)
 - ENIAC (1946)

Atanasoff-Berry Linear Equation Solver (1939)

- Fixed-function calculator for solving up to 29 simultaneous linear equations
- Digital binary arithmetic (50-bit fixed-point words)
- Dynamic memory (rotating drum of capacitors)
- Vacuum tube logic for processing




[Manop, Creative Commons BY-SA 3.0]

In 1973, Atanasoff was credited as inventor of “automatic electronic digital computer” after patent dispute with Eckert and Mauchly (ENIAC)

Zuse Z3 (1941)


- Built by Konrad Zuse in wartime Germany using 2000 relays
- Had normalized floating-point arithmetic with hardware handling of exceptional values (+/- infinity, undefined)
 - 1-bit sign, 7-bit exponent, 14-bit significand
- 64 words of memory
- Two-stage pipeline 1) fetch&execute 2) writeback
- No conditional branch
- Programmed via paper tape



Replica of the Zuse Z3 in the Deutsches Museum, Munich

[Venusianer, Creative Commons BY-SA 3.0]

Harvard Mark I (1944)


- 
- Proposed by Howard Aiken at Harvard, and funded and built by IBM
 - Mostly mechanical with some electrically controlled relays and gears
 - Weighed 5 tons and had 750,000 components
 - Stored 72 numbers each of 23 decimal digits
 - Speed: adds 0.3s, multiplies 6s, divide 15s, trig >1 minute
 - Instructions on paper tape (2-address format)
 - Could run long programs automatically
 - Loops by gluing paper tape into loops
 - No conditional branch
 - Although mentioned Babbage in proposal, was more limited than analytical engine

[Waldir, Creative Commons BY-SA 3.0]

ENIAC (1946)

- First electronic general-purpose computer
- Construction started in secret at UPenn Moore School of Electrical Engineering during WWII to calculate firing tables for US Army, designed by Eckert and Mauchly
- 17,468 vacuum tubes
- Weighed 30 tons, occupied 1800 sq ft, power 150kW
- Twelve 10-decimal-digit accumulators
- Had a conditional branch!
- Programmed by plugboard and switches, time consuming!
- Purely electronic instruction fetch and execution, so fast
 - 10-digit x 10-digit multiply in 2.8ms (2000x faster than Mark-1)
- As a result of speed, it was almost entirely I/O bound
- As a result of large number of tubes, it was often broken (5 days was longest time between failures)

ENIAC




[Public Domain, US Army Photo]

EDVAC

- ENIAC team started discussing stored-program concept to speed up programming and simplify machine design
- John von Neumann was consulting at UPenn and typed up ideas in “First Draft of a report on EDVAC”
- Herman Goldstine circulated the draft June 1945 to many institutions, igniting interest in the stored-program idea
 - But also, ruined chances of patenting it
 - Report falsely gave sole credit to von Neumann for the ideas
 - Maurice Wilkes was excited by report and decided to come to US workshop on building computers
- Later, in 1948, modifications to ENIAC allowed it to run in stored-program mode, but 6x slower than hardwired
 - Due to I/O limitations, this speed drop was not practically significant and improvement in productivity made it worthwhile
- EDVAC eventually built and (mostly) working in 1951
 - Delayed by patent disputes with university

Manchester SSEM “Baby” (1948)

- Manchester University group build small-scale experimental machine to demonstrate idea of using cathode-ray tubes (CRTs) for computer memory instead of mercury delay lines
- *Williams-Kilburn Tubes were first random access electronic storage devices*
- 32 words of 32-bits, accumulator, and program counter
- Machine ran world’s first stored-program in June 1948
- Led to later Manchester Mark-1 full-scale machine
 - Mark-1 introduced *index* registers
 - Mark-1 commercialized by Ferranti



Cambridge EDSAC (1949)

- Maurice Wilkes came back from workshop in US and set about building a stored-program computer in Cambridge
- EDSAC used mercury-delay line storage to hold up to 1024 words (512 initially) of 17 bits (+1 bit of padding in delay line)
- Two’s-complement binary arithmetic
- Accumulator ISA with self-modifying code for indexing
- David Wheeler, who earned the world’s first computer science PhD, invented the subroutine (“Wheeler jump”) for this machine
 - Users built a large library of useful subroutines
- UK’s first commercial computer, LEO-I (Lyons Electronic Office), was based on EDSAC, ran business software in 1951
 - Software for LEO was still running in the 1980s in emulation on ICL mainframes!
- EDSAC-II (1958) was first machine with microprogrammed control unit

Commercial computers: BINAC (1949) and UNIVAC (1951)


- Eckert and Mauchly left U.Penn after patent rights disputes and formed the Eckert-Mauchly Computer Corporation
- World's first commercial computer was BINAC with two CPUs that checked each other
 - BINAC apparently never worked after shipment to first (only) customer
- Second commercial computer was UNIVAC
 - Used mercury delay-line memory, 1000 words of 12 alpha characters
 - Famously used to predict presidential election in 1952
 - Eventually 46 units sold at >\$1M each
 - Often, mistakingly called the IBM UNIVAC

IBM 701 (1952)

- IBM's first commercial scientific computer
- Main memory was 72 Williams's Tubes, each 1Kib, for total of 2048 words of 36 bits each
 - Memory cycle time of 12μs
- Accumulator ISA with multiplier/quotient register
- 18-bit/36-bit numbers in sign-magnitude fixed-point
- Misquote from Thomas Watson Sr/Jr:
 - “I think there is a world market for maybe five computers”***
- Actually TWJr said at shareholder meeting:
 - “as a result of our trip [selling the 701], on which we expected to get orders for five machines, we came home with orders for 18.”***


IBM 650 (1953)

- The first mass-produced computer
- Low-end system with drum-based storage and digit serial ALU
- Almost 2,000 produced



[Cushing Memorial Library and Archives, Texas A&M, Creative Commons Attribution 2.0 Generic]

IBM 650 Architecture



IBM 650 Instruction Set

- Address and data in 10-digit decimal words
- Instructions encode:
 - Two-digit opcode encoded 44 instructions in base instruction set, expandable to 97 instructions with options
 - Four-digit data address
 - Four-digit next instruction address
 - Programmer's arrange code to minimize drum latency!
- Special instructions added to compare value to all words on track

Early Instruction Sets


- Very simple ISAs, mostly single-address accumulator-style machines, as high-speed circuitry was expensive
 - Based on earlier “calculator” model
- Over time, appreciation of software needs shaped ISA
- Index registers (Kilburn, Mark-1) added to avoid need for self-modifying code to step through array
- Over time, more index registers were added
- And more operations on the index registers
- Eventually, just provide general-purpose registers (GPRs) and orthogonal instruction sets
- But some other options explored...

Burrough's B5000 Stack Architecture: Robert Barton, 1960

- Hide instruction set completely from programmer using high-level language (ALGOL)
- Use stack architecture to simplify compilation, expression evaluation, recursive subroutine calls, interrupt handling,...


Evaluation of Expressions

$(a + b * c) / (a + d * c - e)$



Reverse Polish


$a \ b \ c \ * \ + \ a \ d \ c \ * \ + \ e \ - \ /$
 ↑ ↑ ↑ ↑
 push push multiply



Evaluation Stack


Evaluation of Expressions

$(a + b * c) / (a + d * c - e)$



Reverse Polish

$a \ b \ c \ * \ + \ a \ d \ c \ * \ + \ e \ - \ /$
 ↑
 add



27

IBM's Big Bet: 360 Architecture

- By early 1960s, IBM had several incompatible families of computer:

701 → 7094

650 → 7074

702 → 7080

1401 → 7010

- Each system had its own
 - Instruction set
 - I/O system and secondary storage (magnetic tapes, drums and disks)
 - assemblers, compilers, libraries,...
 - market niche (business, scientific, real time, ...)

IBM 360 : Design Premises

Amdahl, Blaauw and Brooks, 1964

- The design must lend itself to growth and successor machines
- General method for connecting I/O devices
- Total performance - answers per month rather than bits per microsecond \Rightarrow programming aids
- Machine must be capable of supervising itself without manual intervention
- Built-in hardware fault checking and locating aids to reduce down time
- Simple to assemble systems with redundant I/O devices, memories etc. for fault tolerance
- Some problems required floating-point larger than 36 bits

Stack versus GPR Organization

Amdahl, Blaauw and Brooks, 1964

1. The performance advantage of push-down stack organization is derived from the presence of fast registers and not the way they are used.
2. "Surfacing" of data in stack which are "profitable" is approximately 50% because of constants and common subexpressions.
3. Advantage of instruction density because of implicit addresses is equaled if short addresses to specify registers are allowed.
4. Management of finite-depth stack causes complexity.
5. Recursive subroutine advantage can be realized only with the help of an independent stack for addressing.
6. Fitting variable-length fields into fixed-width word is awkward.

IBM 360: A General-Purpose Register (GPR) Machine

- Processor State
 - 16 General-Purpose 32-bit Registers
 - may be used as index and base register
 - Register 0 has some special properties
 - 4 Floating Point 64-bit Registers
 - A Program Status Word (PSW)
 - PC, Condition codes, Control flags
- A 32-bit machine with 24-bit addresses
 - But no instruction contains a 24-bit address!
- Data Formats
 - 8-bit bytes, 16-bit half-words, 32-bit words, 64-bit double-words

 *The IBM 360 is why bytes are 8-bits long today!*

IBM 360: Initial Implementations

	<i>Model 30</i>	<i>...</i>	<i>Model 70</i>
<i>Storage</i>	8K - 64 KB		256K - 512 KB
<i>Datapath</i>	8-bit		64-bit
<i>Circuit Delay</i>	30 nsec/level		5 nsec/level
<i>Local Store</i>	Main Store		Transistor Registers
<i>Control Store</i>	Read only 1 μ sec		Conventional circuits

IBM 360 instruction set architecture (ISA) completely hid the underlying technological differences between various models.

Milestone: The first true ISA designed as portable hardware-software interface!

With minor modifications it still survives today!

Acknowledgements

- These course notes were developed by:
 - Krste Asanovic (UCB)
 - Arvind (MIT)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)