

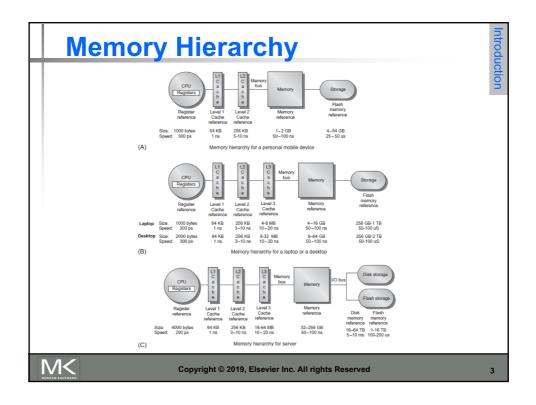
Introduction

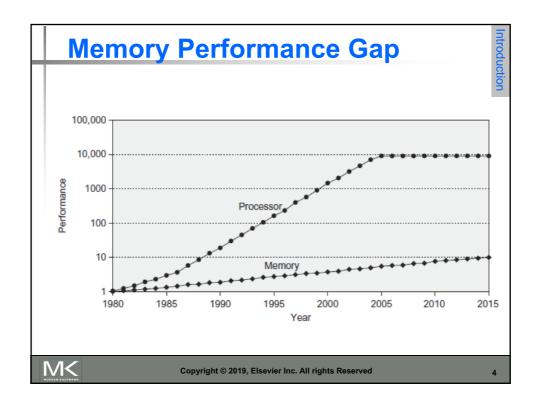
 Programmers want unlimited amounts of memory with low latency

- Fast memory technology is more expensive per bit than slower memory
- Solution: organize memory system into a hierarchy
 - Entire addressable memory space available in largest, slowest memory
 - Incrementally smaller and faster memories, each containing a subset of the memory below it, proceed in steps up toward the processor
- Temporal and spatial locality insures that nearly all references can be found in smaller memories
 - Gives the allusion of a large, fast memory being presented to the processor

M<

Copyright © 2019, Elsevier Inc. All rights Reserved





Memory Hierarchy Design

- Memory hierarchy design becomes more crucial with recent multi-core processors:
 - Aggregate peak bandwidth grows with # cores:
 - Intel Core i7 can generate two references per core per clock
 - Four cores and 3.2 GHz clock
 - 25.6 billion 64-bit data references/second + 12.8 billion 128-bit instruction references/second = 409.6 GB/s!
 - DRAM bandwidth is only 8% of this (34.1 GB/s)
 - Requires:
 - Multi-port, pipelined caches
 - Two levels of cache per core
 - Shared third-level cache on chip

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Performance and Power

High-end microprocessors have >10 MB on-chip cache

Consumes large amount of area and power budget

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Memory Hierarchy Basics

When a word is not found in the cache, a miss occurs:

- Fetch word from lower level in hierarchy, requiring a higher latency reference
- Lower level may be another cache or the main memory
- Also fetch the other words contained within the block
 - Takes advantage of spatial locality
- Place block into cache in any location within its set, determined by address
 - block address MOD number of sets in cache

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Memory Hierarchy Basics

n sets => n-way set associative

- Direct-mapped cache => one block per set
- Fully associative => one set
- Writing to cache: two strategies
 - Write-through
 - Immediately update lower levels of hierarchy
 - Write-back
 - Only update lower levels of hierarchy when an updated block is replaced
 - Both strategies use write buffer to make writes asynchronous

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

3

Chapter 2 — Instructions: Language of the Computer

Memory Hierarchy Basics

- Miss rate
 - Fraction of cache access that result in a miss
- Causes of misses
 - Compulsory
 - First reference to a block
 - Capacity
 - Blocks discarded and later retrieved
 - Conflict
 - Program makes repeated references to multiple addresses from different blocks that map to the same location in the cache

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Memory Hierarchy Basics | Misses | Miss rate × Memory accesses | Miss rate × Memory accesses | Instruction | | Average memory access time = Hit time + Miss rate × Miss penalty | Speculative and multithreaded processors may execute other instructions during a miss | Reduces performance impact of misses

Memory Hierarchy Basics

- Six basic cache optimizations:
 - Larger block size
 - Reduces compulsory misses
 - Increases capacity and conflict misses, increases miss penalty
 - Larger total cache capacity to reduce miss rate
 - Increases hit time, increases power consumption
 - Higher associativity
 - Reduces conflict misses
 - Increases hit time, increases power consumption
 - Higher number of cache levels
 - Reduces overall memory access time
 - Giving priority to read misses over writes
 - Reduces miss penalty
 - Avoiding address translation in cache indexing
 - Reduces hit time

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

11

Memory Technology and Optimizations

- Performance metrics
 - Latency is concern of cache
 - Bandwidth is concern of multiprocessors and I/O
 - Access time
 - Time between read request and when desired word arrives
 - Cycle time
 - Minimum time between unrelated requests to memory
- SRAM memory has low latency, use for cache
- Organize DRAM chips into many banks for high bandwidth, use for main memory

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

12

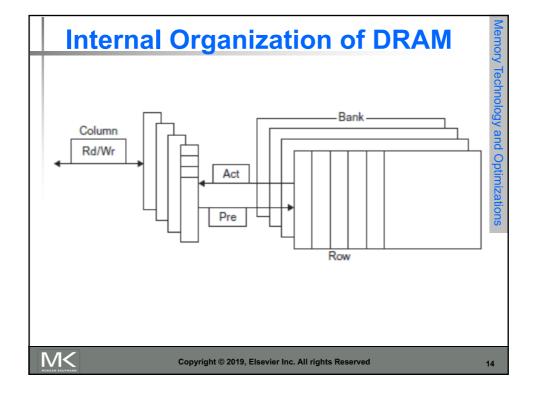
y reciliology and optimization

Memory Technology

- SRAM
 - Requires low power to retain bit
 - Requires 6 transistors/bit
- DRAM
 - Must be re-written after being read
 - Must also be periodically refreshed
 - Every ~ 8 ms (roughly 5% of time)
 - Each row can be refreshed simultaneously
 - One transistor/bit
 - Address lines are multiplexed:
 - Upper half of address: row access strobe (RAS)
 - Lower half of address: column access strobe (CAS)

M<

Copyright © 2019, Elsevier Inc. All rights Reserved



Memory Technology

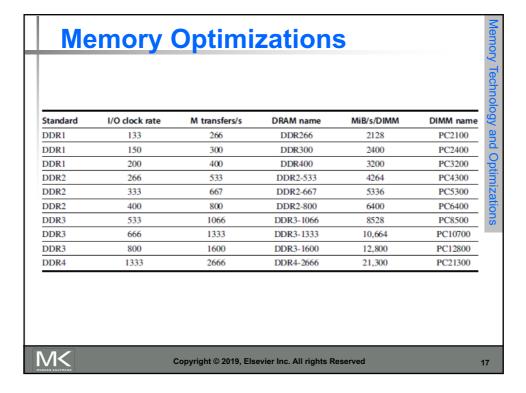
- Amdahl:
 - Memory capacity should grow linearly with processor speed
 - Unfortunately, memory capacity and speed has not kept pace with processors
- Some optimizations:
 - Multiple accesses to same row
 - Synchronous DRAM
 - Added clock to DRAM interface
 - Burst mode with critical word first
 - Wider interfaces
 - Double data rate (DDR)
 - Multiple banks on each DRAM device

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

15

Memory Optimizations Best case access time (no precharge) Precharge needed RAS time (ns) CAS time (ns) Total (ns) Production year Chip size DRAM type Total (ns) 2000 256M bit DDR1 63 512M bit DDR1 30 45 2002 15 15 DDR2 30 45 2004 1G bit 15 15 2006 2G bit DDR2 10 10 20 30 2010 4G bit DDR3 13 13 26 39 DDR4 2016 8G bit 13 26 39 13 M< Copyright © 2019, Elsevier Inc. All rights Reserved 16



Memory Optimizations

- DDR:
 - DDR2
 - Lower power (2.5 V -> 1.8 V)
 - Higher clock rates (266 MHz, 333 MHz, 400 MHz)
 - DDR3
 - 1.5 V
 - 800 MHz
 - DDR4
 - 1-1.2 V
 - 1333 MHz
- GDDR5 is graphics memory based on DDR3

M<

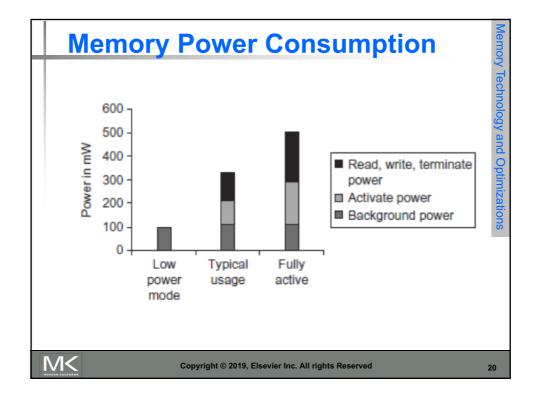
Copyright © 2019, Elsevier Inc. All rights Reserved

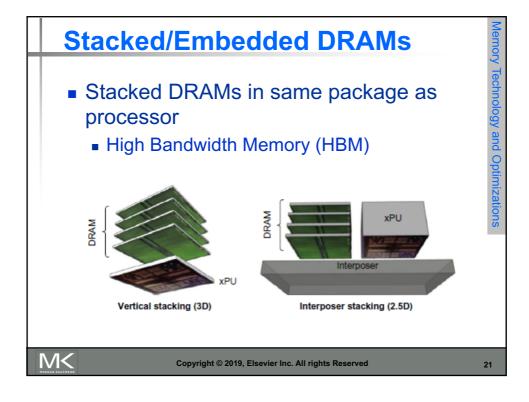
Memory Optimizations

- Reducing power in SDRAMs:
 - Lower voltage
 - Low power mode (ignores clock, continues to refresh)
- Graphics memory:
 - Achieve 2-5 X bandwidth per DRAM vs. DDR3
 - Wider interfaces (32 vs. 16 bit)
 - Higher clock rate
 - Possible because they are attached via soldering instead of socketted DIMM modules

M<

Copyright © 2019, Elsevier Inc. All rights Reserved





Flash Memory

- Type of EEPROM
- Types: NAND (denser) and NOR (faster)
- NAND Flash:
 - Reads are sequential, reads entire page (.5 to 4 KiB)
 - 25 us for first byte, 40 MiB/s for subsequent bytes
 - SDRAM: 40 ns for first byte, 4.8 GB/s for subsequent bytes
 - 2 KiB transfer: 75 uS vs 500 ns for SDRAM, 150X slower
 - 300 to 500X faster than magnetic disk

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

NAND Flash Memory

- Must be erased (in blocks) before being overwritten
- Nonvolatile, can use as little as zero power
- Limited number of write cycles (~100,000)
- \$2/GiB, compared to \$20-40/GiB for SDRAM and \$0.09 GiB for magnetic disk
- Phase-Change/Memrister Memory
 - Possibly 10X improvement in write performance and 2X improvement in read performance

Copyright © 2019, Elsevier Inc. All rights Reserved

Memory Dependability

- Memory is susceptible to cosmic rays
- Soft errors: dynamic errors
 - Detected and fixed by error correcting codes (ECC)
- Hard errors: permanent errors
 - Use spare rows to replace defective rows
- Chipkill: a RAID-like error recovery technique

M<

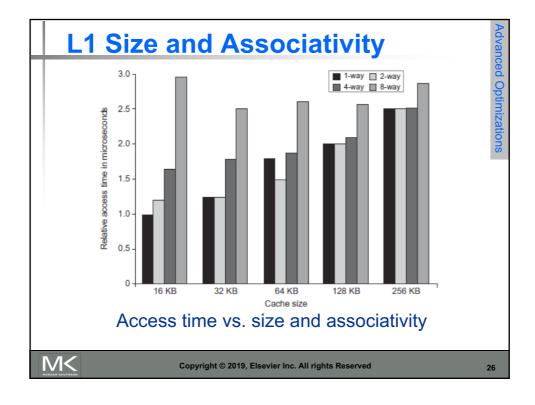
Copyright © 2019, Elsevier Inc. All rights Reserved

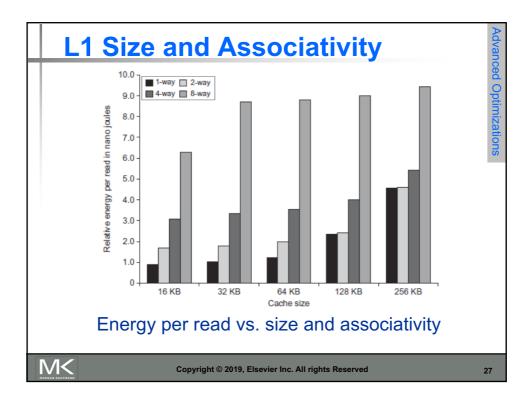
Advanced Optimizations

- Reduce hit time
 - Small and simple first-level caches
 - Way prediction
- Increase bandwidth
 - Pipelined caches, multibanked caches, non-blocking caches
- Reduce miss penalty
 - Critical word first, merging write buffers
- Reduce miss rate
 - Compiler optimizations
- Reduce miss penalty or miss rate via parallelization
 - Hardware or compiler prefetching

M<

Copyright © 2019, Elsevier Inc. All rights Reserved





Way Prediction

- To improve hit time, predict the way to pre-set mux
 - Mis-prediction gives longer hit time
 - Prediction accuracy
 - > 90% for two-way
 - > 80% for four-way
 - I-cache has better accuracy than D-cache
 - First used on MIPS R10000 in mid-90s
 - Used on ARM Cortex-A8
- Extend to predict block as well
 - "Way selection"
 - Increases mis-prediction penalty

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Pipelined Caches

- Pipeline cache access to improve bandwidth
 - Examples:
 - Pentium: 1 cycle
 - Pentium Pro Pentium III: 2 cycles
 - Pentium 4 Core i7: 4 cycles
- Increases branch mis-prediction penalty
- Makes it easier to increase associativity

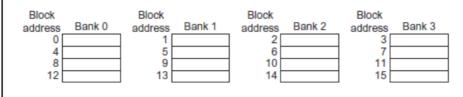
M<

Copyright © 2019, Elsevier Inc. All rights Reserved

29

Multibanked Caches

- Organize cache as independent banks to support simultaneous access
 - ARM Cortex-A8 supports 1-4 banks for L2
 - Intel i7 supports 4 banks for L1 and 8 banks for L2
- Interleave banks according to block address



M<

Copyright © 2019, Elsevier Inc. All rights Reserved

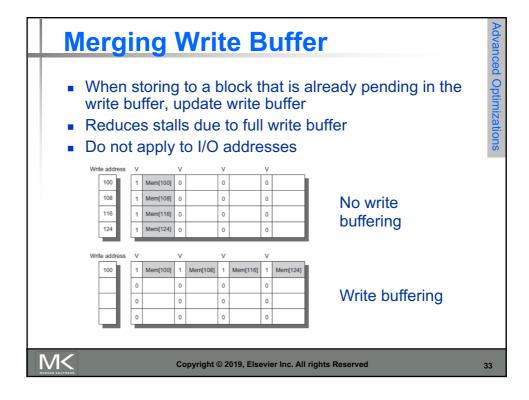
Nonblocking Caches Allow hits before previous misses complete "Hit under miss" "Hit under multiple miss" L2 must support this In general, processors can hide L1 miss penalty but not L2 miss penalty "Hit-under-I-miss — Hit-under-2-misses — Hit-under-64-misses "Hit-under-I-miss — Hit-under-2-misses — Hit-under-64-misses "PHI-under-I-miss — Hit-under-2-misses — Hit

Critical Word First, Early Restart

- Critical word first
 - Request missed word from memory first
 - Send it to the processor as soon as it arrives
- Early restart
 - Request words in normal order
 - Send missed work to the processor as soon as it arrives
- Effectiveness of these strategies depends on block size and likelihood of another access to the portion of the block that has not yet been fetched

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

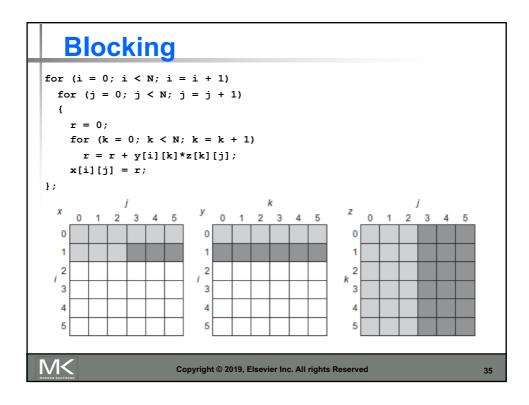


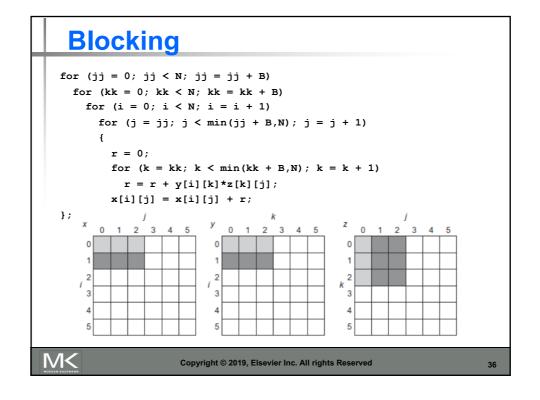
Compiler Optimizations

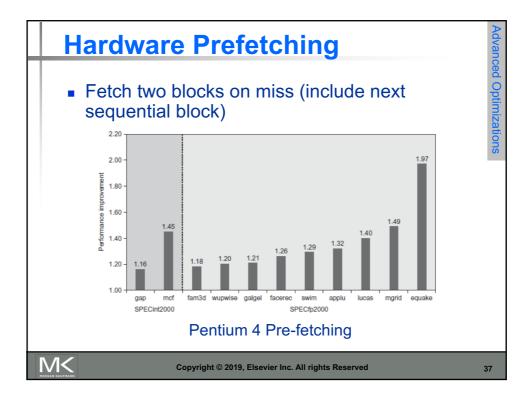
- Loop Interchange
 - Swap nested loops to access memory in sequential order
- Blocking
 - Instead of accessing entire rows or columns, subdivide matrices into blocks
 - Requires more memory accesses but improves locality of accesses

M<

Copyright © 2019, Elsevier Inc. All rights Reserved







Compiler Prefetching

- Insert prefetch instructions before data is needed
- Non-faulting: prefetch doesn't cause exceptions
- Register prefetch
 - Loads data into register
- Cache prefetch
 - Loads data into cache
- Combine with loop unrolling and software pipelining

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Use HBM to Extend Hierarchy

- 128 MiB to 1 GiB
- Smaller blocks require substantial tag storage
- Larger blocks are potentially inefficient
- One approach (L-H):
 - Each SDRAM row is a block index
 - Each row contains set of tags and 29 data segments
 - 29-set associative
 - Hit requires a CAS

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

39

Use HBM to Extend Hierarchy

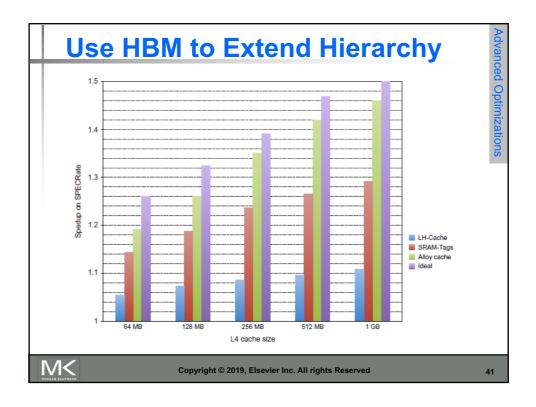
- Another approach (Alloy cache):
 - Mold tag and data together
 - Use direct mapped
- Both schemes require two DRAM accesses for misses
 - Two solutions:
 - Use map to keep track of blocks
 - Predict likely misses

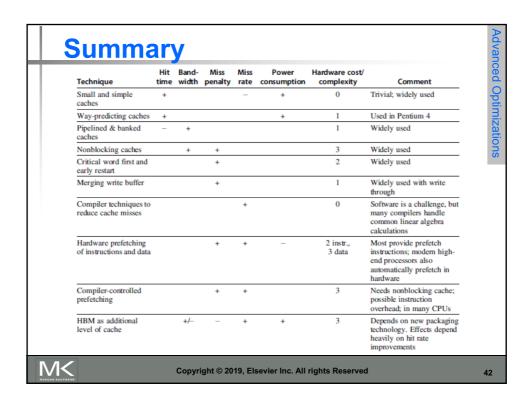
M<

Copyright © 2019, Elsevier Inc. All rights Reserved

40

Advanced Optimizations





Virtual Memory and Virtual Machines

- Protection via virtual memory
 - Keeps processes in their own memory space
- Role of architecture
 - Provide user mode and supervisor mode
 - Protect certain aspects of CPU state
 - Provide mechanisms for switching between user mode and supervisor mode
 - Provide mechanisms to limit memory accesses
 - Provide TLB to translate addresses

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

43

Virtual Machines

- Supports isolation and security
- Sharing a computer among many unrelated users
- Enabled by raw speed of processors, making the overhead more acceptable
- Allows different ISAs and operating systems to be presented to user programs
 - "System Virtual Machines"
 - SVM software is called "virtual machine monitor" or "hypervisor"
 - Individual virtual machines run under the monitor are called "guest VMs"

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

44

Virtual Memory and Virtual Machine

al Melloly and villual Machille

Requirements of VMM

- Guest software should:
 - Behave on as if running on native hardware
 - Not be able to change allocation of real system resources
- VMM should be able to "context switch" guests
- Hardware must allow:
 - System and use processor modes
 - Privileged subset of instructions for allocating system resources

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

45

Impact of VMs on Virtual Memory

- Each guest OS maintains its own set of page tables
 - VMM adds a level of memory between physical and virtual memory called "real memory"
 - VMM maintains shadow page table that maps guest virtual addresses to physical addresses
 - Requires VMM to detect guest's changes to its own page table
 - Occurs naturally if accessing the page table pointer is a privileged operation

M<

Copyright © 2019, Elsevier Inc. All rights Reserved

Extending the ISA for Virtualization

- Objectives:
 - Avoid flushing TLB
 - Use nested page tables instead of shadow page tables
 - Allow devices to use DMA to move data
 - Allow guest OS's to handle device interrupts
 - For security: allow programs to manage encrypted portions of code and data

Copyright © 2019, Elsevier Inc. All rights Reserved

47

Fallacies and Pitfalls

- Predicting cache performance of one program from another
- Simulating enough instructions to get accurate performance measures of the memory hierarchy
- Not deliverying high memory bandwidth in a cache-based system

M<

Copyright © 2019, Elsevier Inc. All rights Reserved