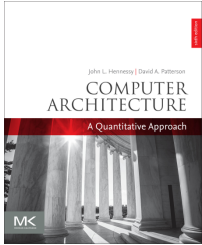

**Computer Architecture**
A Quantitative Approach, Sixth Edition



Chapter 4

Data-Level Parallelism in Vector, SIMD, and GPU Architectures




Copyright © 2019, Elsevier Inc. All rights Reserved

1

Introduction

- SIMD architectures can exploit significant data-level parallelism for:
 - Matrix-oriented scientific computing
 - Media-oriented image and sound processors
- SIMD is more energy efficient than MIMD
 - Only needs to fetch one instruction per data operation
 - Makes SIMD attractive for personal mobile devices
- SIMD allows programmer to continue to think sequentially



Copyright © 2019, Elsevier Inc. All rights Reserved

2

Introduction

SIMD Parallelism

Introduction

- Vector architectures
- SIMD extensions
- Graphics Processor Units (GPUs)
- For x86 processors:
 - Expect two additional cores per chip per year
 - SIMD width to double every four years
 - Potential speedup from SIMD to be twice that from MIMD!



Copyright © 2019, Elsevier Inc. All rights Reserved

3

Vector Architectures

Vector Architectures

- Basic idea:
 - Read sets of data elements into “vector registers”
 - Operate on those registers
 - Disperse the results back into memory
- Registers are controlled by compiler
 - Used to hide memory latency
 - Leverage memory bandwidth



Copyright © 2019, Elsevier Inc. All rights Reserved

4

VMIPS

Vector Architectures

- Example architecture: RV64V
 - Loosely based on Cray-1
 - 32 62-bit vector registers
 - Register file has 16 read ports and 8 write ports
 - Vector functional units
 - Fully pipelined
 - Data and control hazards are detected
 - Vector load-store unit
 - Fully pipelined
 - One word per clock cycle after initial latency
 - Scalar registers
 - 31 general-purpose registers
 - 32 floating-point registers



Copyright © 2019, Elsevier Inc. All rights Reserved

5

VMIPS Instructions

Vector Architectures

- .vv: two vector operands
- .vs and .sv: vector and scalar operands
- LV/SV: vector load and vector store from address
- Example: DAXPY

vsetdcfg	4*FP64	# Enable 4 DP FP vregs
fld	f0,a	# Load scalar a
vld	v0,x5	# Load vector X
vmul	v1,v0,f0	# Vector-scalar mult
vld	v2,x6	# Load vector Y
vadd	v3,v1,v2	# Vector-vector add
vst	v3,x6	# Store the sum
vdisable		# Disable vector regs
- 8 instructions, 258 for RV64V (scalar code)



Copyright © 2019, Elsevier Inc. All rights Reserved

6

Vector Execution Time

Vector Architectures

- Execution time depends on three factors:
 - Length of operand vectors
 - Structural hazards
 - Data dependencies
- RV64V functional units consume one element per clock cycle
 - Execution time is approximately the vector length
- *Convey*
 - Set of vector instructions that could potentially execute together



Copyright © 2019, Elsevier Inc. All rights Reserved

7

Chimes

Vector Architectures

- Sequences with read-after-write dependency hazards placed in same convey via *chaining*
- *Chaining*
 - Allows a vector operation to start as soon as the individual elements of its vector source operand become available
- *Chime*
 - Unit of time to execute one convey
 - m conveys executes in m chimes for vector length n
 - For vector length of n , requires $m \times n$ clock cycles



Copyright © 2019, Elsevier Inc. All rights Reserved

8

Example

```
vld    v0,x5      # Load vector X
vmul   v1,v0,f0    # Vector-scalar multiply
vld    v2,x6      # Load vector Y
vadd   v3,v1,v2    # Vector-vector add
vst    v3,x6      # Store the sum
```

Convoys:

```
1      vld          vmul
2      vld          vadd
3      vst
```

3 chimes, 2 FP ops per result, cycles per FLOP = 1.5

For 64 element vectors, requires $32 \times 3 = 96$ clock cycles



Copyright © 2019, Elsevier Inc. All rights Reserved

9

Vector Architectures

Challenges

- Start up time
 - Latency of vector functional unit
 - Assume the same as Cray-1
 - Floating-point add => 6 clock cycles
 - Floating-point multiply => 7 clock cycles
 - Floating-point divide => 20 clock cycles
 - Vector load => 12 clock cycles
- Improvements:
 - > 1 element per clock cycle
 - Non-64 wide vectors
 - IF statements in vector code
 - Memory system optimizations to support vector processors
 - Multiple dimensional matrices
 - Sparse matrices
 - Programming a vector computer



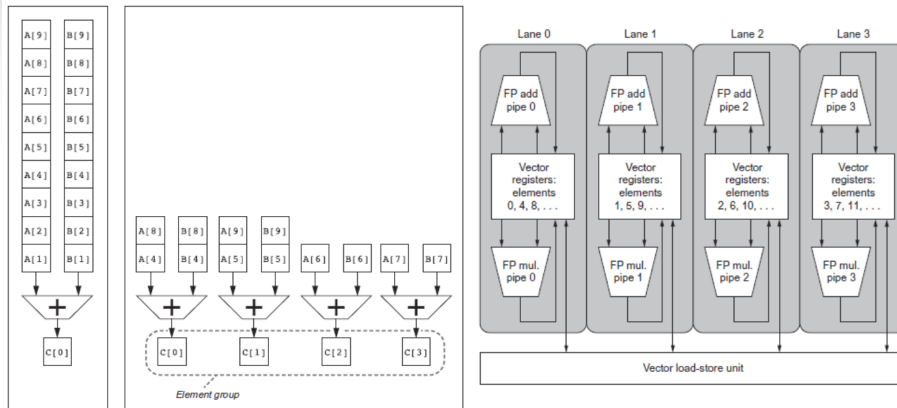
Copyright © 2019, Elsevier Inc. All rights Reserved

10

Vector Architectures

Multiple Lanes

- Element n of vector register A is “hardwired” to element n of vector register B
 - Allows for multiple hardware lanes



Copyright © 2019, Elsevier Inc. All rights Reserved

11

Vector Length Register

```

for (i=0; i < n; i=i+1) Y[i] = a * X[i] + Y[i];
    vsetdcfg 2 DP FP # Enable 2 64b Fl.Pt. registers
    fld f0,a          # Load scalar a
loop: setvl t0,a0      # vl = t0 = min(mvl,n)
    vld v0,x5          # Load vector X
    slli t1,t0,3       # t1 = vl * 8 (in bytes)
    add x5,x5,t1       # Increment pointer to X by vl*8
    vmul v0,v0,f0      # Vector-scalar mult
    vld v1,x6          # Load vector Y
    vadd v1,v0,v1      # Vector-vector add
    sub a0,a0,t0       # n -= vl (t0)
    vst v1,x6          # Store the sum into Y
    add x6,x6,t1       # Increment pointer to Y by vl*8
    bnez a0,loop      # Repeat if n != 0
    vdisable          # Disable vector regs
  
```



Copyright © 2019, Elsevier Inc. All rights Reserved

12

Vector Mask Registers

Vector Architectures

- Consider:
 - for (i = 0; i < 64; i=i+1)
 - if (X[i] != 0)
 - X[i] = X[i] – Y[i];
- Use predicate register to “disable” elements:

```

vsetdcfg      2*FP64    # Enable 2 64b FP vector regs
vsetpcfgi     1         # Enable 1 predicate register
vld           v0,x5      # Load vector X into v0
vld           v1,x6      # Load vector Y into v1
fmv.d.x       f0,x0      # Put (FP) zero into f0
vpne          p0,v0,f0   # Set p0(i) to 1 if v0(i)!=f0
vsub          v0,v0,v1   # Subtract under vector mask
vst           v0,x5      # Store the result in X
vdisable      # Disable vector registers
vpdisable     # Disable predicate registers

```



Copyright © 2019, Elsevier Inc. All rights Reserved

13

Memory Banks

Vector Architectures

- Memory system must be designed to support high bandwidth for vector loads and stores
- Spread accesses across multiple banks
 - Control bank addresses independently
 - Load or store non sequential words (need independent bank addressing)
 - Support multiple vector processors sharing the same memory
- Example:
 - 32 processors, each generating 4 loads and 2 stores/cycle
 - Processor cycle time is 2.167 ns, SRAM cycle time is 15 ns
 - How many memory banks needed?
 - $32 \times (4+2) \times 15 / 2.167 = \sim 1330$ banks



Copyright © 2019, Elsevier Inc. All rights Reserved

14

Stride

- Consider:


```
for (i = 0; i < 100; i=i+1)
  for (j = 0; j < 100; j=j+1) {
    A[i][j] = 0.0;
    for (k = 0; k < 100; k=k+1)
      A[i][j] = A[i][j] + B[i][k] * D[k][j];
  }
```
- Must vectorize multiplication of rows of B with columns of D
- Use *non-unit stride*
- Bank conflict (stall) occurs when the same bank is hit faster than bank busy time:
 - $\#banks / LCM(stride, \#banks) < \text{bank busy time}$



Scatter-Gather

- Consider:


```
for (i = 0; i < n; i=i+1)
  A[K[i]] = A[K[i]] + C[M[i]];
```
- Use index vector:

vsetdcfg	4*FP64	# 4 64b FP vector registers
vld	v0, x7	# Load K[]
vldx	v1, x5, v0	# Load A[K[]]
vld	v2, x28	# Load M[]
vldi	v3, x6, v2	# Load C[M[]]
vadd	v1, v1, v3	# Add them
vstx	v1, x5, v0	# Store A[K[]]
vdisable		# Disable vector registers



Programming Vec. Architectures

Vector Architectures

- Compilers can provide feedback to programmers
- Programmers can provide hints to compiler

Benchmark name	Operations executed in vector mode, compiler-optimized	Operations executed in vector mode, with programmer aid	Speedup from hint optimization
BDNA	96.1%	97.2%	1.52
MG3D	95.1%	94.5%	1.00
FLO52	91.5%	88.7%	N/A
ARC3D	91.1%	92.0%	1.01
SPEC77	90.3%	90.4%	1.07
MDG	87.7%	94.2%	1.49
TRFD	69.8%	73.7%	1.67
DYFESM	68.8%	65.6%	N/A
ADM	42.9%	59.6%	3.60
OCEAN	42.8%	91.2%	3.92
TRACK	14.4%	54.6%	2.52
SPICE	11.5%	79.9%	4.06
QCD	4.2%	75.1%	2.15



Copyright © 2019, Elsevier Inc. All rights Reserved

17

SIMD Extensions

SIMD Instruction Set Extensions for Multimedia

- Media applications operate on data types narrower than the native word size
 - Example: disconnect carry chains to “partition” adder
- Limitations, compared to vector instructions:
 - Number of data operands encoded into op code
 - No sophisticated addressing modes (strided, scatter-gather)
 - No mask registers



Copyright © 2019, Elsevier Inc. All rights Reserved

18

SIMD Implementations

- Implementations:
 - Intel MMX (1996)
 - Eight 8-bit integer ops or four 16-bit integer ops
 - Streaming SIMD Extensions (SSE) (1999)
 - Eight 16-bit integer ops
 - Four 32-bit integer/fp ops or two 64-bit integer/fp ops
 - Advanced Vector Extensions (2010)
 - Four 64-bit integer/fp ops
 - AVX-512 (2017)
 - Eight 64-bit integer/fp ops
 - Operands must be consecutive and aligned memory locations



Copyright © 2019, Elsevier Inc. All rights Reserved

19

SIMD Instruction Set Extensions for Multimedia

Example SIMD Code

- Example DXPY:

fld	f0,a	# Load scalar a
splat.4D	f0,f0	# Make 4 copies of a
addi	x28,x5,#256	# Last address to load
Loop: fld.4D	f1,0(x5)	# Load X[i] ... X[i+3]
fmul.4D	f1,f1,f0	# a x X[i] ... a x X[i+3]
fld.4D	f2,0(x6)	# Load Y[i] ... Y[i+3]
fadd.4D	f2,f2,f1	# a x X[i]+Y[i]...
		# a x X[i+3]+Y[i+3]
fsd.4D	f2,0(x6)	# Store Y[i]... Y[i+3]
addi	x5,x5,#32	# Increment index to X
addi	x6,x6,#32	# Increment index to Y
bne	x28,x5,Loop	# Check if done




Copyright © 2019, Elsevier Inc. All rights Reserved

20

SIMD Instruction Set Extensions for Multimedia

Roofline Performance Model

- Basic idea:
 - Plot peak floating-point throughput as a function of arithmetic intensity
 - Ties together floating-point performance and memory performance for a target machine
- Arithmetic intensity
 - Floating-point operations per byte read



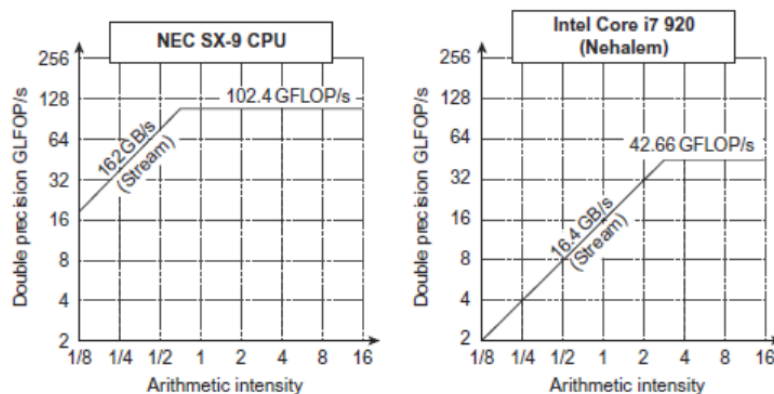
Copyright © 2019, Elsevier Inc. All rights Reserved

21

SIMD Instruction Set Extensions for Multimedia

Examples

- Attainable GFLOPs/sec = (Peak Memory BW × Arithmetic Intensity, Peak Floating Point Perf.)



Copyright © 2019, Elsevier Inc. All rights Reserved

22

SIMD Instruction Set Extensions for Multimedia

Graphical Processing Units

Graphical Processing Units

- Basic idea:
 - Heterogeneous execution model
 - CPU is the *host*, GPU is the *device*
 - Develop a C-like programming language for GPU
 - Unify all forms of GPU parallelism as *CUDA thread*
 - Programming model is “Single Instruction Multiple Thread”



Copyright © 2019, Elsevier Inc. All rights Reserved

23

Threads and Blocks

Graphical Processing Units

- A thread is associated with each data element
- Threads are organized into blocks
- Blocks are organized into a grid
- GPU hardware handles thread management, not applications or OS



Copyright © 2019, Elsevier Inc. All rights Reserved

24

NVIDIA GPU Architecture

Graphical Processing Units

- Similarities to vector machines:
 - Works well with data-level parallel problems
 - Scatter-gather transfers
 - Mask registers
 - Large register files
- Differences:
 - No scalar processor
 - Uses multithreading to hide memory latency
 - Has many functional units, as opposed to a few deeply pipelined units like a vector processor



Copyright © 2019, Elsevier Inc. All rights Reserved

25

Example

Graphical Processing Units

- Code that works over all elements is the grid
- Thread blocks break this down into manageable sizes
 - 512 threads per block
- SIMD instruction executes 32 elements at a time
- Thus grid size = 16 blocks
- Block is analogous to a strip-mined vector loop with vector length of 32
- Block is assigned to a multithreaded SIMD processor by the thread block scheduler
- Current-generation GPUs have 7-15 multithreaded SIMD processors



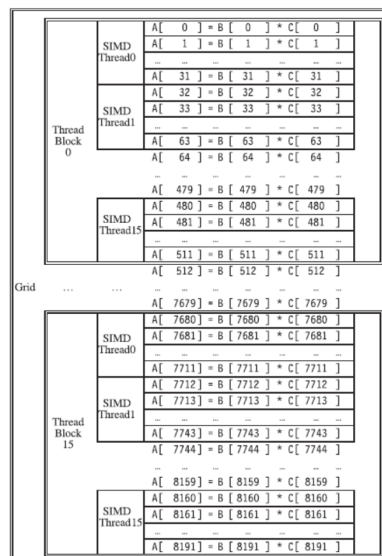
Copyright © 2019, Elsevier Inc. All rights Reserved

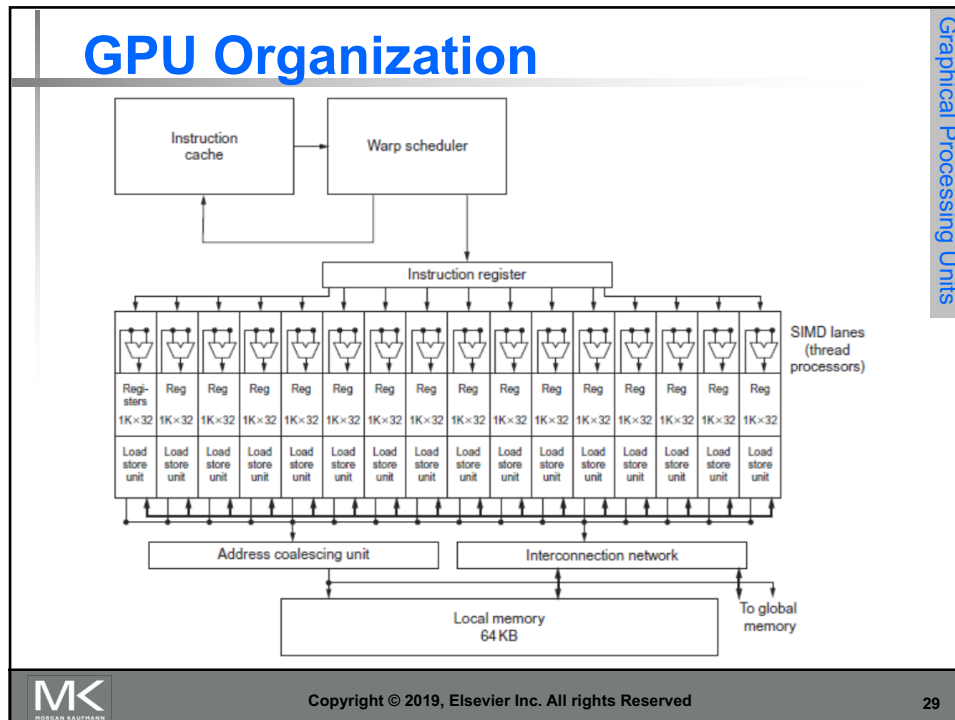
26

Terminology

- Each thread is limited to 64 registers
- Groups of 32 threads combined into a SIMD thread or “warp”
 - Mapped to 16 physical lanes
- Up to 32 warps are scheduled on a single SIMD processor
 - Each warp has its own PC
 - Thread scheduler uses scoreboard to dispatch warps
 - By definition, no data dependencies between warps
 - Dispatch warps into pipeline, hide memory latency
- Thread block scheduler schedules blocks to SIMD processors
- Within each SIMD processor:
 - 32 SIMD lanes
 - Wide and shallow compared to vector processors

Example





NVIDIA Instruction Set Arch.

- ISA is an abstraction of the hardware instruction set
 - “Parallel Thread Execution (PTX)”
 - opcode.type d,a,b,c;
 - Uses virtual registers
 - Translation to machine code is performed in software
 - Example:


```
shl.s32      R8, blockIdx, 9      ; Thread Block ID * Block size (512 or 29)
add.s32      R8, R8, threadIdx    ; R8 = i = my CUDA thread ID
ld.global.f64 RD0, [X+R8]        ; RD0 = X[i]
ld.global.f64 RD2, [Y+R8]        ; RD2 = Y[i]
mul.f64 RD0, RD0, RD4             ; Product in RD0 = RD0 * RD4 (scalar a)
add.f64 RD0, RD0, RD2             ; Sum in RD0 = RD0 + RD2 (Y[i])
st.global.f64 [Y+R8], RD0        ; Y[i] = sum (X[i]*a + Y[i])
```

Graphical Processing Units

Copyright © 2019, Elsevier Inc. All rights Reserved

30

Conditional Branching

- Like vector architectures, GPU branch hardware uses internal masks
- Also uses
 - Branch synchronization stack
 - Entries consist of masks for each SIMD lane
 - I.e. which threads commit their results (all threads execute)
 - Instruction markers to manage when a branch diverges into multiple execution paths
 - Push on divergent branch
 - ...and when paths converge
 - Act as barriers
 - Pops stack
- Per-thread-lane 1-bit predicate register, specified by programmer



Example

```

if (X[i] != 0)
    X[i] = X[i] - Y[i];
else X[i] = Z[i];

ld.global.f64    RD0, [X+R8]           ; RD0 = X[i]
setp.neq.s32     P1, RD0, #0           ; P1 is predicate register 1
@!P1, bra        ELSE1, *Push          ; Push old mask, set new mask bits
                                           ; if P1 false, go to ELSE1

ld.global.f64    RD2, [Y+R8]           ; RD2 = Y[i]
sub.f64          RD0, RD0, RD2          ; Difference in RD0
st.global.f64    [X+R8], RD0           ; X[i] = RD0
@P1, bra         ENDIF1, *Comp          ; complement mask bits
                                           ; if P1 true, go to ENDIF1

ELSE1:           ld.global.f64 RD0, [Z+R8] ; RD0 = Z[i]
                  st.global.f64 [X+R8], RD0 ; X[i] = RD0
ENDIF1: <next instruction>, *Pop        ; pop to restore old mask
  
```



NVIDIA GPU Memory Structures

Graphical Processing Units

- Each SIMD Lane has private section of off-chip DRAM
 - “Private memory”
 - Contains stack frame, spilling registers, and private variables
- Each multithreaded SIMD processor also has local memory
 - Shared by SIMD lanes / threads within a block
- Memory shared by SIMD processors is GPU Memory
 - Host can read and write GPU memory



Copyright © 2019, Elsevier Inc. All rights Reserved

33

Pascal Architecture Innovations

Graphical Processing Units

- Each SIMD processor has
 - Two or four SIMD thread schedulers, two instruction dispatch units
 - 16 SIMD lanes (SIMD width=32, chime=2 cycles), 16 load-store units, 4 special function units
 - Two threads of SIMD instructions are scheduled every two clock cycles
- Fast single-, double-, and half-precision
- High Bandwidth Memory 2 (HBM2) at 732 GB/s
- NVLink between multiple GPUs (20 GB/s in each direction)
- Unified virtual memory and paging support



Copyright © 2019, Elsevier Inc. All rights Reserved

34

Pascal Multithreaded SIMD Proc.



Copyright © 2019, Elsevier Inc. All rights Reserved

35

Vector Architectures vs GPUs

- SIMD processor analogous to vector processor, both have MIMD
- Registers
 - RV64V register file holds entire vectors
 - GPU distributes vectors across the registers of SIMD lanes
 - RV64 has 32 vector registers of 32 elements (1024)
 - GPU has 256 registers with 32 elements each (8K)
 - RV64 has 2 to 8 lanes with vector length of 32, chime is 4 to 16 cycles
 - SIMD processor chime is 2 to 4 cycles
 - GPU vectorized loop is grid
 - All GPU loads are gather instructions and all GPU stores are scatter instructions



Copyright © 2019, Elsevier Inc. All rights Reserved

36

SIMD Architectures vs GPUs

Graphical Processing Units

- GPUs have more SIMD lanes
- GPUs have hardware support for more threads
- Both have 2:1 ratio between double- and single-precision performance
- Both have 64-bit addresses, but GPUs have smaller memory
- SIMD architectures have no scatter-gather support



Copyright © 2019, Elsevier Inc. All rights Reserved

37

Loop-Level Parallelism

Detecting and Enhancing Loop-Level Parallelism

- Focuses on determining whether data accesses in later iterations are dependent on data values produced in earlier iterations
 - Loop-carried dependence
- Example 1:
for (i=999; i>=0; i=i-1)
 x[i] = x[i] + s;
- No loop-carried dependence



Copyright © 2019, Elsevier Inc. All rights Reserved

38

Loop-Level Parallelism

Detecting and Enhancing Loop-Level Parallelism

- Example 2:

```
for (i=0; i<100; i=i+1) {
    A[i+1] = A[i] + C[i]; /* S1 */
    B[i+1] = B[i] + A[i+1]; /* S2 */
}
```

- S1 and S2 use values computed by S1 in previous iteration
- S2 uses value computed by S1 in same iteration



Copyright © 2019, Elsevier Inc. All rights Reserved

39

Loop-Level Parallelism

Detecting and Enhancing Loop-Level Parallelism

- Example 3:

```
for (i=0; i<100; i=i+1) {
    A[i] = A[i] + B[i]; /* S1 */
    B[i+1] = C[i] + D[i]; /* S2 */
}
```

- S1 uses value computed by S2 in previous iteration but dependence is not circular so loop is parallel
- Transform to:


```
A[0] = A[0] + B[0];
for (i=0; i<99; i=i+1) {
    B[i+1] = C[i] + D[i];
    A[i+1] = A[i+1] + B[i+1];
}
B[100] = C[99] + D[99];
```



Copyright © 2019, Elsevier Inc. All rights Reserved

40

Loop-Level Parallelism

Detecting and Enhancing Loop-Level Parallelism

- Example 4:


```
for (i=0;i<100;i=i+1) {
    A[i] = B[i] + C[i];
    D[i] = A[i] * E[i];
}
```
- Example 5:


```
for (i=1;i<100;i=i+1) {
    Y[i] = Y[i-1] + Y[i];
}
```



Copyright © 2019, Elsevier Inc. All rights Reserved

41

Finding dependencies

Detecting and Enhancing Loop-Level Parallelism

- Assume indices are affine:
 - $a \times i + b$ (i is loop index)
- Assume:
 - Store to $a \times i + b$, then
 - Load from $c \times i + d$
 - i runs from m to n
 - Dependence exists if:
 - Given j, k such that $m \leq j \leq n, m \leq k \leq n$
 - Store to $a \times j + b$, load from $a \times k + d$, and $a \times j + b = c \times k + d$



Copyright © 2019, Elsevier Inc. All rights Reserved

42

Finding dependencies

Detecting and Enhancing Loop-Level Parallelism

- Generally cannot determine at compile time
- Test for absence of a dependence:
 - GCD test:
 - If a dependency exists, $\text{GCD}(c,a)$ must evenly divide $(d-b)$
- Example:

```
for (i=0; i<100; i=i+1) {  
    X[2*i+3] = X[2*i] * 5.0;  
}
```



Copyright © 2019, Elsevier Inc. All rights Reserved

43

Finding dependencies

Detecting and Enhancing Loop-Level Parallelism

- Example 2:

```
for (i=0; i<100; i=i+1) {  
    Y[i] = X[i] / c; /* S1 */  
    X[i] = X[i] + c; /* S2 */  
    Z[i] = Y[i] + c; /* S3 */  
    Y[i] = c - Y[i]; /* S4 */  
}
```
- Watch for antidependencies and output dependencies



Copyright © 2019, Elsevier Inc. All rights Reserved

44

Finding dependencies

Detecting and Enhancing Loop-Level Parallelism

- Example 2:


```
for (i=0; i<100; i=i+1) {
    Y[i] = X[i] / c; /* S1 */
    X[i] = X[i] + c; /* S2 */
    Z[i] = Y[i] + c; /* S3 */
    Y[i] = c - Y[i]; /* S4 */
}
```
- Watch for antidependencies and output dependencies



Copyright © 2019, Elsevier Inc. All rights Reserved

45

Reductions

Detecting and Enhancing Loop-Level Parallelism

- Reduction Operation:


```
for (i=9999; i>=0; i=i-1)
    sum = sum + x[i] * y[i];
```
- Transform to...


```
for (i=9999; i>=0; i=i-1)
    sum [i] = x[i] * y[i];
for (i=9999; i>=0; i=i-1)
    finalsum = finalsum + sum[i];
```
- Do on p processors:


```
for (i=999; i>=0; i=i-1)
    finalsum[p] = finalsum[p] + sum[i+1000*p];
```
- Note: assumes associativity!



Copyright © 2019, Elsevier Inc. All rights Reserved

46

Fallacies and Pitfalls

Fallacies and Pitfalls

- GPUs suffer from being coprocessors
 - GPUs have flexibility to change ISA
- Concentrating on peak performance in vector architectures and ignoring start-up overhead
 - Overheads require long vector lengths to achieve speedup
- Increasing vector performance without comparable increases in scalar performance
- You can get good vector performance without providing memory bandwidth
- On GPUs, just add more threads if you don't have enough memory performance



Copyright © 2019, Elsevier Inc. All rights Reserved

47