

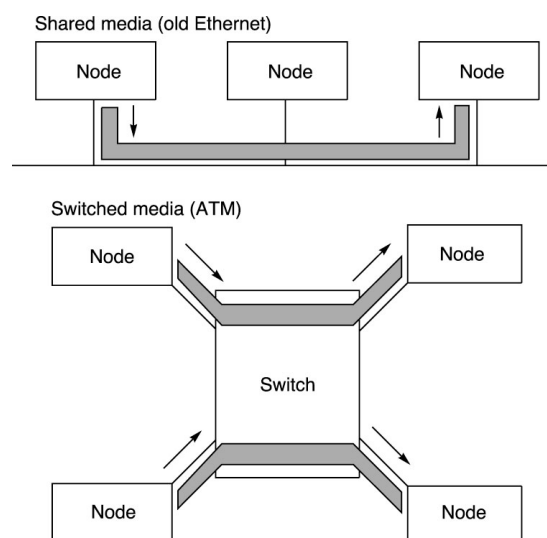
CSC 631: High-Performance Computer Architecture

Spring 2017

Lecture 12: Many-core & Interconnect

Interconnection Networks

- Classification: Shared Medium or Switched



Shared Media Networks

- Need arbitration to decide who gets to talk
- Arbitration can be centralized or distributed
- Centralized not used much for networks
 - Special arbiter device (or must elect arbiter)
 - Good performance if arbiter far away? Nah.
- Distributed arbitration
 - Check if media already used (carrier sensing)
 - If media not used now, start sending
 - Check if another also sending (collision detection)
 - If collision, wait for a while and retry
 - “For a while” is random (otherwise collisions repeat forever)
 - Exponential back-off to avoid wasting bandwidth on collisions

Switched Networks

- Need switches
 - Introduces switching overheads
- No time wasted on arbitration and collisions
- Multiple transfers can be in progress
 - If they use different links, of course
- Circuit or Packet Switching
 - Circuit switching: end-to-end connections
 - Reserves links for a connection (e.g. phone network)
 - Packet switching: each packet routed separately
 - Links used only when data transferred (e.g. Internet Protocol)

Routing

- Shared media has trivial routing (broadcast)
- In switched media we can have
 - Source-based (source specifies route)
 - When connection made, set up route
 - Switches forward packets along the route
 - Destination-based (source specifies destination)
 - Switches must route packet toward destination
- Also can be classified into
 - Deterministic (one route from a source to a destination)
 - Adaptive (different routes can be used)

Routing Methods for Switches

- Store-and-Forward
 - Switch receives entire packet, then forwards it
 - If error occurs when forwarding, switch can re-send
- Wormhole routing
 - Packet consists of flits (a few bytes each)
 - First flit contains header w/ destination address
 - Switch gets header, decides where to forward
 - Other flits forwarded as they arrive
 - Looks like packet worming through network
 - If an error occurs along the way, sender must re-send
 - No switch has the entire packet to re-send it

Cut-Through Routing

- What happens when link busy?
 - Header arrives to switch, but outgoing link busy
 - What do we do with the other flits of the packet?
- Wormhole routing: stop the tail when head stops
 - Now each flit along the way blocks the a link
 - One busy link creates other busy links => traffic jam
- Cut-Through Routing
 - If outgoing link busy, receive and buffer incoming flits
 - The buffered flits stay there until link becomes free
 - When link free, the flits start worming out of the switch
 - Need packet-sized buffer space in each switch
 - Wormhole Routing switch needs to buffer only one flit

Routing: Network Latency

- Switch Delay
 - Time from incoming to outgoing link in a switch
- Switches
 - Number of switches along the way
- Transfer time
 - Time to send the packet through a link
- Store-and-Forward end-to-end transfer time
 - $(\text{Switches} * \text{SwitchDelay}) + (\text{TransferTime} * (\text{Switches} + 1))$
- Wormhole or Cut-Through end-to-end transfer time
 - $(\text{Switches} * \text{SwitchDelay}) + \text{TransferTime}$
 - Much better if there are many switches along the way
 - See the example on page 811

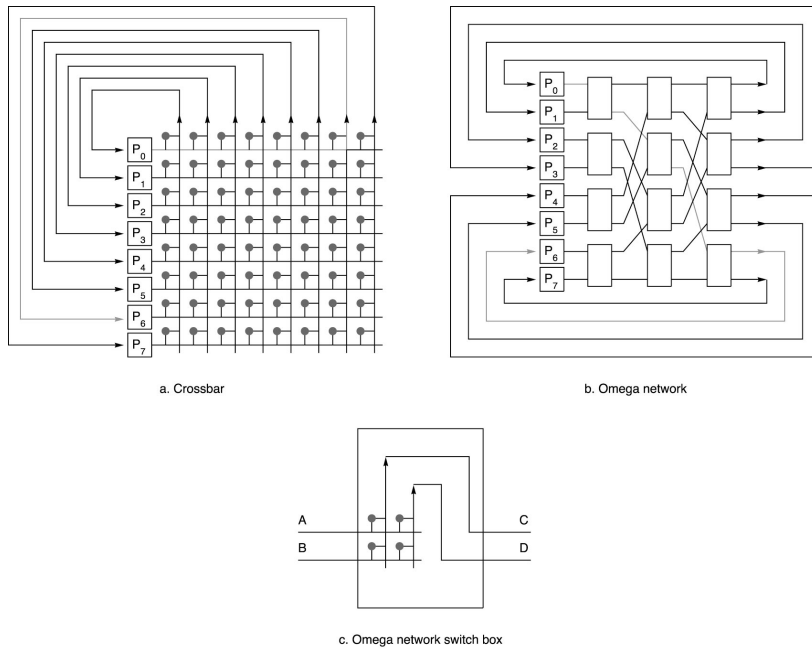
Switch Technology

- What do we want in a switch
 - Many input and output links
 - Usually number of input and output links the same
 - Low contention inside the switch
 - Best if there is none (only external links cause contention)
 - Short switching delay
- Crossbar
 - Very low switching delay, no internal contention
 - Complexity grows as square of number of links
 - Can not have too many links (e.g. up to 64 in and 64 out)

Switch Technology

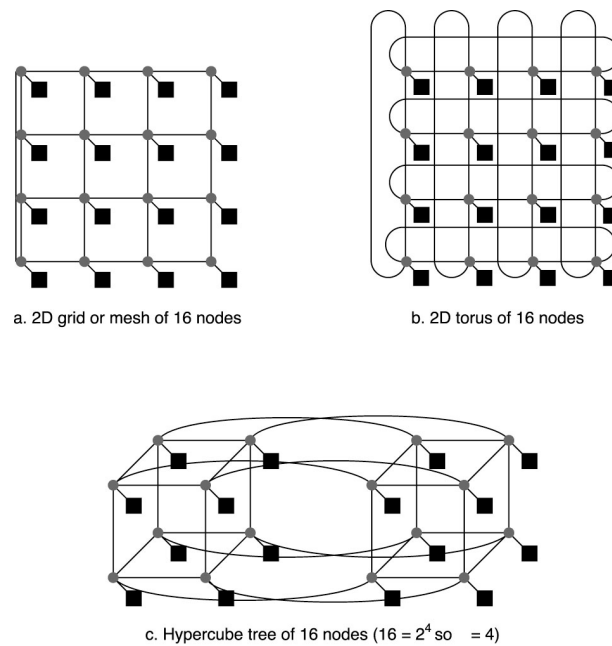
- What do we want in a switch
 - Many input and output links
 - Usually number of input and output links the same
 - Low contention inside the switch
 - Best if there is none (only external links cause contention)
 - Short switching delay
- Crossbar
 - Very low switching delay, no internal contention
 - Complexity grows as square of number of links
 - Can not have too many links (e.g. up to 64 in and 64 out)
- Omega Network
 - Build switches with more ports using small crossbars
 - Lower complexity per link, but longer delay and more contention

Switch Technology



© 2003 Elsevier Science (USA). All rights reserved.

Network Topology



© 2003 Elsevier Science (USA). All rights reserved.

Network Topology

- What do we want in a network topology
 - Many nodes, high bandwidth, low contention, low latency
 - Low latency: few switches along any route
 - For each (src, dest) pair, we choose shortest route
 - Longest such route over all (src,dst) pairs: *network diameter*
 - We want networks with small diameter!
 - Low contention: high aggregate bandwidth
 - Divide network into two groups, each with half the nodes
 - Total bandwidth between groups is *bisection bandwidth*
 - Actually, we use the minimum over all such bisections

On-Chip Networks

- We'll have many cores on-chip
 - Need switched network to provide bandwidth
- Need to map well onto chip surface
 - E.g. hypercube is not great
- Mesh or grid should work well, torus OK too
 - Limited ports per switch (CPU & 4 neighbors)
 - All links short (going to neighbors)
 - Many parallel algorithms map well onto grids
 - Matrices, grids, etc.

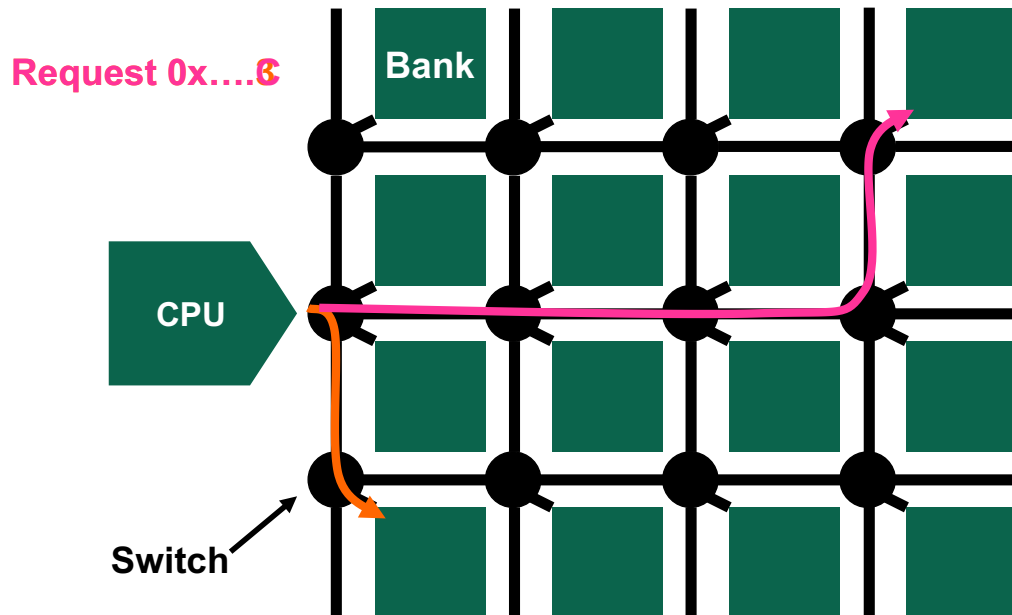
Trouble with shared caches

- Private caches OK
 - Each placed with its own processor
- We want a shared cache, too
 - Fits more data than if broken into private caches
 - Private caches replicate data
 - Dynamically shared
 - Threads that need more space get more space
- But how do we make a shared cache fast

Trouble with shared caches

- Private caches OK
 - Each placed with its own processor
- We want a shared cache, too
 - Fits more data than if broken into private caches
 - Private caches replicate data
 - Dynamically shared
 - Threads that need more space get more space
- But how do we make a shared cache fast

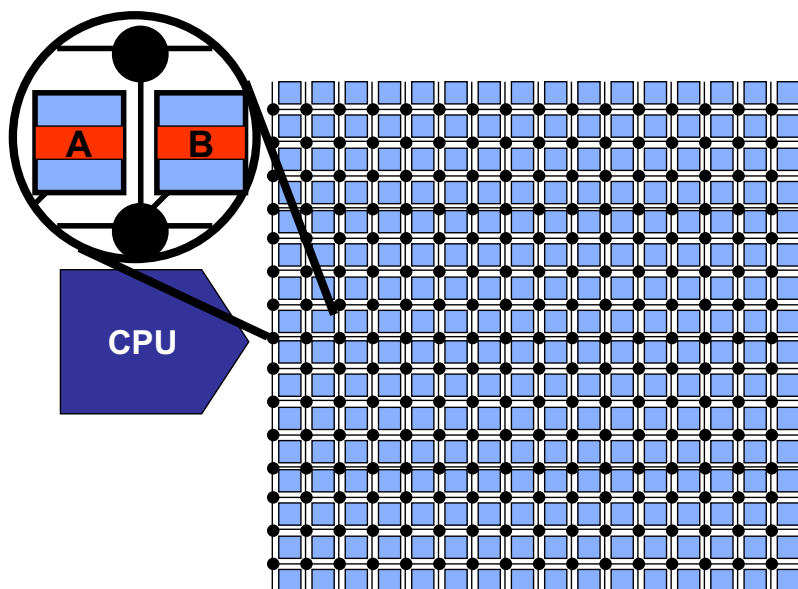
Non-uniform Cache Arch. (NUCA)



S-NUCA Performance

- Fast access to nearby banks
- Slow access to far-away banks
- Average better than worst-case

D-NUCA Solution



D-NUCA Performance

- Fast access to nearby banks
- Slow access to far-away banks
- Average much better than worst-case
- But we keep moving blocks
 - Lots of power-hungry activity
- Need smart policies for block migration
 - Move blocks less frequently
 - But get most of the benefit of being able to move

D-NUCA Issues

- Blocks keep moving, how do we find them?
- One solution: Use an on-chip directory!
 - Use direct mapping to assign a home bank
 - If we don't know where the block is, ask the home bank
 - If we move the block, tell the home bank
 - If we think we know where the block is, look there. If it's been moved, ask home bank